

# MPI 2.2

William Gropp



## Scope of MPI 2.2

- Small changes to the standard. A small change is defined as one that does not break existing *correct* MPI 2.0 or 2.1 user code - either by interface changes or semantic changes - and does not require large implementation changes.
  - ♦ What this often means is *adding* something to MPI, such as a new routine or a new basic datatype. Resolving ambiguities can be considered (but don't break programs)
- Note that some topics may be moved to the MPI 2.1 (obvious fix) or MPI 3 discussions (more than minor change)



# Process

- Discussion
  - ♦ Formation of ad hoc working group to create the ...
- Written Proposal
  - ♦ Must identify specific pages/lines
  - ♦ Must be comprehensive (all relevant lines, including bindings)
  - ♦ Should identify possible impact on *incorrect* MPI programs
  - ♦ Readings and votes according to rules
- Adopted into MPI 2.1 document to create MPI 2.2 document



3

## Topics for MPI 2.2 From Current Errata and Mail Discussions

- Language bindings
  - ♦ Use of const
  - ♦ Use of restrict
  - ♦ New C datatypes
  - ♦ Java? Fortran 2003?
  - ♦ Info constructor in C++
  - ♦ Missing C++ null functions (e.g., `CONVERSION_FN_NULL`)
- Extensions of Routines
  - ♦ `MPI_Alltoallx` (generalize `MPI_Alltoallv`)
  - ♦ `MPI_Alloc_mem` behavior on failure
  - ♦ Constant blocksize version of `reduce_scatter`
  - ♦ Messages larger than 2GB (e.g., `int`->`MPI_Aint` for count?)
  - ♦ `MPI_IN_PLACE` for `Reduce_scatter`
- Extensions of "constants"
  - ♦ `MPI_ERR_INFO`?
  - ♦ C++ `SEEK_SET` etc and `stdio`
  - ♦ Last error code and last error class
- Open Questions
  - ♦ `MPI_GROUP_EMPTY`
  - ♦ `MPI_LONG_LONG_INT` and `MINLOC/MAXLOC`
  - ♦ Does `FILE_GET_VIEW` return copies of datatypes?
  - ♦ MPI Datatype for sending `MPI_Aint` and/or `MPI_Offset`
  - ♦ Datatypes and `MPI_Probe/Iprobe`
  - ♦ C++ and interlanguage compatibility
  - ♦ Send Buffer Access



4

## Additional Topics

- Topics Suggested on Monday
  - ♦ Globalization (UTF-8)
  - ♦ Wchar\_t strings
  - ♦ Secure API
  - ♦ Require C++ Namespace
  - ♦ Collective “v” ops with bound on count
  - ♦ Allow more MPI functions to be macros
  - ♦ Interaction with multicore/threads
  - ♦ Memory allocation
  - ♦ Profiling interface (part. Fortran vs C)
  - ♦ Mpiexec
  - ♦ Synchronized and aggregated print
  - ♦ MPI\_Offset type extents for datatypes



5

## Making Progress

- We need to resolve the open questions (even if that means forwarding them to MPI 3 discussions)
  - ♦ Need volunteers to take on open items
- We need to identify self-organizing working groups for larger projects
- Jeff Squyres has volunteered to write up his errata items
- I'll write up mine
- A history of proposals (including failed ones) will be maintained (but not part of the standard)
- We have substantial proposals (see errata page) for
  - ♦ Alltoallx
  - ♦ MPI\_PROC\_NULL and epochs
  - ♦ Rationale for no MPI\_IN\_PLACE in MPI\_Exscan
  - ♦ MPI\_IN\_PLACE for MPI\_Reduce\_scatter
  - ♦ Send buffer semantics
  - ♦ Adding const



6



## C++ complex type and interlanguage compatibility

- E.g., Send a complex in Fortran (or C) and receive it in C++
- Proposed text
  - ♦ MPI-2, page 276, after line 4, add
  - ♦ Advice to users.
  - ♦ Most but not all datatypes in each language have corresponding datatypes in other languages. For example, there is no C or Fortran counterpart to the MPI::BOOL or the the MPI::COMPLEX, MPI::DOUBLE\_COMPLEX, or MPI:LONG\_DOUBLE\_COMPLEX. End of advice to users.
- Extending the C++ datatypes to C and Fortran needs to include MPI::BOOL as well as the complex types, and should define what the equivalent types are in C and Fortran. The real issue here is the MPI:F\_COMPLEX and completing the list of such routines.



# MPI\_Alltoallx

- MPI-2, page 164, line 16-30 should read:  
7.3.5. Generalized All-to-all Functions  
One of the basic data movement operations needed in parallel signal processing is the 2-D matrix transpose. This operation has motivated two generalizations of the MPI\_ALLTOALLV function. These new collective operations are MPI\_ALLTOALLW and MPI\_ALLTOALLX; the ``W'' indicates that it is an extension to MPI\_ALLTOALLV, and ``X'' indicates that it is an extension to MPI\_ALLTOALLW.  
MPI\_ALLTOALLX is the most general form of All-to-all. Like MPI\_TYPE\_CREATE\_STRUCT, the most general type constructor, MPI\_ALLTOALLW and MPI\_ALLTOALLX allow separate specification of count, displacement and datatype. In addition, to allow maximum flexibility, the displacement of blocks within the send and receive buffers is specified in bytes. In MPI\_ALLTOALLW, these displacements are specified as integer arguments and in MPI\_ALLTOALLX they are specified as address integer.
- Rationale. The MPI\_ALLTOALLW function generalizes several MPI functions by carefully selecting the input arguments. For example, by making all but one process have sendcounts[i] = 0, this achieves an MPI\_SCATTERW function. MPI\_ALLTOALLX allows the usage of MPI\_BOTTOM as buffer argument and defining the different buffer location via the displacement arguments rather than only via different

