



MPI 2.1 at MPI Forum Chicago, April 28-30, 2008 Version 2.1 – Combined Document

Rolf Rabenseifner
rabenseifner@hlrs.de
(Chairman of MPI 2.1 Task)

University of Stuttgart
High-Performance Computing-Center Stuttgart (HLRS)
www.hlrs.de



MPI 2.1
Slide 1

Höchstleistungsrechenzentrum Stuttgart



MPI-2.1 – Working Plan

Monday:

- All reviewer are asked for their (formalized) reviews
- Discussions (only if needed)

Tuesday:

- Decisions (straw voted) on open issues from the discussion
- Summary = finishing the official reading



Where we are – a statistical summary

- MPI 1.3 + 2.0, merging 335 distinguished parts
- Ballot 1+2 40 items implemented
- Ballot 3 14 items / 12 accepted + implemented
- Ballot 4 22 items / 22 accepted + implemented
+ a lot of e-mails
- Reviews until Apr. 19 477 items analyzed
232 items implemented by myself
- 14 chapter authors 21 chapters / annexes / frontm. / bibl. / index
read and corrected - **great work!**
57 review items implemented
- I checked all modifications so far done by the chapter authors (“diff”)

It's time to make the cut !

**Version Apr. 19, 2008 + all YYYYYY modifications + TYPO-corrections
will be my final MPI-2.1**

(No chance that I can continue this process any longer)



Reviewer Question & Answer

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by

- MPI-1.1 (June 1995) and MPI-2 (July 1997),
- MPI-1 errata (Oct. 12, 1998),
- Ballot 1+2 (May 15, 2002),
- Ballot 3 (protocol from March 2008 meeting, 1st vote), and
- Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications

(they are already implemented in the current draft)

are accepted by the MPI Forum:

- XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting:

- YYYYYY.

Further comments (not official part of this review):

- ZZZZZ



We need reviewers for: (**bold=large**) : (**red=responsible chapter author**) (Reviewer, green=@meeting)

Frontmatter & Acknowledgements	Bill Gropp & Rich Graham, Rusty Lusk
Chap. 1: Introduction to MPI	Bill Gropp, Rusty Lusk, Karl Feind, Adam Moody, Traeff
Chap. 2: MPI-2 Terms and Conventions	Tony Skjellum, Bill Gropp, Richard Barrett, Traeff
Chap. 3: Point-to-Point Communication (incl. sections from MPI-2 Misc. + 8.9)	Rich Graham, Jespar Larsson Traeff, George Bosilca, Steve Poole, Kannan Narasimhan, David Solt, B. Gropp Matt Koop, Adam Moody
Chap. 4: Collective Communication (incl. sections from MPI-2 Ext. Collect.)	Adam Moody, Steven Ericsson-Zenith, Edgar Gabriel, R. Thakur, B. Gropp, G. Bosilca, Th. Hoefler, J. Traeff
Chap. 5: Groups, Context, and Communicators (incl. sections from MPI-2 Ext.Col. + 8.8)	Richard Treumann, Steven Ericsson-Zenith, Edgar Gabriel, Tony Skjellum, Bill Gropp, G. Bosilca, Robert Blackmore
Chap. 6: Process Topologies	Jesper L. Traeff, Rusty Lusk, Bill Gropp, Richard Barrett
Chap. 7: MPI Environmental Management (incl. sections from MPI-2 Misc.)	George Bosilca, Rich Graham, Jespar Larsson Traeff, Steve Poole, Kannan Narasimhan, David Solt, B. Gropp
Chap. 8: Miscellany	Jesper L. Traeff, Rich Graham, George Bosilca, Steve Poole, Kannan Narasimhan, B. Gropp
Chap. 9: Process Creation and Management	David Solt, Dries Kimpe, Rusty Lusk, George Bosilca, Bill Gropp, Kalem Karian,
Chap. 10: One-Sided Communication	Jespar Larsson Traeff, Ericsson-Zenith, Martin Schulz, Bill Gropp, Darius Buntinas,
Chap. 11: External Interfaces	Bronis de Supinski, Bill Gropp, Rainer Keller
Chap. 12: I/O	Rajeev Thakur, Joachim Worringer, Bill Gropp, Koziol
Chap. 13: Language Bindings	Jeff Squyres, Steve Poole, Purushotham Bangalore, Bill Gropp, Erez Haba, Alexander Supalov
Chap. 14: Profiling Interface	Bronis de Supinski, Bill Gropp, Jeff Brown
Chap. 15: Deprecated Functions	Rolf Rabenseifner
Bibliography	Bill Gropp, Rusty Lusk
Annex A Language Bindings	A. Supalov, J. Squyres, St. Poole, P. Bangalore, B. Gropp
Annex B Change Log / Indexes	Rolf Rabenseifner

Reviewer Statements

Question:

Do you accept the B.1 Item 33 modification listed on this slide?

Yes:

16

No:

0

Abstain:

11

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	State-ment is true	XXXXX (Modifications, to be voted by the Forum)	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
All	Rolf Rabenseifner	YES	<ul style="list-style-type: none"> B.1, Item 33. See also slide 9 Remove of 4.1.1 & 4.1.2 YES: , NO: , Abstain: → missing Refs on 462 & 486 Mod: Non-negative integers YES: 19, NO: 3, Abstain: 5 	Page ii line 14: date → Apr. 19 Plus all on the sheet "STILL TODO"	Thank you to all involved persons. Page XX line 41: NOT: >300 E-mail authors

Non-negative: MPI_Type_create_indexed_block: count, blocklength

MPI_Pack: incount,

+ additionally: outsize , but not position


MPI_unpack: insize

(outcount maybe handled by MPI-2.2)

MPI_Pack_size: incount

+ additionally: size

grey = not present in the April meeting, review statement per e-mail



Question:

Do you accept the modifications listed on this slide

And on the slide 6?

Yes:

19

No:

3

Abstain:

5

Chap 3(4): integer -> nonnegative integer

Count arguments in:

MPI_Recv (p.29), MPI_Bsend (p.39), MPI_Ssend (p.40), MPI_Rsend (p.40), MPI_Isend (p.49), MPI_Ibsend (p.50), MPI_Issend (p.50), MPI_Irsend (p.51), MPI_Irecv (p.51), MPI_Send_init (p.68), MPI_Bsend_init (p.69), MPI_Ssend_init (p.69), MPI_Rsend_init (p.70), MPI_Recv_init (p.70), MPI_Sendrecv (p.73), MPI_Sendrecv_replace (p.74)

~~MPI_Get_count (p.32), MPI_Get_elements (p.100)~~

Editor's comment:

~~*NO* → OUT argument, MPI_UNDEFINED can be returned~~

Count argument in:

MPI_Waitany (p.57), MPI_Testany (p.58), MPI_Waitall (p.59), MPI_Testall (p.60), MPI_Waitsome (p.60), MPI_Testsome (p.61), MPI_Startall (p.71)

Rank arguments (do we want?) in:

~~MPI_Send (p.26), MPI_Recv (p.29), MPI_Bsend (p.39), MPI_Ssend (p.40),
MPI_Rsend (p.40), MPI_Isend (p.49), MPI_Ibsend (p.50), MPI_Issend
(p.50), MPI_Irsend (p.51), MPI_Irecv (p.51), MPI_Send_init (p.68),
MPI_Bsend_init (p.69), MPI_Ssend_init (p.69), MPI_Rsend_init (p.70),
MPI_Recv_init (p.70), MPI_Sendrecv (p.73), MPI_Sendrecv_replace (p.74)~~

Editor's comment:

No, because valid values MPI_ANY_SOURCE, MPI_PROC_NULL,
and MPI_ROOT may be implemented as negative integers.

Datatypes

- Proposal 1: (p<page>.<line> according to MPI-2.1 Draft Apr. 19, 2008)
 - P494.3: duplicate line 3-4 with
 - **MPI_CHAR – signed char**
 - **MPI::CHAR – char**
 - Reason: this is identical to the inconsistent definition in
 - MPI-1.1 page 28? (same as p28.2 – Table 3.2), and
 - MPI-2.0 page 271? (same as p454.5 – Table 16.1)
- Proposal 2:
 - P494.3 is kept, p454.5 is changed to both signed char
 - Reason: resolve inconsistency by keeping MPI-1 definition
- Proposal 3:
 - P494.3 and P28.2: char (instead of signed char)
- Results of voting:
 - Proposal 2 (3 votes) versus Proposal 3 (14 votes), 5 abstain
→ Proposal 2 is therefore removed.
 - Proposal 1 (13 votes) versus Proposal 3 (6 votes), 3 abstain

Should go to
MPI-2.2.

MPI-2.2 should
revisit all the
“signed”
datatypes and
check whether
“signed” should
be removed.

Datatypes

- P494.4: (treated as printable character)
- P494.11+13: (treated as integral value)



Reviewer Statements

Question:

Do you accept the XXXXX modifications listed on this slide?

Yes:

No:

Abstain:

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	Statement is true	XXXXX (Modifications, to be voted by the Forum)	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
Frontmatter & Acknowl.	Bill Gropp				
	Rich Graham	YES			
	Rusty Lusk				
EVERYWHERE			Xxx_xxx: the “_” does not work with search and cut&paste in pdf files		



Reviewer Statements

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: VVVVV Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	Statement is true	XXXXX	YYYYY (Bugs, to be done)	ZZZZZ
Chapter 1 Introduction to MPI	Bill Gropp				
	Rusty Lusk				
	Bronis			P4.12 are solving → solved, upto → up to P4.18+19 It is expexted that .. Will be → The .. was, Straight → just	
	Adam Moody	YES		P3.21 0-copy → zero-copy P4.1 July → On July P6.34 delete sentence P6.47 + p7.20 MPI-2 → MPI Additional small ... → file	
	Jesper Traeff	YES		Email-file Message-passing with hypen	
	Rainer Keller	YES		Typos on email, p7.7 new sentence, see file	

Reviewer Statements

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	Statement is true	XXX XX	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
Chapter 2 MPI Terms and Conventions	Tony Skjellum				
	Bill Gropp				
	Richard Barrett				
	Jesper Traeff	YES		P9.14 MPI-2 → MPI Email-file	
	Bronis de Supinski	YES		P16.10 “but” → “, as documented in Chapter ..., but” P16.28 MPI-2 → MPI P16.38 Chapter 8 ---- P23.2+3 remove Chap7 of MPI-1 doc. p23.32 MPI-2 → MPI	

Reviewer Statements

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	Statement is true	XXXXX (Modifications, to be voted by the Forum)	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
Chapter 3 Point-to-Point Communication	Richard Graham	YES			
	Jesper Traeff	YES	See file jesper...Nonnegative.ppt	see email-file	
	George Bosilca	YES			
	Steve Poole				
	Kannan Narasimhan Kaushik	YES			

Reviewer Statements

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	Statement is true	XXXXX (Modifications, to be voted by the Forum)	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
Chapter 3 Point-to-Point Communication (Continued)	David Solt				
	Bill Gropp				
	Matt Koop	YES			
	Adam Moody	YES		P33.4-5 remove P33.31+37 all & ANY → MPI_{WAIT TREST}{ ALL SOME } P49.7 Objects → Request Objects P52.37 2 times “,” → “ ” Further issues, see file	

Reviewer Statements

Question:

Do you accept the XXXXX modifications listed on this slide?

Yes:

No:

Abstain:
(see slide non non-negatives)

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	State- ment is true	XXXXX (Modifications, to be voted by the Forum)	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
Chapter 4 Datatypes	Richard Graham	YES		See file	
	Jesper Traeff	YES	See file jesper...Nonnegative. ppt	See file traeff_chap.3+4 _mpi-2.1- reviews.txt	
	George Bosilca	YES			
	Steve Poole				
	Kannan Narasimhan Kaushik	(YES)		See file	

Reviewer Statements

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	Statement is true	XXXXX (Modifications, to be voted by the Forum)	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
Chapter 4 Datatypes (Continued)	David Solt				
	Bill Gropp				
	Matt Koop				
	Adam Moody				
	Rolf Rabenseifner	YES		Bring in again 4.1.1 & 4.1.2	



Reviewer Statements

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	Statement is true	XXXX X <small>(Modifications, to be voted by the Forum)</small>	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
Chapter 5 Collective Communication	Adam Moody	YES		p136.24-24: The content of root's buffer is copied to all other processes.	
	Edgar Gabriel				
	Rajeev Thakur	YES			
	Alexander Supalov			P167.37 space after (is unnecessary. Please check everywhere	
	George Bosilca	YES			
	Jeff Squyres			P168.21 see jsquyres-chap5_mpi- 2.1-review.txt	
	Torsten Hoefler	YES		See file	

Reviewer Statements

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	State -ment is true	XXXXX (Modifications, to be voted by the Forum)	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
Chapter 5 Collective Communication	Jesper Traeff	YES	<ul style="list-style-type: none"> • Add MPI_LONG_LONG(_INT) + MPI_UNSIGNED_LONG_ LONG to 162.12 	<ul style="list-style-type: none"> • Sec. → Section • section \ref → Section \ref • See file 	MAXLOC and LONG_LO NG in MPI-2.2



Reviewer Statements

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	Statement is true	XXXXX (Modifications, to be voted by the Forum)	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
Chapter 6 Groups, Contexts, Communicators, and Caching	Richard Treumann	YES (email)			
	Edgar Gabriel	YES		220.27 MPI_Comm_free 221.44-46 remove Title+Paragraph 222.7-9 remove "in MPI-1", 2 x "new" 223.17 remove "new functions that are replacements for the MPI-1"	
	Tony Skjellum				
	Bill Gropp				
	George Bosilca	YES			
	Robert Blackmore	YES			

Reviewer Statements

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	Statement is true	XXXXX (Modifications, to be voted by the Forum)	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
Chapter 7 Process Topologies	Jesper Traeff	YES			
	Rusty Lusk				
	Bill Gropp				
	Richard Barrett				



Reviewer Statements

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	State- ment is true	XXXXX (Modifications, to be voted by the Forum)	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
Chapter 8 MPI Environmental Management	George Bosilca	YES			
	Richard Graham	YES		See file (no mods)	
	Jesper Traeff				
	Steve Poole				
	Kannan Narasimhan Kaushik	YES		41.a-f,	
	David Solt				
	Bronis de Supinski			12.4 after 8.4	

Reviewer Statements

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	Statement is true	XXXXX (Modifications, to be voted by the Forum)	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
Chapter 9 Miscellany	Jesper Traeff	YES		Title Info Object	
	Richard Graham			P283.16 "Many MPI routines take the argument info"	
	George Bosilca	YES			
	Steve Poole				
	Kannan Narasim- han Kaushik	YES		41.f	
	Bill Gropp				

Reviewer Statements

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	Statement is true	XXXXX (Modifications, to be voted by the Forum)	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
Chapter 10 Process Creation and Management	David Solt				
	Rusty Lusk				
	George Bosilca	YES			
	Bill Gropp				
	Yann Kalemkarian				
	Dries Kimpe (by e-mail)	YES	See file	See file	See file

Reviewer Statements

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	Statement is true	XXXXX (Modifications, to be voted by the Forum)	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
Chapter 11 One-sided Communications	Jesper Traeff	YES			
	Bill Gropp				
	Darius Buntinas				
	Martin Schulz (by email)	YES		Some formatting 46.b-d	See file 46.a+e → MPI-2.2

Reviewer Statements

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	Statement is true	XXXXX (Modifications, to be voted by the Forum)	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
Chapter 12 External Interfaces	Bronis de Supinski	YES		See file	
	Bill Gropp				
	Rainer Keller	YES		file	



Reviewer Statements

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	Statement is true	XXXXX (Modifications, to be voted by the Forum)	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
Chapter 13 I/O	Rajeev Thakur	YES		42.a (p432.46: page 4.1.4 → page 85)	
	Bill Gropp				
	Quincey Koziol	YES		See file	



Reviewer Statements

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	Statement is true	XXXXX (Modifications, to be voted by the Forum)	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
Chapter 14 Profiling Interface	Bronis de Supinski	YES			
	Bill Gropp				
	Jeff Brown	YES			



Reviewer Statements

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	Statement is true	XXXXX (Modifications, to be voted by the Forum)	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
Chapter 15 Deprecated Functions	Rolf Rabenseifner	YES			
	Rainer Keller	YES		<ul style="list-style-type: none"> • p441, l38 -- language --> language • Many locations, see file 	
	Bill Gropp				
	Jeff Squyres	YES		Jeff about MPI_NULL_C OPY_FN ... (Jeff does the Latex for this)	

Reviewer Statements

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	State- ment is true	XXXXX (Modifications, to be voted by the Forum)	YYYYY (Bugs, to be done)	ZZZZZ (Further commen ts)
Chapter 16 Language Bindings	Jeff Squyres	YES		44.a (p479.2: MPI_THREAD_INIT → MPI_INIT_THREAD)	
	Steve Poole				
	Puri Bangalore	YES		file	
	Bill Gropp				
	Erez Haba	YES			
	Alexander Supalov	YES		Email-file	

Reviewer Statements

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	State- ment is true	XXXXX (Modifications, to be voted by the Forum)	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
Annex A Language Bindings Summary	Alexander Supalov	YES		40.d (p553.44) 40.e (p554.35) 40.f (p557)	
	Jeff Squyres	YES		See file	
	Steve Poole				
	Puri Bangalore	YES		A.4.20 removed (Jeff's idea)	
	Bill Gropp				

Additional modifications (from Jeff Squyres)

- Chap. 15 Deprecated Functions
 - Predefined Callback routines must be added (C and Fortran):
 - **MPI_NULL_COPY_FN**
 - **MPI_DUP_FN**
 - **MPI_NULL_DELETE_FN**
- A.1.2 split into (always C / Fortran / C++)
 - Types
 - Prototypes
 - **C++: with namespace notation**
 - Deprecated prototypes
- A2.x, A3.x, A4.x
 - Predefined callback routines must be also reported by the script that generates automatically these sections.
This applies to
 - **X=4 Groups, Contexts, Communicators, and Caching (title must be corrected)**
 - **X=14 Deprecated**
- Index sections: Some missing entries will be added by Jeff Squyres

Reviewer Statements

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Chapter	Reviewer	Statement is true	XXXXX (Modifications, to be voted by the Forum)	YYYYY (Bugs, to be done)	ZZZZZ (Further comments)
Annex B Change-Log	Rolf Rabenseifner	(YES, see slide 6)		Remove Ballot numbers (keep in .tex files)	
	Puri	YES		See file	
Bibliography	Bill Gropp			See MPI-1.3	
	Rusty Lusk				
Indexes	Rolf Rabenseifner	(YES, see slide 6)		<ul style="list-style-type: none"> • 44.I (list from Jeff must be fixed) • Rename: "MPI Constants and Predefined Handles Index" 	
	Bill Gropp				

Still TO DO = Must be implemented in final MPI-2.1

Rolf Rabenseifner, HLRS, Apr. 20, 2008

40.a OK In Frontmatter, several committee --> committee (twin t)

40.b OK p.27.10 / p.27.25: below --> in Table 3.1 / 3.2

40.c OK page ii line 14: date --> Apr. 19

Alexander Supalov, Intel, Apr. 22, 2008

Page numbers according to Draft April 19, 2008:

40.d OK 553/44.5-46.5

Overcomplicated sentence, should be split in two, with the first rephrased. Note that "<" and ">" may be represented differently in the standard sources.

Reads

The bindings prototype functions using the type <CLASS> rather than listing each function for every MPI class; the token <CLASS> can be replaced with valid MPI-2 class names, such as Group, Datatype, etc., except when noted.

Should read

The bindings prototype functions are using the type <CLASS> rather than listing each function for every MPI class. The token <CLASS> can be replaced with valid MPI-2 class names, such as Group, Datatype, etc., except when noted.

40.e OK 554/35

Reads "results", should read "result".

40.f OK 557/all

Unfortunate odd/even page layout - the right margin overlaps with the line numbers. Can be fixed by adding an empty page somewhere, or by another small shift of the prior main text.

Editor's comment:

May be solved with smaller font size, or move 17 mm left

Kannan Narasimhan, HP, Apr. 26, 2008

Page numbers according to Draft April 19, 2008:

Chapter 8 (MPI Environment Management)

41.a OK Pg 265, 8.3 Error Handling (Line 12): Remove the phrase "On the other hand" and reword the sentence as follows: C and C++ have distinct typedefs for user defined error handling callback functions that accept communicator, file, and window arguments.

41.b OK Pg 266, 8.3.1 Error Handlers for Communicators (Lines 32-36): Change all references to HANDLER_FUNCTION to {COMM,WIN,FILE}_ERRHANDLER_FN in the "Advice to users" para.

41.c OK Pg 268, 8.3.2 Error Handlers for Windows (Line [13] 9):
ADD: Creates an error handler that can be attached to a window object.
Editor's comment: okay, but must be added after line 9.

41.d OK Pg 269, 8.3.3 Error Handlers for Files (Line [24] 20):
ADD: Creates an error handler that can be attached to a file object.

41.e OK Pg 270, 8.3.4 Freeing Error handlers and Retrieving Error Strings (Line 27-28): Change the sentence to:
The error handler will be deallocated after all of the objects associated with it (communicator, window, or file) are deallocated.

Chapter 9 (Miscellany)

41.f OK Pg 283, 9.1 The Info Object (Line 16): Change the first sentence to:
Many of the routines in MPI take an info argument.
(i.e., MPI instead of MPI-2)

In the frontmatter:

- *NO* Integrate e-mail authors of Ballot 1-4 into credits.tex
Too many authors (>300) with partially small input to the standardization

In the backmatter:

- DONE Using all CONST:MPI.* entries in the Index to build a Constant Index
- TODO Using all TYPEDEF:MPI.* entries in the Index to build a Typedef Index
It may be possible to add also references to the C types mentioned at 486/39-48 and 523/14, 525/21-23, and 503/19-32, 507/40
- DONE Using all MPI_* entries in the Index to build the Function Index

MPI-1.3 and MPI-2.1:

- D1.3 Title in same format as on MPI-2.1
- TODO References MPI-2.1 [8] and [9] must be corrected
- D1.3 27.x must be implemented also in MPI-1.3 (the second "the" was removed)

Additional Review Files = Describing XXXXX and YYYYY, i.e., must be implemented in final MPI-2.1

- graham_chap.4+8_comments.txt
- hoefler_chap.5_mpi-2.1-review.txt
- jsquyres-chap5_mpi-2.1-review.txt
- jsquyres-for-rolf-2.1.txt
- kannan-chapter4--review.txt
- kannan-chapter8.txt
- kannan-chapter9.txt
- keller_chap.12_mpi-2.1-review.txt
- keller_chap.15_mpi-2.1-review.txt
- keller_chap.1_mpi-2.1-review.txt
- kimpe_chap.10-mpi-2.1-reviews.txt
- koziol-chapter13-mpi21.txt
- moody_chap.1_additional_comments.txt
- moody_chap.3_mpi-2.1-review.txt
- moody_chap.5-mpi-2.1-review.txt
- puri-changelog-mpi-2.1-review.txt
- puri-mpi2.1-chapter16-review.txt
- schulz_chap.11_mpi-2.1-review.txt
- supalov_chap.16_mpi-2.1-review.txt
- supinski_chap.12_mpi-2.1-review.txt
- traeff_chap.1+2_mpi-2.1-reviews.txt
- traeff_chap.3+4_mpi-2.1-reviews.txt
- traeff_chap.3+4_Nonnegative.ppt
- traeff_chap.5_mpi-2.1-reviews.txt



Review Summary

Chapter	#YES
Frontmatter & Ackn.	2
1. Introduction	4
2. Terms & Conventions	3
3. Point-to-Point Comm.	6
4. Datatypes	3+(1)
5. Collective Comm.	6
6. Groups, Contexts, ...	5
7. Process Topologies	2
8. Environmental Manag.	4
9. Miscellany	4

Chapter	#YES
10. Process Creation ...	3
11. One-Sided Comm.	3
12. External Interfaces	3
13. I/O	3
14. Profiling Interface	3
15. Deprecated Funct.	3
16. Language Bindings	6
A. Lang. Bind. Summary	4
B. Change-Log	2
Bibliography	1
Index	1



Reviewer Statements – final straw vote over total MPI-2.1

Chap. of MPI-2.1 Draft Apr. 19, 2008 represents fully and exactly (as I can overlook) the text that was given by MPI-1.1 (June 1995) and MPI-2 (July 1997), MPI-1 errata (Oct. 12, 1998), Ballot 1+2 (May 15, 2002), Ballot 3 (protocol from March 2008 meeting, 1st vote), and Ballot 4 (protocol from March 2008 meeting, official reading)

if the following text-modifications (they are already implemented in the current draft) are accepted by the MPI Forum: XXXXX

and if the following bugs will be corrected until first vote in June 2008 meeting: YYYYYY. Further comments (not official part of this review): ZZZZZ

Question:

Is the MPI-2.1 document, including the XXXXX modifications, and if all YYYYYY modifications are done, ready for 1st vote in the next meeting?

Yes:

27

No:

0

Abstain:

0

MPI 2
Slide 3



Next steps

- MPI-2.1 → implementing review items until May 30, 2008, still colored
(no further corrections, only items that are now on the table)
 - Reviewers should make a final check
 - Final version of MPI-2.1 – only black&white
June 13, 2008
 - July 1, 2008: 1st vote
 - Sep. 4, 2008: 2nd vote = done
- Text is technically finished
 - implementation can start now!
(you need not to wait until June 13!)



Implementation of MPI-2.1

- Do you **expect** that your company/organization is able to implement MPI-2.1 until EuroPVM/MPI 2008 (Sep. 7-10, 2008)?
- If no, please can you give an estimate.
(33 small issues that you should check in Annex B, Change-Log)

Library	Yes mm/yy	asked
MPICH2	YES	Rajeev Thakur
Open MPI		Jeff Squyres
IBM		Dick Treumann, Robert Blackmore
Sun Microsystems		Terry Dontje
Microsoft		Erez Haba

Library	Yes mm/yy	asked
Bull		Yann Kalemkarian
HP		Kannan Kaushik
Cray	Q1 2009	Mark Pagel
NEC	YES	Hubert Ritzdorf, Jesper Traeff
Intel		Alexander Suplaov, Keith Underwood
SiCortex	YES	David Gingold

H L R I S

Thanks a lot for all your cooperation!

MPI-2.1 standardization is now nearly done ☺

→ Have a lot of fun with MPI-2.2 and MPI-3.0



Annex: Slides from Mar. 2008 meeting



MPI 2.1
Slide 41

Rolf Rabenseifner
Höchstleistungsrechenzentrum Stuttgart

H L R I S 

MPI-2.1 – Working Plan

Monday:

- All reviewer are asked
 - to check whether all editor's decisions about their review are okay:
 - OK** = Obvious, typo, correction must be done, ...
or good idea, correction should be done, ...
 - ?** = should be decided by the MPI Forum
 - 22** = should be deferred to MPI-2.2.
It is mainly a suggestion to the content of MPI and not to the merging of MPI-1.1 with MPI-2.
 - Q?** = Question & answer only, no need for corrections
 - = Comment, can be ignored
 - *NO*** = Correction must or should not be done
 - OK-items = same opinion of editor and reviewer (~100 items)
→ expectation: no discussion need → should be okay for the forum
 - Especially the ***NO*** and **"?"** must be checked and discussed, probably off-line

Tuesday: Remaining **"?"** must be decided by the Forum, see **review.txt**


Open topics on the path to a fully merged MPI-2.1 (1)

- Chap. 1 – Intro chapter
 - Responsible: Bill Gropp (together with Rusty Lusk)
- Chap. Collectives – Intro Chapter
 - Responsible: Adam Moody
- Move MPI_LONG_LONG, MPI_WCHAR, ... into the original text
Remove the additional MPI-2 sections.
 - Responsible: Rolf Rabenseifner
- Move all deprecated functions into an Annex B
 - Responsible: Rolf Rabenseifner
- Remake the type_ub / type_lb stuff (text is needed, only functions are deprecated)
 - Responsible: Rolf Rabenseifner



Open topics on the path to a fully merged MPI-2.1 (2)

- Producing C++ interfaces that can be included into the chapters
 - Responsible: Jeff Squyres
- Including them into the chapters:
 - Responsible: Rolf Rabenseifner
- Merge of all Language Binding Annexes to one Annex A
 - Responsible: Alexander Supalov
- Extracting Const Index from current total Index
 - Responsible: Rolf Rabenseifner
- Extracting Typedef Index from current total Index
 - Responsible: Rolf Rabenseifner
- Extracting Function Index from current total Index
 - Responsible: Rolf Rabenseifner



Open topics on the path to a fully merged MPI-2.1 (3)

- Producing List of Figures and List of Tables
 - Responsible: Rolf Rabenseifner (trivial)
- Producing List of Examples
 - Responsible: Rainer Keller
- Include Ballots 1-4 into MPI-2.1
 - Responsible: Rolf Rabenseifner
- Include all the OK-marked items from the MPI-2.1 reviews into MPI-2.1
 - Responsible: Rolf Rabenseifner
- Change-Log
 - Responsible: Rolf Rabenseifner (Ballot 1-4)



Open topics on the path to a fully merged MPI-2.1 (4)

- Correcting the Examples – substituting deprecated functions
 - Responsible for examples in Fortran: David Solt
 - Additional Volunteers: Bill Gropp (Datatype)
 - Responsible for examples in C: George Bosilca
 - Additional Volunteers: Pavan Balati (Datatypes)
- Adding a full acknowledgement list in the Frontmatter:
 - Currently MPI-2.1
 - Will be extended to MPI-2.2 and MPI-3.0
 - For MPI 2.2: all institutions and persons who have contributed to January and March 2008 meetings
 - All persons who have contributed to the email lists in the Ballots 1-4 and the review 1.3 and 2.1 process
 - Responsible: Richard Graham




Timescale – MPI-2.1 – Goals

- I would like to have finished all in the next week or two?
- Now:
 - Open Review 2.1 items – see review.txt “?”
- Additional text-merge round with dedicated chapter authors.
Goal:
 - **Remove references to MPI-1 and MPI-2.**
 - **Substitute by MPI.**
 - **Necessary modifications to achieve this goal.**
 - Responsible authors – see next slide



Annex: Slides from Jan. 2008 meeting



MPI 2.1
Slide 48

Rolf Rabenseifner
Höchstleistungsrechenzentrum Stuttgart

H L R I S 

MPI 2.1 Goals

Scope of Effort:

- Clarification to the MPI standards document,
- resulting in a single document describing the full MPI 2.1 standard.
- This includes merging of documents, text corrections, and added clarifying text.

Working plan:

- MPI 1.1 + Chap. 3 of MPI-2 (Version 1.2 of MPI) + some errata will be combined to
→ **MPI 1.3**
- MPI 1.2.1 + rest of MPI-2 (MPI 2.0) will be combined to
→ **MPI 2.1 draft (without clarifications)**
- adopted MPI 2.1 Ballots 1&2 + new MPI 2.1 ballots 3&4 are combined to the
→ **Ballot 1-4 of MPI 2.1 adopted errata**
(with references still based on MPI 1.1 and MPI-2 documents)
- MPI 2.1 draft + MPI 2.1 adopted errata
→ **MPI 2.1**



Why combined documents?

The goals behind this combining of the documents have been already expressed in the MPI-1.1 standard:

"Sect. 1.2 Who should use this standard?

This standard is intended for use by all those who want to write portable message-passing programs in Fortran 77 and C.

This includes individual application programmers, developers of software designed to run on parallel machines, and creators of environments and tools. ..."

It is more efficient that the MPI Forum combines the documents once than every user of the MPI documents has to do this in his/her daily work based on the combination of MPI-1.1 and the several updating documents, i.e., MPI-2, and the future updates 2.1, 2.2,

Schedule based on official rules

Rules and Procedures

1. Here is a reminder of the traditional MPI voting rules, which have served us well. These rules have been extended to the email discussion of MPI erratas and have been applied to the errata ballots. We expect to adapt these rules, preserving their spirit, as we go forward.
2. One vote per organization
3. To vote, an organization must have been present at the last two MPI Forum meetings.
4. **Votes are taken twice, at separate meetings. Votes are preceded by a reading at an earlier meeting, to familiarize everyone with the issues.**
5. Measures pass on a simple majority.
6. Only items consistent with the charter can be considered.

From http://www.mpi-forum.org/mpi2_1/index.htm

For **MPI x.x combined documents**:

This reading at the MPI Forum meetings will be substituted by a review report through a review group. Each Forum member can be part of this group.

With the 1st official vote on a combined document (at next meeting), this modification of the voting rules is accept for that document.

MPI 2.1 Review Procedure (Voting plan) – Step 1

action
point

- Straw **vote** on the working plan (see 4 steps on previous slide)
- MPI 1.1 + Chap. 3 of MPI-2 (Version 1.2 of MPI) + some errata will be combined to
→ **MPI 1.3**

action
point

- Jan.08 meeting:
Short discussion and **defining a review group** who is reviewing the **MPI 1.3 merging plan** (printed copies available) and the **MPI 1.3 combined document**
- See e-mail: From: Rainer Keller, Subject: Re: [mpi-21] Documents
Date: Mon, 7 Jan 2008 12:13:14 +0100
- Reporting by e-mail on mpi-21 reflector
- Corrections if necessary (until Jan. 31, 2008)
→ final version of **MPI 1.3 merging plan** and **MPI 1.3**
- Final report of the reviewers at March 2008 meeting (=substitutes the **reading**)
- **1st vote** by the MPI Forum at April 2008 meeting
- **2nd (final) vote** by the MPI Forum at June 2008 meeting

MPI 2.1 Review Procedure (Voting plan) – Step 2

- MPI 1.3 combined document + rest of MPI-2 (MPI 2.0) will be combined to
→ **MPI 2.1 draft**
 - **Discussion of the 11 major merging decisions** and finishing them with **straw votes** (**Jan.2008** meeting)
based on the distributed text (printed copies available)
 - **Defining a review group** (**Jan.2008** meeting)
 - First draft of combined document (Feb 22, 2008, to be done by Rolf Rabenseifner)
 - Reviewing process and report of the reviewers (until March 10-12, 2008 meeting)
 - Discussion and further corrections if necessary (March 2008 meeting)
 - All necessary straw votes should be done at end of March 2008 meeting.
 - April 1, 2008, the final document should be available for twice voting.
 - Final report of the reviewers at April 2008 meeting (=substitutes the **reading**)
 - **1st vote** by the MPI Forum at June 2008 meeting
 - **2nd (final) vote** by the MPI Forum at Sep. 2008 meeting

action
point


action
point



MPI 2.1 Review Procedure (Voting plan) – Step 3

- adopted MPI 2.1 Ballots 1&2 + new MPI 2.1 ballots 3&4 are combined to the
→ **MPI 2.1 adopted errata**
(with references still based on MPI 1.1 and MPI-2 documents)
 - Ballots 1&2 are done (Chapter 1, Errata for MPI-2, May 15, 2002)
<http://www.mpi-forum.org/docs/errata-20-2.pdf>
 - Ballot 3
<http://www.cs.uiuc.edu/homes/wgropp/projects/parallel/MPI/mpi-errata/index.html>
<http://www.cs.uiuc.edu/homes/wgropp/projects/parallel/MPI/mpi-errata/ballot3.html>
 - **Official Reading** and straw votes (Jan. 14, 2008)
Re-reading if modifications are necessary (Jan. 15 or 16, 2008)
 - **First official vote** (on all topics with positive straw votes) (March, 2008)
 - Ballot 4 at March 2008 meeting with official reading (1st day and next days).
Unresolved items are moved to MPI 2.2.
 - → **Final text of Ballot 3+4 is available at end of March meeting, i.e., it can be already included into MPI 2.1 combined document.**
 - April 2008: Ballot 3: 2nd vote; Ballot 4: 1st vote
 - June 2008: Ballot 4: 2nd vote (all second votes are based on text in the combined documents) → **MPI 2.1 adopted errata** finished

action
point



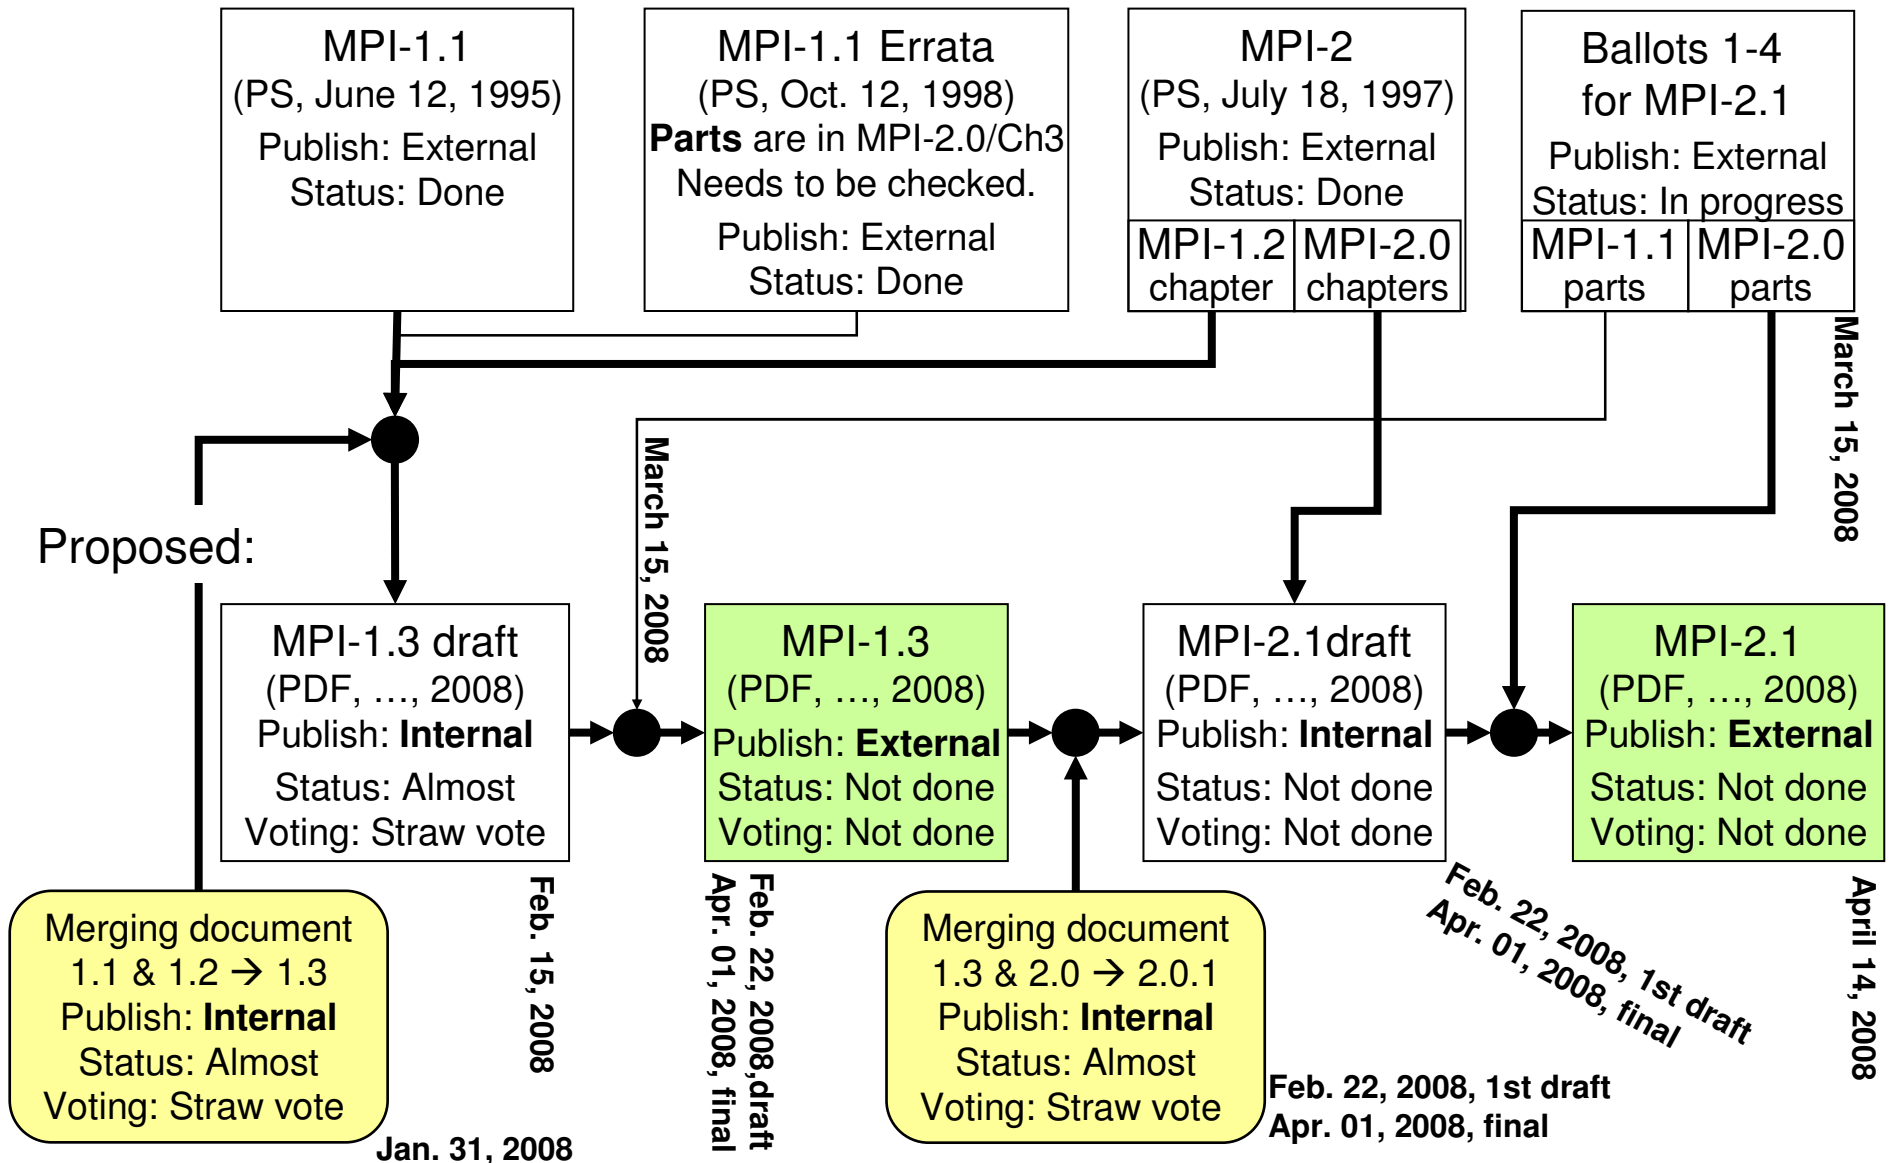
MPI 2.1 Review Procedure (Voting plan) – Step 4

- MPI 2.1 draft + MPI 2.1 adopted errata
→ **MPI 2.1 combined document**
 - The MPI 2.1 Ballots 1-4 (as after final reading on April 2008 meeting) are included into the MPI 2.1 draft (from April 1, 2008) (as prepared for final review/reading at April 2008 meeting)
→ MPI 2.1 combined document (April 14, 2008)
 - Defining the reviewing group (on March 2008 meeting) (may be smaller as for the MPI2.1 draft)
 - Reporting by e-mail on mpi-21 reflector until April 18, 2008
 - Corrections if necessary until April 23, 2008
 - Final report of the reviewers at April 2008 meeting (=substitutes the **reading**)
 - **First vote** on June 2008 meeting
 - **Second (final) vote** on Sep. 2008 meeting


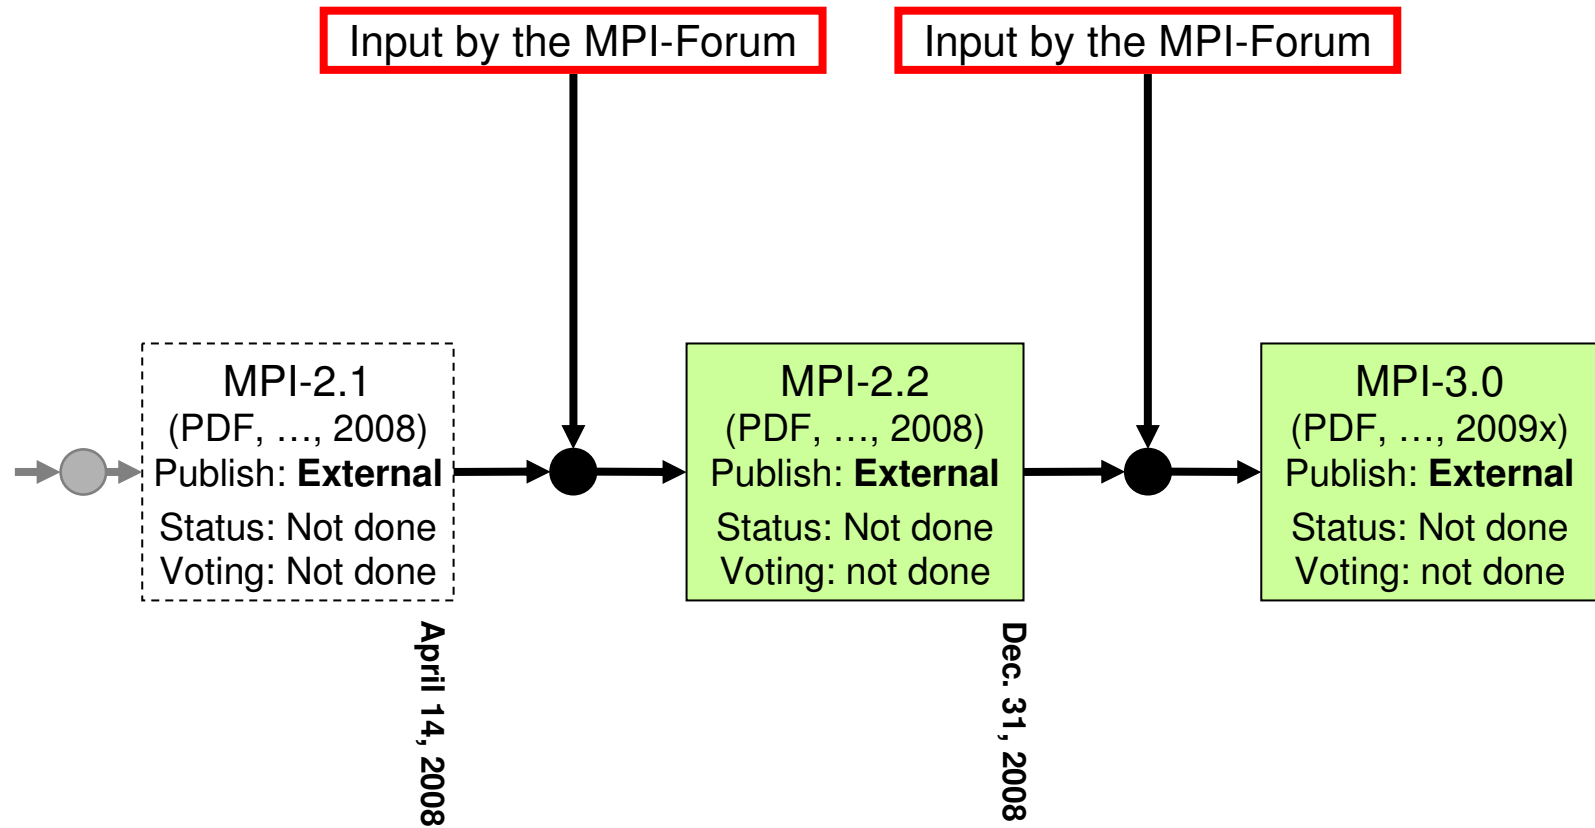


MPI Standards document plan

Existing and in progress:



MPI Standards document plan



Change-Log as Annex

- Who wants in MPI 2.1 an Annex telling effective (user and implementors visible) mods between MPI 2.0 and MPI 2.1?
- Yes: all=43
- No: 0
- Abstain: 0
- Should we keep old such histories in future standards?
- Yes: 8
- No: 13
- Abstain: $43-8-13=22$
- Changbars for all in the final combined document?
- Yes: 1
- No: 28
- Abstain: $43-28-1=14$



MPI 2.0 combined document

- There are 11 topics that must be decided or confirmed by the MPI Forum → next slides
- For further details we need a review group → slide at the end



MPI 2.0 merging decisions

Question:

The MPI 2.0 combined document title-page should be as stated here?

Yes:

All=41

No:

0

Abstain:

0

1.) The title of the combined document:

**MPI: A Message-Passing Interface Standard
Version 2.1**



Question:

MPI 2.0 merging decisions

The MPI
2.0
combined
document
title-page
should be
as stated
here in
2.+3.?

Yes:

all

No:

0

Abstain:


0

2.) The date of the merged document is fixed when it is released (in 2008).

3.) Ackno on the title page:

"This work was supported in part by ARPA, NSF and DARPA under grant ASC-9310330, the National Science Foundation Science and Technology Center Cooperative Agreement No. CCR-8809615, and the NSF contract CDA-9115428, and by the Commission of the European Community through Esprit project P6643 and under project HPC Standards (21111)."

4.) Do we add on 2.1 already new supporters? Yes – offline per e-mail



MPI 2.0 merging decisions

Question:

The MPI 2.0 combined document abstract should be as stated here?

Yes:

No:

Abstain:

4.) Abstract

"This document describes the MPI standard version 2.1 in one combined document. This document combines the content from the previous standards "*MPI: A Message-Passing Interface Standard, June 12, 1995*" (MPI-1.1) and "*MPI-2: Extensions to the Message-Passing Interface, July, 1997*" (MPI-1.2 and MPI-2.0). The standard MPI-1.1 includes point-to-point message passing, collective communications, group and communicator concepts, process topologies, environmental management, and a profiling interface. Language bindings for C and Fortran are defined. The MPI-1.2 part of the MPI-2 document contains clarifications and corrections to the MPI-1.1 standard and defines MPI-1.2. The MPI-2 part of the MPI-2 document describes additions to the MPI-1 standard and defines the MPI standard version 2.0. These include miscellaneous topics, process creation and management, one-sided communications, extended collective operations, external interfaces, I/O, and additional language bindings (C++). Additional clarifications and errata corrections are included."

→ offline e-mail : be specific on errata doc. And include MPI 1.3



MPI 2.0 merging decisions

5.) Copyright years

1993, 1994, 1995, 1996, 1997, 2008

Question:

The MPI 2.0 combined document copyright years should be as stated here?

Yes:


all

No:

0

Abstain:

0



MPI 2.0 merging decisions

Question:

The MPI 2.0 combined document Versions-list should be as stated here?

Yes:

No:

Abstain:

6.) New entries on the history page → **Offline per e-mail**

Still
TODO!

Version 2.1: <date>, 2008. This document combines the previous documents MPI 1.3 (????, 2008) and MPI-2.0 (July 18, 1997). Certain parts of MPI 2.0, such as some sections of Chapter 4, Miscellany, and Chapter 7, Extended Collective Operations have been merged into the Chapters of MPI 1.3. Additional errata and clarifications collected by the MPI Forum are also included in this document.

Version 1.3: <date>, 2008. This document combines the previous documents MPI 1.1 (June 12, 1995) and the MPI 1.2 Chapter in MPI-2 (July 18, 1997). Additional errata collected by the MPI Forum referring to MPI 1.1 and MPI 1.2 are also included in this document.

Version 2.0: <date>, 1997. Beginning after the release of MPI 1.1, the MPI Forum began meeting to consider corrections and extensions. MPI-2 has been focused on process creation and management, one-sided communications, extended collective communications, external interfaces and parallel I/O. A miscellany chapter discusses items that don't fit elsewhere, in particular language interoperability."

Version 1.2: July, 18 1997. The MPI-2 Forum introduced MPI 1.2 as Chap.3 in the standard "MPI-2: Extensions to the Message-Passing Interface", July 18, 2007." ...

Version 1.1: June, 1995. Beginning in March, 1995, the Message ...

Version 1.0: June, 1994. The Message Passing Interface Forum ...

As
already
voted for
MPI 1.2.1

MPI 2.0 merging decisions

Question:

The MPI
2.0
combined
document
Ackno-
pages
should be
as stated
here?

Yes:

all

No:

0

Abstain:

0

7.) The acknowledgment pages of both documents are printed one after the next.

The first list of persons is introduced with the following new line:

"Those who served as primary coordinators in MPI 1.0 and MPI 1.1 are:"

And the MPI-2 list of persons is introduced with:

"Those who served as primary coordinators in MPI 1.2 and MPI 2.0 are:"

For the merging process I would add:

"The editors of the combined documents have been:

-- Rainer Keller (MPI 1.3 - combined document)

-- Rolf Rabenseifner (MPI 2.1 - combined document)"



MPI 2.0 merging decisions

Question:

The MPI 2.0 combined document entries of deprecated routines should be as stated here?

Yes:

all

No:

0

Abstain:

0

8.) Before the definition of each deprecated interface, the following sentence is added:

**"The following function is deprecated
and is superseded by in MPI 2.0"**

MPI 2.0 merging decisions

Question:

The MPI 2.0 combined document structure should be as stated here?

Yes:

All ?

No:

0 ?

Abstain:


0 ?

9.) Sequence of all chapters:

- First all MPI-1, except Profiling
- All MPI-2, except Extended Collective (it is fully merged into corresponding MPI-1.1 chapters)
- Profiling Chapter
- Only one merged Appendix

The following MPI-2.0 sections are merged into the corresponding MPI-1.1 sections:

- MPI-2.0 Chap. 4 Miscellany, except
 - MPI-2.0 Sect. 4.10 (The Info Object) (Move? Yes:no:abstain)
 - MPI-2.0 Sect. 4.11 (Memory Allocation) → MPI-1 Env. Manag. (5:2:?)
 - MPI-2.0 Sect. 4.12 (Language Interoper.) → End of Lang.Bind 10.3 (17:0:?)
 - MPI-2.0 Sect. 4.17 (Functions and Macros) → Terms & Conven. 2.6.5 (5:0:?)
- MPI-2.0 Chap. 5 Extended Collective Operations
- MPI-2.0 Sect. 8.8 New Attribute Caching Functions
- MPI-2.0 Sect. 8.9 Duplicating a Datatype



Describe the chapter moving in a change-log

- Describe the chapter moving in a change-log?
 - Yes: all
 - No: 0
 - Abstain: 0



MPI 2.0 merging decisions

Question:

The MPI 2.0 combined document ...

should be as stated here?

Yes:

No:

Abstain:

Not decided by the Forum

10.) The new "Version Number" section was put at the beginning of the MPI-1.1 Chap. "Environmental Management" (done in MPI 1.2)

The Section "Portable MPI Process Startup" is put at the end of the MPI-1.1 Chapter "Environmental Management".

Okay?

Where to include
Section
„Memory Allocation“?

7 MPI Environmental Management

7.1 Implementation information

7.1.1 Version number

7.1.2 Environmental Interface

7.2 Error handling

7.3 Error codes and classes

7.4 Timers and synchronization

7.5 Startup

7.6 Portable MPI Process Startup

from
MPI 1.2

from
MPI 2.0

MPI 2.0 merging decisions

Question:

The MPI 2.0 combined document C++ interfaces should be handled as stated here?

Yes:

all

No:

0

Abstain:

0

11.) The MPI-1.2 C++ interfaces in MPI-2.0 App. B.3 - B.12 are written in a syntax that is different to the syntax used in all other C++ declarations.

For consistency, in B.3 - B.12, the

```
"namespace MPI {  
    standard-type mpi-routine(standard-type arg1, mpi-type arg2...)  
    mpi-type      mpi-routine(standard-type arg1, mpi-type arg2...)  
}"
```

→ this type of declaration should be used in Annex A

should be modified into

```
"standard-type MPI::mpi-routine(standard-type arg1, MPI::mpi-type arg2...)"  
"MPI::mpi-type MPI::mpi-routine(standard-type arg1, MPI::mpi-type arg2...),,"
```

Examples: `static MPI::Intercomm MPI::Comm::Join(const int fd)`
`MPI::Intercomm MPI::Intercomm::Create(const Group& group) const`
`MPI::Intercomm MPI::Intracomm::Create(const Group& group) const`
`void MPI::Comm::Barrier() const = 0`



Straw votes okay for merging details?

With MPI 1.3 and MPI 2.1 documents:

- Is it okay to have only a final „reading“ (=review report) and two official votes, instead of already doing official votes on some details?

Official (institutional) votes:

- Yes:
 - No:
 - Abstain:
-
- Reason: The merging does not modify the standard. Only formatting and editorial wording is rarely modified.
-
- **(This slide was skipped at January 2008 meeting.)**



MPI 2.0 Merging Review Group

We need reviewers for: (**bold=large**)

Reviewers: (**next slide**)

- **Frontmatter** (large)
- **Chapter 1: Introduction to MPI** (large)
- Chapter 2: MPI-2 Terms and Conventions (simple)
- **Chapter 3: Point-to-Point Communication** (large)
- **Chapter 4: Collective Communication** (large)
- **Chapter 5: Groups, Context, and Communicators** (large)
- Chapter 6: Process Topologies (simple)
- **Chapter 7: MPI Environmental Management** (large)
- Chapter 8: Miscellany (simple)
- Chapter 9: Process Creation and Management (simple)
- Chapter 10: One-Sided Communication (simple)
- Chapter 11: External Interfaces (simple)
- Chapter 12: I/O (simple)
- **Chapter 13: Language Bindings** (large)
- Chapter 14: Profiling Interface (simple)
- Bibliography (simple)
- **Annex A** (large)



We need reviewers for: (bold=large)	Reviewers: (green=@meeting)
Frontmatter	Rusty Lusk, Bill Gropp
Chap. 1: Introduction to MPI	Rusty Lusk, Bill Gropp, Karl Feind, Adam Moody
Chap. 2: MPI-2 Terms and Conventions	Tony Skjellum, Bill Gropp, Richard Barrett
Chap. 3: Point-to-Point Communication (incl. sections from MPI-2 Misc. + 8.9)	Rich Graham , Jespar Larsson Traeff, George Bosilca, Steve Poole, Kannan Narasimhan , David Solt, B. Gropp Matt Koop
Chap. 4: Collective Communication (incl. sections from MPI-2 Ext. Collect.)	Steven Ericsson-Zenith, Edgar Gabriel, Rajeev Thakur, Bill Gropp, Adam Moody, George Bosilca
Chap. 5: Groups, Context, and Communicators (incl. sections from MPI-2 Ext.Col. + 8.8)	Steven Ericsson-Zenith, Edgar Gabriel, Bill Gropp, George Bosilca, Robert Blackmore
Chap. 6: Process Topologies	Rusty Lusk, Bill Gropp, Richard Barrett
Chap. 7: MPI Environmental Management (incl. sections from MPI-2 Misc.)	Rich Graham , Jespar Larsson Traeff, George Bosilca, Steve Poole, Kannan Narasimhan , David Solt, B. Gropp
Chap. 8: Miscellany	Rich Graham, George Bosilca, Steve Poole, Kannan Narasimhan, B. Gropp
Chap. 9: Process Creation and Management	Dries Kimpe, Rusty Lusk, George Bosilca, Bill Gropp, Kalem Karian
Chap. 10: One-Sided Communication	Ericsson-Zenith, Jespar Larsson Traeff, Martin Schulz, Bill Gropp, Darius Buntinas
Chap. 11: External Interfaces	Bronis de Supinski, Bill Gropp
Chap. 12: I/O	Rajeev Thakur , Joachim Worringer, Bill Gropp
Chap. 13: Language Bindings	Jeff Squyres, Steve Poole, Purushotham Bangalore, Bill Gropp, Erez Haba, Alexander Supalov
Chap. 14: Profiling Interface	Bronis de Supinski, Bill Gropp, Jeff Brown
Bibliography	Rusty Lusk, Bill Gropp
Annex A	Jeff Squyres, Steve Poole, Purushotham Bangalore, Bill Gropp, Alexander Supalov

MPI 2.1 - Telecon

- I'll not be available on Feb. 11-13 (scheduled dates)
- There is no need because all Ballot 4 work will be done by e-mail
- And MPI 1.3, MPI 2.1 merging and Ballot 3 is done.
- → No telecon for the MPI 2.1 chapter
- Okay?
- (Accepted by the Forum due to the reasons mentioned above)



MPI 2.1

- Thank you very much for your strong support on MPI 2.1!