



MPI Collectives and Topology Workgroup

Atlanta, GA, January 2010



Torsten Hoefler

Agenda

- ▶ Topological Collectives
 - ▶ A concrete interface proposal
 - ▶ Problems
 - ▶ Solutions and Discussion
- ▶ Graph Topology issues

Motivation

- ▶ Stencil is probably the most important scalable communication pattern used!
- ▶ Different stencil operations
 - ▶ Regular (e.g., Cartesian) – each process is identical
 - ▶ e.g., simple CFD
 - ▶ Irregular – each process can be different
 - ▶ e.g., sparse matrix vector, AMR
- ▶ MPI supports optimization for both operations through graph topologies
 - ▶ Only allows to optimize mapping (which is important!)
 - ▶ Comm. scheduling can improve performance further
 - ▶ e.g., DCMF's (BG/P) Multisend interface (IPDPS'10); IPDPS'09



Overview and Terminology

- ▶ Establish directed communication graph between nodes
 - ▶ Each node has set of incoming and outgoing neighbors
 - ▶ IN = set of incoming neighbors (receive from)
 - ▶ OUT = set of outgoing neighbors (send to)
 - ▶ Topology can be optimized statically (offline)
 - ▶ Reorder ranks for network
 - ▶ Establish/optimize routing tables for minimum congestion
 - ▶ Optimize communication schedules for neighborhood collectives!
 - ▶ Similar to persistent requests
 - ▶ Buffers are not bound statically though (more flexible, network-centric)



Neighbor Gather

- ▶ `MPI_Neighbor_gather(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, comm)`
 - ▶ Gather data from all incoming neighbors to local buffer
 - ▶ single item in sendbuf
 - ▶ $|IN|$ items in recvbuf
 - ▶ Broadcast single item to all OUT processes
 - ▶ Receive item from $IN[i]$ at position i
- ▶ Both can be optimized as a tree if $|IN|$ is large enough
- ▶ Intelligent scheduling



Neighbor Alltoall

- ▶ `MPI_Neighbor_alltoall(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, comm)`
 - ▶ Exchange of all data
 - ▶ |OUT| items in sendbuf
 - ▶ |IN| items in recvbuf
 - ▶ Send item i in sendbuf to $OUT[i]$
 - ▶ Receive item I in recvbuf from $IN[i]$
- ▶ Dissemination- or Bruck-like algorithms
- ▶ Communication scheduling



Neighbor Reduce

- ▶ `MPI_Neighbor_reduce(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvttype, op, comm)`
 - ▶ Reduce items from neighbors
 - ▶ Single item in sendbuf
 - ▶ Single item in recvbuf
 - ▶ Apply op to all items from IN
 - ▶ Send item to OUT
- ▶ Enables tree-based implementations
- ▶ Intelligent scheduling



Vector Versions

- ▶ `MPI_Neighbor_gatherv()`
- ▶ `MPI_Neighbor_alltoallv()`
- ▶ `MPI_Neighbor_alltoallw()`
- ▶ `MPI_Neighbor_reducev()`
 - ▶ Special case!
 - ▶ Different connected neighborhoods have different numbers of elements
 - ▶ User needs to supply the correct values (non-trivial)!



Discussion

- ▶ Similar to persistent point-to-point
 - ▶ Persistent communication pattern
 - ▶ Build your own collective 😊
 - ▶ Non-persistent data
 - ▶ Does not conflict with persistent point-to-point
 - ▶ Should be discussed by persistence WG (is it active?)
- ▶ Nonblocking variants are easy to add
 - ▶ `MPI_Ineighbor_*(..., request)`

Example 1: Regular Stencil (MPI-2.2 p285)

```
int ndims=2, dims[2];
MPI_Dims_create(&comm, &ndims, &dims);
MPI_Comm comm_cart;
Int periods[2] = {1,1};
MPI_Cart_create(&comm, &ndims, &dims, &periods, 1, &comm_cart);
double sbuf[4], rbuf[4]; // exchange single double with each neighbor
for(int i=0; i<100; ++i) {
    prepare_buffers(u, sbuf, rbuf);
    MPI_Request req;
    MPI_Ineighbor_alltoall(sbuf, 1, MPI_DOUBLE, rbuf, 1, MPI_DOUBLE,
                           comm_cart, &req)
    relax(u);
    MPI_Wait(&req, MPI_STATUS_IGNORE);
}
```



Example 2: Irregular Stencil (MPI-2.2 p272)

```
// create graph topology (See example 7.3)
int in, out, weighted;
MPI_Dist_graph_neighbors_count(gcomm, &in, &out, &weighted);
int ine[in], oute[out], inw[in], outw[out];
MPI_Dist_graph_neighbors(gcomm, in, ine, inw, out, oute, outw);
double sbuf[out], rbuf[in]; MPI_Request req;
for(int i=0; i<100; ++i) {
    prepare_buffers(u, in, ine, out, oute, sbuf, rbuf);
    MPI_Ineighbor_alltoall(sbuf, 1, MPI_DOUBLE, rbuf, 1,
        MPI_DOUBLE, gcomm, &req)
    relax(u);
    MPI_Wait(&req, MPI_STATUS_IGNORE);
}
```

Biggest Problem!

- ▶ Most current MPI programs process incoming data as soon as it arrives
 - ▶ Synchronizing this with a collective adds additional and unnecessary overhead!
 - ▶ Will not be adopted by power-users if this problem persists!
- ▶ Do we see more problems?

Solutions!

- ▶ Streaming completion Version 1:
 - ▶ Also applicable to dense collectives 😊
 - ▶ Gather, Allgather, Alltoall
 - ▶ MPI_Cwaitany(count, request, rank)
- ▶ Streaming completion Version 2:
 - ▶ MPI_Sneighbor_alltoall (... , requests)
 - ▶ |IN| requests
 - ▶ Normal MPI requests
- ▶ Streaming completion also for send neighbors?
 - ▶ If not, we need an extra request for them



Example 3: Streaming Irregular Stencil

```
// same preamble as in example 2
MPI_Request reqs[in+1]; // one for all outgoing edges
for(int i=0; i<100; ++i) {
    prepare_buffers(u, in, ine, out, oute, sbuf, rbuf);
    MPI_Sneighbor_alltoall(sbuf, 1, MPI_DOUBLE, rbuf, 1,
        MPI_DOUBLE, gcomm, &reqs)
    relax(u);
    for(int j=0; j<in+1; j++) {
        int idx;
        MPI_Waitany(in+1, reqs, &idx, MPI_STATUS_IGNORE);
        if(idx!=in) merge(u, idx)
    }
}
```



Concrete Proposal

- ▶ Option 1: Addition to Collectives Chapter 5
 - 1. Suboptimal because Topologies are Chapter 7!
- ▶ Option 2: Addition to Topology Chapter 7
 - 1. Defining text in Section 7.6
 - 2. Examples in Section 7.7
- ▶ Option 3: Addition to Topology Chapter 7
 - 1. Defining text in Section 7.7
 - 2. Examples in Section 7.7
- ▶ Two examples:
 - 1. Regular stencil
 - 2. Irregular stencil/Arbitrary graph



MPI_Cart_create (#194/#195)

- ▶ See tickets (D'oh)!