

Collective Operations and Topologies WG

May 2010 Meeting

Torsten Hoefler



UNIVERSITY OF **ILLINOIS**
AT URBANA-CHAMPAIGN

Update on NBC

- Received three thorough reviews
 - Thanks to Fab, Alexander, and Dave
- Will “clean” the document and get it ready to be released asap
 - Will use pdf diff tools to ensure consistency
 - Do we have editing guidelines for MPI-3 yet?



NBC: Implementation Issues

Rich mentioned some problems?

Discussion



UNIVERSITY OF ILLINOIS AT URBANA-CHAMPAIGN

illinois.edu

(Non)Scalable Graph Topologies

- We added the distributed graph interface to MPI-2.2
 - Effectively replaces the old one
 - Is better in **all cases!**
- We postponed the deprecation
 - We should do it now
 - What is the Forum's opinion?



Sparse Collectives

- Worked with IBM (Sameer Kumar) on optimized implementation in DCMF
 - Trying to separate communication pattern (sparse colls) from buffer binding (persistent colls) and one-sided semantics
 - DCMF does all in one, synchronization based on collectives between iterations



DCMF Talk

Present IBM's DCMF talk here!



What fits MPI best

- Three separate principles:
 1. Sparse collectives (SC)
 - Enables neighborhood definitions
 2. Persistent collectives (PC)
 - Enables static buffer binding
 3. One Sided collectives (OSC)
 - Enables looser synchronization
 - Didn't really talk about this yet and it might migrate to the (too busy?) OS working group



Step by step

- Dependency chain (for neighborhoods)
 - $SC \leftarrow PC \leftarrow OSC$
 - (PC could stand on it's own)
- Step-wise process:
 - Propose SC
 - Propose PC
 - Propose OSC
- Forum decides where to stop ☺



SC: Progress

- Reference implementation exists in LibNBC
 - Well understood, non-blocking, proved useful in one application (TDDFT/Octopus)
- Proposal-draft finished
 - Soliciting initial feedback from Forum
 - Will be up for public comments soon



SC: Proposal

- Add as extension to Chapter 7
 - 7.6 Sparse Collective Communication on Process Topologies
- Why not in Chap 5 (Collectives)?
 - SC depends on Process topologies (PTs)
 - PTs are not introduced in Chap 5
 - Will back-reference to Chap 5



SC: MPI_Neighbor_gather

- `int MPI_Neighbor_gather(void* sendbuf, int sendcount, MPI_Datatype sendtype, void* recvbuf, int recvcount, MPI_Datatype recvttype, MPI_Comm comm)`
- Sends the **same** data to each neighbor
 - Vector variant for receiving different sizes



SC: MPI_Neighbor_alltoall

- `int MPI_Neighbor_alltoall(void* sendbuf, int sendcount, MPI_Datatype sendtype, void* recvbuf, int recvcount, MPI_Datatype recvttype, MPI_Comm comm)`
- Sends **personalized** data to each neighbor
 - Vector variant for different size comms
 - W-variant for optimized DDT layouts



SC: MPI_Neighbor_reduce

- `int MPI_Neighbor_reduce(void* sendbuf, int sendcount, MPI_Datatype sendtype, void* recvbuf, int recvcount, MPI_Datatype recvttype, MPI_Op op, MPI_Comm comm)`
- Vector variant for overlapping neighborhoods
- It's not directly applicable to stencils
 - Clear use-case needs to be found
 - We might decide to exclude it for now (?)



PC: Persistent Collectives

- They seem very useful – Tony?
 - Exploit temporal locality
 - Most (iterative) applications re-use buffers often
 - Really easy to specify
 - `MPI_Alltoall_init(..., &req)`
`MPI_Start(&req)`
`/* computation */`
`MPI_Wait(&req, MPI_STATUS_IGNORE)`
 - Trivial to implement w/o optimization



OC: One Sided Collectives

- Relax synchronization (cf. NBC)
 - Synch. Cannot be avoided!
- Buffers have to be ready all the time
 - Sender can just RDMA into recv-buffers
- Synchronization is tricky
 - When are the buffers valid? When can they be overwritten?
- DCMF implementation is very specific



OC: Discussion

- Want to solicit feedback from Forum if we should pursue this path
 - Discuss w/ OS WG



Questions/Discussion?



UNIVERSITY OF ILLINOIS AT URBANA-CHAMPAIGN
illinois.edu