

# MPI-3 RMA Proposal 1

Torsten Hoefler and William Gropp

On behalf of the RMA working group



UNIVERSITY OF **ILLINOIS**  
AT URBANA-CHAMPAIGN

# New Features at a Glance

- MPI\_Win\_allocate
- MPI\_Win\_create\_allmem
  - MPI\_Win\_register
  - MPI\_Win\_deregister
- MPI\_Get\_accumulate
  - MPI\_Accumulate\_get
- MPI\_Compare\_and\_swap





# New Features Continued

- MPI\_Win\_lock\_all
  - MPI\_Win\_unlock\_all
- MPI\_Win\_flush
  - MPI\_Win\_flush\_all
- MPI\_Win\_flush\_local
  - MPI\_Win\_flush\_local\_all
- MPI\_RMA\_query
- erroneous->undefined



# Win Allocate

- Collective Window Allocation
  - enables symmetric allocation for simple offset calculation
- MPI\_WIN\_ALLOCATE(size, disp\_unit, info, comm, base, win)
  - int MPI\_Win\_allocate(MPI\_Aint size, int disp\_unit, MPI\_Info info, MPI\_Comm comm, void \*base, MPI\_Win \*win)



# Win Create Allmem

- Creates a window that includes all the process' memory
  - Memory needs to be registered before being accessed remotely
- `int MPI_Win_create_allmem(MPI_Info info, MPI_Comm comm, MPI_Win *win)`
  - `int MPI_Win_register(MPI_Win win, void *base, MPI_Aint size)`
  - `int MPI_Win_deregister(MPI_Win win, void *base)`





# Get Accumulate

- Similar to fetch&add
- `int MPI_Get_accumulate(void *origin_addr, void *result_addr, MPI_Datatype datatype, int target_rank, MPI_Aint target_disp, MPI_Op op, MPI_Win win)`
  - Supports `MPI_NO_OP`



# Accumulate Get

- Reverse of Get Accumulate
  - Could easily be emulated by applying the op locally
  - Straw-vote for/against inclusion
- `int MPI_Accumulate_get(void *origin_addr, void *result_addr, MPI_Datatype datatype, int target_rank, MPI_Aint target_disp, MPI_Op op, MPI_Win win)`



# Compare and Swap

- `int MPI_Compare_and_swap(void *origin_addr, void *compare_addr, void *result_addr, int target_rank, MPI_Aint target_disp, MPI_Win win)`
- No datatype and op argument
  - Have special datatype and only bitwise equal as op ☹
  - Limited by today's hardware implementations
- Needs query function for size of the type
  - Maybe create and free too?





# Win Lock All

- Locks all target processes in a window
  - Optimization for a simple loop
- `int MPI_Win_lock_all(int assert, MPI_Win win)`
  - Needs `int MPI_Win_unlock_all(MPI_Win win)`
- Did not include `MPI_Win_lock_group`
  - No use-case



# Win Flush

- Blocks until all operations completed remotely
- `int MPI_Win_flush(int rank, MPI_Win win)`
  - Flush a specific rank
- `int MPI_Win_flush_all(MPI_Win win)`
  - Flush all ranks



# Win Flush Local

- Blocks until all operations completed locally
  - Local buffers can be re-used
- `int MPI_Win_flush_local(int rank, MPI_Win win)`
  - Operations targeted to a specific rank
- `int MPI_Win_flush_local_all(MPI_Win win)`
  - All previous operations on window win





# RMA Query

- Query remote memory consistency
  - Separate for each operation
- `int MPI_RMA_query(int optype, MPI_Win win, int *model)`
  - Returns either `MPI_RMA_SEPARATE` (MPI-2) or `MPI_RMA_ONE` (public window = private window)
  - `MPI_RMA_ONE` changes semantic rules 5+6



# RMA Query Continued

Process A:

MPI\_Barrier

MPI\_Win\_lock(EXCLUSIVE,B)

MPI\_Get(X) /\* ok, read from window \*/

MPI\_Win\_unlock(B)

Process B:

window location X

MPI\_Win\_lock(EXCLUSIVE,B)

store X /\* local update to copy of B \*/

MPI\_Win\_unlock(B)

MPI\_Barrier



# RMA Query Continued

Process A:

MPI\_Barrier

**MPI\_Win\_lock(EXCLUSIVE,B)**

MPI\_Get(X) /\* ok, read from window \*/

**MPI\_Win\_unlock(B)**

Process B:

window location X

**MPI\_Win\_lock(EXCLUSIVE,B)**

store X /\* local update to copy of B \*/

**MPI\_Win\_unlock(B)**

**MPI\_Win\_flush(B)**

MPI\_Barrier





# RMA Query Continued

Process A:

**MPI\_Win\_lock(EXCLUSIVE,B)**

MPI\_Put(X) /\* update to window \*/

**MPI\_Win\_unlock(B)**

**MPI\_Win\_flush(B)**

MPI\_Barrier

Process B:

window location X

MPI\_Barrier

**MPI\_Win\_lock(EXCLUSIVE,B)**

load X

**MPI\_Win\_unlock(B)**



# Bonachea's and Duell's criticism

- Window creation is collective
  - hinders efficient exposure for local objects
  - no “sparse” communication
- MPI\_Win\_create\_allmem
  - Well suited for automatic memory management
  - Good as compilation target



# Bonachea's and Duell's criticism

- Exposed memory must be `MPI_Alloc_mem()`'d
  - no exposure of static memory or stack-variables
  - `alloc_mem` might be limited by the implementation
- Also addressed in `MPI_Win_create_allmem`
  - Can register stack
  - Dangerous though! Undefined results if stack is out of focus.





# Bonachea's and Duell's criticism

- Forbids conflicting get/put (or local load/store) accesses to same memory
  - really hard to track for compilers (halting problem?)
  - Easy source of bugs in user codes
- Addressed in the proposal 1
  - outcome is undefined



# Bonachea's and Duell's criticism

- Window's memory may not be updated by remote gets and local stores concurrently
  - simplifies MPI implementation significantly
  - seems very artificial and suboptimal from user's perspective
- Addressed in the proposal 1
  - outcome is undefined



# Bonachea's and Duell's criticism

- Overlapping memory regions of multiple windows can be created but not be used
  - “concurrent communications may lead to erroneous results”
- now also “undefined”
  - [is this much better than erroneous?]
  - But these applications are likely to use a single “allmem” window, so problem avoided





# Bonachea's and Duell's criticism

- Passive target RMA ops only lock a single process during an epoch
  - ops from one source to different targets are serialized
  - one window for each target to enable concurrent access?
    - scalability limitation
- MPI\_Win\_lockall



# Comparison to ARMCI

- Similar to MPI\_RMA\_ONE (if available)
  - ARMCI does not support non-CC systems without an additional activity (thread/process/callback)
- Local completion is different
  - Either blocking, implicit handles or Test/Wait(all)
  - Proposal 1 is similar to “implicit handles”
- Remote completion similar
  - (all)Fence == Flush(\_all)
- Ordering with collectives
  - ARMCI\_Barrier combines barrier + allfence



# Comparison to ARMCI cont.

- ARMCI\_Malloc for exposed memory
  - Collective and local; no register (in an undocumented version)
- ARMCI\_RMW
  - Similar to Get\_accumulate
  - No compare-and-swap
- ARMCI\_Lock/Unlock
  - Lock special synch. objects
  - Proposal 1 doesn't include user-defined locks
- Limited set of collectives
  - Supports MPI collectives





# Co-Array Fortran

- RMA through co-arrays
  - Explicit collective allocation
- Execution divided in segments (delimited by “image control statements”)
  - (non-volatile) operations in a segment are unordered
  - Exceptions are values with “atom” or volatile argument
- Image control statements:
  - Sync all, sync images, sync memory, lock/unlock
- Can be implemented with proposal 1
  - Active messages might be necessary



# Unified Parallel C

- Explicit collective allocation
- Relaxed and strict access
  - Relaxed accesses are unordered
    - Exception are conflicting accesses to the same memory which appear in program order
  - Strict accesses are ordered (program order)
- Synch operations:
  - `upc_fence`, `upc_(notify,wait)`, `upc_barrier`, `upc_(un)lock`
  - `flush`, `(nonblocking) barrier`, `barrier`, `<no user locks>`



# Comparison Summary

- Proposal 1 covers most features from existing RMA systems
  - No remote completion of individual operations
    - Can often be emulated with flush
  - No user-defined locks
    - Can be emulated with compare and swap
- Proposal 1 can efficiently simulate other RMA systems
  - AMs seem to be missing but they can be simulated with point-to-point





# Executive Summary

- Proposal 1 fixes known deficits of MPI-2 RMA and can be implemented on non-CC architectures
  - Part of the complexity is pushed to the user
    - Enables the use of MPI-RMA for new important application domains
      - E.g., Graph computation, well, AMs are missing ☺
- Proposal 2 has further extensions
  - Individual operation completions among others



# Thanks!

## Please review the proposal!

### Questions?



<https://svn.mpi-forum.org/trac/mpi-forum-web/wiki/RmaWikiPage>

