

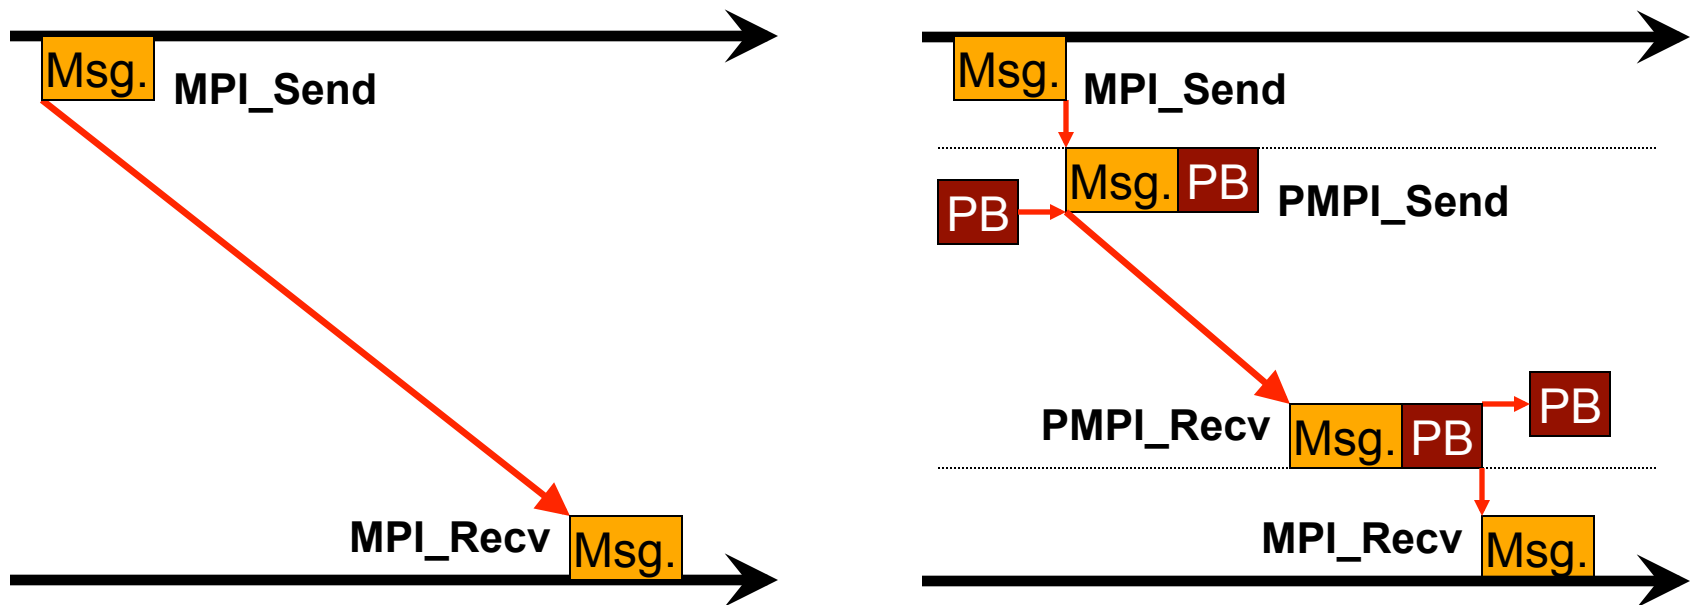
Piggybacking in MPI-3

MPI Forum Meeting, October 2010

Martin Schulz with input from
George Bosilica, Greg Bronevetsky, Darius Buntinas,
Bronis de Supinski, Kathryn Mohror, Anh Vo

What do we mean by piggybacking?

- ▶ Transparently attach data to messages
 - ▶ Uniquely associated with a specific message
 - ▶ Piggyback data added/removed without influencing application
 - ▶ Typically used for small amounts of data



Use Cases

- ▶ **Fault Tolerance Protocols**
 - ▶ Transparent propagation of checkpoint interval IDs
 - ▶ Message logging
 - ▶ **Performance Analysis**
 - ▶ Overhead propagation/detection
 - ▶ Critical path analysis
 - ▶ Late sender instrumentation
 - ▶ **Debugging**
 - ▶ MPI correctness tools
 - ▶ **General**
 - ▶ Transparent state propagation in libraries
 - ▶ Message matching
-

Requirements

- ▶ **Must work for the complete MPI standard**
 - ▶ All point to point operations, including wild card receives
 - ▶ Collective operations, including reductions and barriers
 - ▶ Need to look into all corner cases (e.g., lssend matching lrecv(*))
 - ▶ **Thread safety**
 - ▶ **Low overhead, minimal impact on performance**
 - ▶ Raw bandwidth and latency
 - ▶ Application performance
 - ▶ **“Selective piggyback”, ideally per message**
 - ▶ Ability to do piggybacking on a subset of messages
 - ▶ Receiver needs to know/query whether piggyback is present
 - ▶ Selective piggybacking infeasible for PMPI wrappers
-

Solution on top of MPI are insufficient (1)

- ▶ Three main options
 - ▶ A) Two messages
 - ▶ B) MPI datatypes (as they are now)
 - ▶ C) Manual packing into a separate buffer
 - ▶ A) Two messages
 - ▶ Fails for wildcard receives (can lead to incorrect matchings)
 - ▶ Transparent data return difficult (need to attach to requests/ requires memory allocations)
 - ▶ Thread safety unclear (how to handle two sends from two threads in the same process)
-

Solution on top of MPI are insufficient (2)

- ▶ **B) MPI datatypes (as they are now)**
 - ▶ Requires creation of new datatype for each message (expensive and hard to optimize)
 - ▶ Use of MPI_BOTTOM may reduce optimization potential
 - ▶ Large overhead on applications (see next slides)
 - ▶ Partial receives cause problems (how to correctly count number of elements and bytes)
 - ▶ Does not work for reduce/barrier (no datatypes used), difficult for other collectives
 - ▶ **C) Manual packing into a separate buffer**
 - ▶ High overhead (memory copy for each messages)
 - ▶ Partial receives hard (how to correctly count number of elements and bytes)
 - ▶ Does not work for reduce (different ops)/barrier (no data), difficult for other collectives
-

Solution on top of MPI are insufficient (3)

- ▶ In all cases: no selective piggybacking
 - ▶ Receiver does not know whether sender included piggyback
 - ▶ Sender does not know whether receiver expects piggyback
 - ▶ Only solution: always send piggyback data
 - ▶ Even if they would work, performance is problematic
 - ▶ Significant overheads in latency and bandwidth
 - ▶ Performance not portable across implementations
-

Measurement Setup

- ▶ **MPI implementations and platforms**
 - ▶ MVAPICH on LLNL's TLCC clusters
 - ▶ Open MPI on LLNL's TLCC clusters
 - ▶ Blue Gene/P (dawndev) at LLNL
 - ▶ [IBM AIX (up) and IA-64/Quadrix (thunder) at LLNL]
 - ▶ **Experiments**
 - ▶ Raw bandwidth and latency for piggybacking
 - ▶ Three naïve implementations mentioned before
 - ▶ Ignoring correctness problems
 - ▶ Sending with complex datatypes
 - ▶ Data type creation (incl. complex datatypes)
 - ▶ Application performance (Sweep3D and SMG2000)
-

Latency

- ▶ Hera 4 byte (MVAPICH 1)
 - ▶ Base version: 2.163us
 - ▶ Datatype: 3.306us (**52.8%**)
 - ▶ Two messages: 3.154us (**45.8%**)
 - ▶ Packing: 2.926us (**35.3%**)
- ▶ Hera 4 byte (Open MPI 1.4.2)
 - ▶ Base version: 1.804us
 - ▶ Datatype: 3.444us (**90.9%**)
 - ▶ Two messages: 3.623us (**100.8%**)
 - ▶ Packing: 3.411us (**89.1%**)
- ▶ BG/P 4 byte (MPICH-2)
 - ▶ Base version: 3.573us
 - ▶ Datatype: 53.959us (**1410.2%**)
 - ▶ Two messages: 8.099us (**126.7%**)
 - ▶ Packing: 14.509us (**306.1%**)

Bandwidth



- ▶ Hera 4 byte (MVAPICH 1)
 - ▶ Base version: 1238 MB/s
 - ▶ Datatypes: 579 MB/s (**-53.2%**)
 - ▶ Two messages: 1231 MB/s (**-0.6%**)
 - ▶ Packing: 572 MB/s (**-53.8%**)
- ▶ Hera 4 byte (Open MPI 1.4.2)
 - ▶ Base version: 1429 MB/s
 - ▶ Datatype: 1106 MB/s (**-22.6%**)
 - ▶ Two messages: 1104 MB/s (**-22.7%**)
 - ▶ Packing: 1105 MB/s (**-22.7%**)
- ▶ BG/P 4 byte (MPICH-2)
 - ▶ Base version: 355 MB/s
 - ▶ Datatypes: 239 MB/s (**-32.7%**)
 - ▶ Two messages: 351 MB/s (**-1.1%**)
 - ▶ Packing: 175 MB/s (**-50.7%**)

Sending with Wrapped Datatypes

- ▶ Latency for sending with struct datatypes
 - ▶ Pre-create one struct datatype and send repeatedly with it
 - ▶ Hera/MVAPICH: **17.4%** overhead (2.540us vs. 2.163us)
 - ▶ Hera/Open MPI: **8.8%** overhead (1.964us vs. 1.804us)
 - ▶ BG/P: **77.9%** overhead (6.357us vs. 3.573us)
 - ▶ Latency for sending with continuous datatype creation
 - ▶ Create one struct datatype for each send and receive operation
 - ▶ Hera/MVAPICH: **52.3%** overhead (3.295us vs. 2.163us)
 - ▶ Hera/Open MPI: **73.8%** overhead (3.136us vs. 1.804us)
 - ▶ BG/P: **1220.2%** overhead (47.171us vs. 3.573us)
-


Datatype Creation

- ▶ Continuous creation and freeing of datatypes
 - ▶ Single node execution
 - ▶ Increasing depth of structured datatypes (from “struct” to “struct of struct” to “struct of struct of struct”)




Application Performance

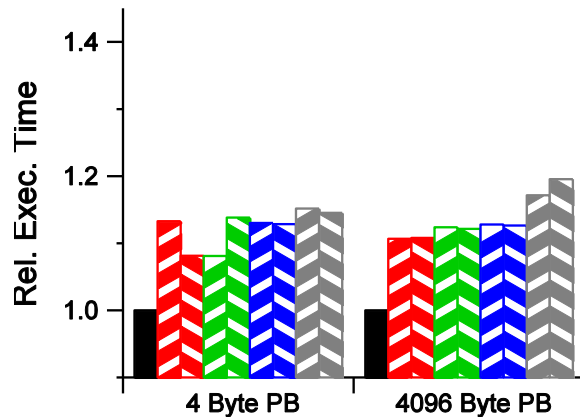
■ Baseline ▨ Pack/Before ▨ Datatype/Before ▨ SepMsg/Before ▨ SepMsg/Before (WC)
 ▨ Pack/After ▨ Datatype/After ▨ SepMsg/After ▨ SepMsg/After (WC)



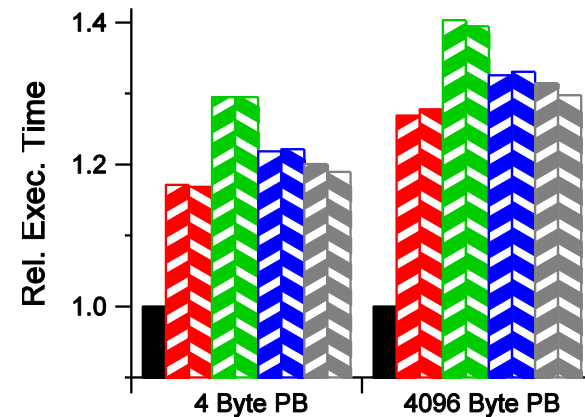
Sweep3D on Thunder



Sweep3D on uP



SMG2000 on Thunder



SMG2000 on uP

Pack = Use memory copy to add PB to message
 Datatype = Use new struct datatype for each message
 SepMsg = Use two messages
 SepMsg(WC) = SepMsg with using MPI_ANY_SOURCE in the code

We Need a New API for Piggybacking

- ▶ Requirement for many tools, but
 - ▶ Existing solutions don't work or are too expensive
 - ▶ Solutions using PMPI are infeasible
 - ▶ However, this is not a difficult requirement
 - ▶ Can easily be implemented within MPI library
 - ▶ Portable access to this functionality requires standardized API
 - ▶ Design guidelines
 - ▶ Clean semantics
 - ▶ Specific for the intended use, not too general
 - ▶ Allow optimizations in MPI for better piggybacking
-

Ideal Solution: Separate Set of API Calls

- ▶ Explicit PB buffers as arguments to send calls
 - ▶ **MPI_SendPB**(buf,count,datatype,dest,tag,comm,
pb_buf,pb_count,pb_datatype)
 - ▶ pb_count and pb_datatype could be restricted
 - ▶ Receiving PB data is done through the status object
 - ▶ Accessor routine maintains compatibility
 - ▶ **MPI_CheckPB**(status, flag, buf,count,datatype)
returns flag=1 if PB has been received
buf,count, datatype act like a receive operation
 - ▶ Separate arguments in collective calls
 - ▶ **MPI_BcastPB**(buf,count,datatype,root,comm,
pb_buf,pb_count,pb_datatype)
 - ▶ Full functionality by allowing many cross combinations (e.g., gather/reduce)
(reduction operation for PB depends on PB user and is independent of the host operation in the MPI code)
 - ▶ **MPI_BcastPBReduce**(buf,count,datatype,root,comm,
pb_buf,pb_count,pb_datatype,reduction_operator)
-

Pros and Cons

▶ Advantages

- ▶ Clean semantics of adding an orthogonal aspect
- ▶ Directly/Only implements the required functionality

▶ Disadvantages

- ▶ More API calls
- ▶ Cross product of all collectives and all PB collectives necessary
- ▶ Wrapped PB messages a bit more difficult

▶ Open questions

- ▶ Which collective operations should be supported for PB?
 - ▶ Updates to message matching rules (?)
-

Alternatives

- ▶ **Prototype datatypes**
 - ▶ Precreate “datatype prototypes”
 - ▶ Populate prototypes before sending/receiving
 - ▶ **Overloading parameters**
 - ▶ Reuse existing messaging APIs
 - ▶ “Attach” piggyback information to a parameter
 - ▶ Datatypes
 - ▶ Communicators
-

Datatype Prototypes for Point to Point

- ▶ **Precreate datatype**
 - ▶ Create a two element struct
 - ▶ Specify `pb_count`, `pb_datatype`, but not `pb_buf`
 - ▶ Don't specify user buffer
 - ▶ Before message send/receive
 - ▶ Add `pb_buf` information
 - ▶ Instantiate prototype with user buffer information
 - ▶ **Idea: expensive step only once**
 - ▶ MPI has chance to preallocate datatype
 - ▶ Per message operation simple and cheap
 - ▶ Maintain datatype optimization principle
-

Example Part 1: Setup (once)

- ▶ **MPI_Aint** offsets[2];
int counts[2];
MPI_Datatype types[2];
MPI_Datatype mine;
 - ▶ offsets[0] = **MPI_UNKNOWN**;
offsets[1] = **MPI_UNKNOWN**;
 - ▶ types[0] = pb_dtype;
types[1] = **MPI_TYPE_UNKNOWN**;
 - ▶ counts[0] = pb_count;
counts[1] = **MPI_CNT_UNKNOWN**;
 - ▶ MPI_Type_create_pstruct(2, counts, offsets, types, &mine);
 - ▶ MPI_Type_commit(&mine);
-

Example Part 2: Communication

- ▶ **MPI_Datatype mine_complete;**
 - ▶ offsets[0] = pb_buf;
offsets[1] = user_buf;
 - ▶ types[1] = user_dtype;
counts[1] = user_count;
 - ▶ **PMPI_Type_complete_pstruct**(2, counts, offsets, types,
mine, &mine_complete);
 - ▶ **PMPI_Send**(MPI_BOTTOM, 1, mine_complete, dest, tag,
comm);
 - ▶ **PMPI_Type_free**(&mine_complete);
-

Discussion

▶ Pros

- ▶ Small API extension
- ▶ Builds on top of familiar concepts

▶ Cons

- ▶ Still a complex datatype creation
- ▶ Piggyback too generic
- ▶ Requires memory allocation/free for each send/receive
- ▶ Difficult for collectives
- ▶ Does not apply to barriers

▶ Open questions

- ▶ Performance
 - ▶ Support for collectives
 - ▶ API to support grouping of collectives?
 - ▶ Support for selective piggybacking
-

Attach to Existing Parameters

- ▶ **We can't change existing APIs**
 - ▶ Assumption: we don't want to duplicate API
 - ▶ Need to use existing parameters to convey new information
 - ▶ Overload arguments
 - ▶ **Can be done with any parameter; attractive choices are**
 - ▶ Datatypes
 - ▶ Communicators
 - ▶ **Main idea**
 - ▶ Pre-create template to notify MPI that this may happen
 - ▶ Dynamically attach content to parameter before call
 - ▶ Automatic detach after call / use once
-

Attach PB to Datatypes

- ▶ Create a template ahead of time
 - ▶ `MPI_Datatype_PB_Template(MPI_Datatype pb_dt, int pb_count, MPI_Datatype_template *pb_template)`
 - ▶ Optionally provide collective operation (default none)
 - ▶ `MPI_Datatype_PB_ReduceOp(MPI_Op pb_op, MPI_Datatype_template *pb_template)`
 - ▶ Create new sending datatype
 - ▶ `MPI_Datatype_PB_Instantiate(void *buf, MPI_Datatype_template pb_template, MPI_Datatype dt, MPI_Datatype *dt_new)`
 - ▶ Communication with new datatype
 - ▶ `MPI_Send(<use new datatype>)`
 - ▶ `MPI_Recv(<use new datatype>)`
 - ▶ Check whether we received a piggyback operation
 - ▶ `MPI_Datatype_PB_Filled(MPI_Datatype dt_new, int flag)`
 - ▶ Note: no free operation required
-

Attach PB to Communicator

- ▶ Create a template ahead of time
 - ▶ `MPI_Comm_PB_Template(MPI_Comm comm, int count, MPI_Comm_template *pb_template)`
 - ▶ Optionally provide collective operation (default none)
 - ▶ `MPI_Comm_PB_ReduceOp(MPI_Ops pb_ops, MPI_Comm_template *pb_template)`
 - ▶ Create new sending datatype
 - ▶ `MPI_Comm_PB_Instantiate(void *buf, MPI_Comm_template pb_template, MPI_Comm comm, MPI_Comm *comm_new)`
 - ▶ Communication with new communicator
 - ▶ `MPI_Send(<use new communicator>)`
 - ▶ `MPI_Recv(<use new communicator>)`
 - ▶ Check whether we received a piggyback operation
 - ▶ `MPI_Comm_PB_Filled(MPI_Comm comm_new, int flag)`
 - ▶ Note: no free operation required
-

Example for Usage

```
/* One time setup */
MPI_Comm_template tcw;
MPI_Comm_PB_Template(MPI_COMM_WORLD,I,&tcw);

...
/* Send wrapper */
MPI_Send(buf,c,dt,target,tag,comm)
    int pbbuf;
    MPI_Comm newcomm;
    MPI_Comm_PB_instantiate(&pbbuf,tcw,comm,&newcomm);
    PMPI_Send(buf,c,dt,target,tag,newcomm);

...
/* Receive wrapper */
MPI_Recv(buf,c,dt,target,tag,comm,status)
    int pbbuf,flag;
    MPI_Comm newcomm;
    MPI_Comm_PB_instantiate(&pbbuf,tcw,comm,&newcomm);
    PMPI_Recv(buf,c,dt,target,tag,newcomm,status);
    MPI_Comm_PB_filled(newcomm,&flag);
    if (flag) { <use data in pbbuf> }
```

Conclusions

- ▶ **We need a piggybacking API in MPI**
 - ▶ Many use cases for a wide range of tools
 - ▶ Existing functionality insufficient
 - ▶ **Semantically clean solution**
 - ▶ Duplicate messaging API
 - ▶ Add new parameters to specify piggybacking
 - ▶ **Alternative with small API footprint**
 - ▶ Datatype prototypes
 - ▶ Overloading datatypes
 - ▶ Overloading communicators
-