

MPI 3.0 Fortran Bindings:

a high-level summary for non-Fortran
programmers

Craig Rasmussen (Los Alamos)
Rolf Rabenseifner (HLRS)
Jeff Squyres (Cisco Systems)

A very brief overview of the requirements for new MPI 3.0 Fortran bindings

- Requirements

- **comply** with Fortran standard (for the first time)
- enhance **type safety**
- suppress argument checking for choice buffers
- guarantee of **correct** asynchronous operations
- for user convenience
 - provide users with convenient migration path
 - allow some optional arguments (e.g., ierror)
 - support sub-arrays
- for vendor convenience
 - allow vendors to take advantage of the C interoperability standard

Requirement 1: comply with fortran standard

- Requirement
 - **comply** with Fortran standard (for the first time)
- Example of non-compliance
 - call MPI_Send(real_buffer_argument, ...)
 - call MPI_Send(integer_buffer_argument, ...)
- Do you agree with this requirement
 - ☐ yes
 - ☐ no (why not)

Requirement 2: enhance type safety

- Requirement
 - enhance **type safety** (for user)
- Example
 - call MPI_Send(buf,count,datatype,dest,**comm**,**tag**,ierror)
 - call MPI_Send(buf,count,datatype,dest,**tag**,**comm**,ierror)
 - call MPI_Send(buf,count,datatype,dest,**comm**,ierror)
- Do you agree with this requirement
 - ☐ yes
 - ☐ no (why not)

Requirement 3: solve mapping to void* buffers

- Requirement

- have standardized possibility to suppress argument checking for choice buffers (for implementors)

- Example

- currently no good standards-based solution with explicit interfaces
 - Fortran interface is equivalent to a C prototype
 - must define interface for **all combinations** of (type X, kind X dim)
(if the compiler does not support a non-standardized IGNORE option)

- Do you agree with this requirement

☐

yes

☐

no (why not)

Requirement 4: provide ability to write correct programs using asynchronous operations

- Requirement
 - guarantee of **correct** asynchronous operations
- Example of incorrect behavior
 - call MPI_Irecv(buf, ...)
 - compiler may be required to provide temporary copy of buf
 - **problem: temporary copy deallocated on return from MPI_Irecv**
- Do you agree with this requirement
 - ☐ yes
 - ☐ no (why not)

Requirement 4: provide ability to write correct programs using asynchronous operations (2)

- Requirement
 - guarantee of **correct** asynchronous operations
- Example of incorrect behavior
 - call MPI_Irecv(buf, ..., rq ,...)
 - ...
 - call MPI_Wait(rq)
 - a = buf(1)
 - **Problem: code movement of reference to buf moved above MPI_Wait()**

Requirement 5: provide a convenient migration path for users

- Requirement
 - provide users with convenient migration path
- Examples
 - allow “include mpif.h” and “use mpi_f08” in the same program
 - allow easy conversion between new and old Fortran handles
 - make usage as similar as possible
 - subroutines not functions
- Do you agree with this requirement
 - ☐ yes
 - ☐ no (why not)

Requirement 6: take advantage of some features in Fortran as a convenience to users

- Requirement
 - allow some optional arguments (e.g., ierror)
- Example of common error pattern
 - call MPI_Send(buf,count,datatype,dest,tag,comm)
 - call MPI_Send(buf,count,datatype,dest,tag,comm,ierror)
 - forgetting the ierror return argument will no longer be an error
- Do you agree with this requirement
 - ☐ yes
 - ☐ no (why not)

Requirement 7: take advantage of sub-arrays for convenience and performance

- Requirement

- support sub-arrays

- Example

- call `MPI_Irecv(Array(1,:), ...)`

- instead:

set up all the arguments for creating subarray dt
call `MPI_Type_create_subarray(..., dt, ierror)`
call `MPI_Irecv(Array(:, :), 1, dt, ...)`

- Do you agree with this requirement

☐ yes

☐ no (why not)

Requirement 8: allow vendors access to existing C implementations of Fortran wrappers

- Requirement

- allow vendors access to existing C implementations of Fortran wrappers

- Example

- Existing mpif.h and mpi module routine is, e.g., `mpi_send__`
- For the new `mpi_f08` module with existing Fortran 2003 compilers one should only do
 - ▶ substitute `*ierror = ...` by `if (ierror /= NULL) *ierror = ...`
 - ▶ add an alias `mpi_send_f08__` to `mpi_send__`

- Do you agree with this requirement

☐

yes

☐

no (why not)

A very brief overview of design tradeoffs

- Now we take a look at some of the design tradeoffs made in implementing the requirements

Requirement: comply with fortran standard

- Requirement
 - **comply** with Fortran standard (for the first time)
- Example of non-compliance
 - call MPI_Send(real_buffer_argument, ...)
 - call MPI_Send(integer_buffer_argument, ...)
- Design decision
 - take advantage of new features in Fortran
 - TYPE(*), DIMENSION(..)
- Tradeoffs
 - wait for compiler support for full implementation of nonblocking routines

Requirement: enhance type safety

- Requirement
 - enhance **type safety**
- Example
 - call MPI_Send(buf,count,datatype,dest,**comm,tag**,ierror)
 - call MPI_Send(buf,count,datatype,dest,**tag,comm**,ierror)
- Design decision
 - MPI handles are user-defined types
 - TYPE(MPI_Comm)
- Tradeoffs
 - users must convert code to use new handle types

Requirement: provide ability to write correct programs using asynchronous operations

- Requirement
 - guarantee of **correct** asynchronous operations
- Example of incorrect behavior
 - call `MPI_Irecv(buf, ...)`
 - compiler may be required to provide temporary copy of buf
 - **problem: temporary copy deallocated on return from `MPI_Irecv`**
- Design decision
 - use `DIMENSION(..)` and `ASYNCHRONOUS` attribute for choice buffers
- Tradeoffs
 - vendors must convert implementation to access array descriptor

Requirement: provide a convenient migration path for users

- Requirement
 - provide users with convenient migration path
- Examples
 - allow easy conversion between new and old Fortran handles
 - make usage as similar as possible
- Design decision
 - handles are user defined types (think C structs) with an integer component
 - subroutines not functions
- Tradeoffs
 - wrapper functions still needed because of choice to use subroutines

Requirement: take advantage of some features in Fortran as a convenience to users

- Requirement
 - allow some optional arguments (e.g., ierror)
- Example of common error pattern
 - call `MPI_Send(buf,count,datatype,dest,tag,comm)`
 - call `MPI_Send(buf,count,datatype,dest,tag,comm,ierror)`
 - forgetting the ierror return argument will no longer be an error
- Design decision
 - ierror argument made optional
- Tradeoffs
 - None

Requirement: take advantage of sub-arrays for convenience and performance

- Requirement
 - support sub-arrays
- Example
 - call `MPI_Irecv(Array(1,:), ...)`
 - receive into top row of array
- Design decision
 - choice buffers declared with `DIMENSION(..)`
- Tradeoffs
 - vendors must convert implementation to access array descriptor

Decision for C programmers: to support or not to support MPI 3.0 Fortran bindings

- Do you agree with all or most of the requirements?
- Do the technical tradeoffs seem reasonable?
- Since you may not understand the technical details, do you trust the Fortran experts to make good decisions regarding the low-level technical details?
- Fortran Experts
 - WG5 Fortran committee agrees with general changes but don't follow details
 - Nick, Steve Lionel
 - Rolf (and contacts...)
 - Craig and Jeff