

The MPI-Forum Organization and Call for Participation

Martin Schulz

LLNL / CASC

Chair of the MPI Forum

MPI Forum BOF @ SC14



<http://www.mpi-forum.org/>

This work was performed under the auspices of the U.S. Department of Energy by Lawrence Livermore National Laboratory under Contract DE-AC52-07NA27344.



The MPI Forum Drives MPI

- **Standardization body for MPI**
 - Discusses additions and new directions
 - Oversees the correctness and quality of the standard
 - Represents MPI to the community
- **Organization consists of chair, secretary, convener, steering committee, and member organizations**
- **Open membership**
 - Any organization is welcome to participate
 - Consists of working groups and the actual MPI forum
 - Physical meetings 4 times each year (3 in the US, one with EuroMPI/Asia)
 - Working groups meet between forum meetings (via phone)
 - Plenary/full forum work is done mostly at the physical meetings
 - Voting rights depend on attendance
 - An organization has to be present two out of the last three meetings (incl. the current one) to be eligible to vote

Typical Way to Add Items to MPI

1. **New items should be brought to a matching working group for discussion**
 - Creation of preliminary proposal
 - Simple (grammar) changes are handled by chapter committees



Current Working Groups and Topics

- **Collectives & Topologies**
 - Torsten Hoefler, ETH
 - Andrew Lumsdaine, Indiana
- **Fault Tolerance**
 - Wesley Bland, ANL
 - Aurelien Bouteiller, UTK
 - Rich Graham, Mellanox
- **Fortran**
 - Craig Rasmussen, U. of Oregon
- **Generalized Requests**
 - Fab Tillier, Microsoft
- **Hybrid Models**
 - Pavan Balaji, ANL
- **I/O**
 - Quincey Koziol, HDF Group
 - Mohamad Chaarawi, HDF Group
- **Large count**
 - Jeff Hammond, Intel
- **Persistence**
 - Anthony Skjellum, U. of Alabama
- **Point to Point Comm.**
 - Dan Holmes, EPCC
 - Rich Graham, Mellanox
- **Remote Memory Access**
 - Bill Gropp, UIUC
 - Rajeev Thakur, ANL
- **Tools**
 - Kathryn Mohror, LLNL
 - Marc-Andre Hermans, RWTH Aachen
- **New working groups**
 - Added on demand
 - Support of 4 organizations at a physical MPI forum meeting

Typical Way to Add Items to MPI



1. **New items should be brought to a matching working group for discussion**
 - Creation of preliminary proposal
 - Simple (grammar) changes are handled by chapter committees
2. **Socializing of idea driven by the WG**
 - Could include plenary presentation to gather feedback
 - Focused on concepts not details like names or formal text
 - Make proposal easily available through WG wiki
 - Important to keep overall standard in mind
3. **Development of full proposal**
 - Latex version that fits into the standard
 - Creation of ticket to track voting
4. **MPI forum reading/voting process**

The MPI Voting Process

- **Quorum**

- 2/3 of eligible organizations have to be present
- 3/4 of present organization have to vote yes
- Goal: standardize only if there is consensus

- **Steps**

1. Reading: “Word by word” presentation to the forum
2. First vote
3. Second vote

- **Each step has to be at a separate physical meeting**

- Ensure people have time to think about additions
- Avoid hasty mistakes, which are hard to fix
- Prototypes are encouraged and helpful to convince people



Summary

- **MPI Forum is an open forum**
 - Everyone / every organization can join
 - Want/Need/Encourage community feedback
- **Major work in the next few years on MPI 4**
 - Currently three major initiatives
 - Fault Tolerance
 - Better support for hybrid programming
 - New models for point-to-point communication
 - Many smaller proposals as well
- **Get involved**
 - Let us know what you or your applications need
 - Let us know where MPI is lacking for your needs
 - Help close these gaps!

How To Get Involved

- **Submit comments to the MPI forum**
 - mpi-comments@mpi-forum.org
 - Feedback on prototypes / proposals as well as the existing standard
- **Subscribe to email lists to see what's going on**
 - Each working group has its own mailing list
- **Join a working group**
 - Check out the respective Wiki pages
 - Participate in WG meetings (typically phone conference)
 - Contact the WG chairs to introduce yourself
- **Participate in physical MPI forum meetings**
 - December 2014, San Jose, CA, USA
 - March 2015, Portland, OR, USA
 - Logistics and agendas available through the MPI forum website
 - Drop me an email if you have questions or are interested
- **More information at: <http://www.mpi-forum.org/>**

Summary

<http://www.mpi-forum.org/>

- **MPI Forum is an open forum**
 - Everyone / every organization can join
 - Want/Need/Encourage community feedback
- **Major work in the next few years on MPI 4**
 - Currently three major initiatives
 - Fault Tolerance
 - Better support for hybrid programming
 - New models for point-to-point communication
 - Many smaller proposals as well
- **Get involved**
 - Let us know what you or your applications need
 - Let us know where MPI is lacking for your needs
 - Help close these gaps!