

pt2pt wg

**FP16**

**MPI Forum Meeting  
Dec., 2018 in San Jose**

**Atsushi Hori  
RIKEN Center for Computational Science**

# Tickets and Goal

- #65 16-bit floating-point support for C/C++
  - ➔ Introducing FP16 Datatype
- #74 Datatype Naming Rule
  - ➔ Defining naming rules so that implementors can introduce new datatypes, and
  - ➔ new datatype aliases following the naming rule

★Which way to go ?

# Catch-up: Datatype Naming Rule

- Rationale of Naming Rule
  - One-to-one mapping between language datatype and MPI datatype
  - Users can easily identify MPI datatype from language datatype and vice versa.
  - Implementors can create a unique MPI datatype name from the language datatype
- Proposed Rule

**MPI\_<KIND>\_<TYPENAME>[\_<EXTRA>]**

  - Ex) MPI\_C\_FLOAT16, MPI\_F\_REAL2, ...

# Proposed aliases of the current datatype names

C datatype names	Aliases following the naming rule
MPI_CHAR	MPI_C_CHAR
MPI_SHORT	MPI_C_SHORT
MPI_SHORT_FLOAT	MPI_C_SHORT_FLOAT
MPI_INT	MPI_C_INT
MPI_LONG	MPI_C_LONG
MPI_LONG_LONG_INT	MPI_C_LONG_LONG
MPI_LONG_LONG	MPI_C_LONG_LONG
MPI_SIGNED_CHAR	MPI_C_SINGED_CHAR
MPI_UNSIGNED_CHAR	MPI_C_UNSIGNED_CHAR
MPI_UNSIGNED_SHORT	MPI_C_UNSIGNED_SHORT
MPI_UNSIGNED	MPI_C_UNSIGNED
MPI_UNSIGNED_LONG	MPI_C_UNSIGNED_LONG
constMPI_UNSIGNED_LONG_LONG	MPI_C_UNSIGNED_LONG_LONG
MPI_FLOAT	MPI_C_FLOAT
MPI_DOUBLE	MPI_C_DOUBLE
MPI_LONG_DOUBLE	MPI_C_LONG_DOUBLE
MPI_WCHAR	MPI_C_WCHAR
MPI_C_BOOL	MPI_C_C_BOOL
MPI_INT8_T	MPI_C_INT8_T
MPI_INT16_T	MPI_C_INT16_T
MPI_INT32_T	MPI_C_INT32_T
MPI_INT64_T	MPI_C_INT64_T
MPI_UINT8_T	MPI_C_UINT8_T
MPI_UINT16_T	MPI_C_UINT16_T
MPI_UINT32_T	MPI_C_UINT32_T
MPI_UINT64_T	MPI_C_UINT64_T
MPI_C_COMPLEX	-
MPI_C_FLOAT_COMPLEX	-
MPI_C_DOUBLE_COMPLEX	-
MPI_C_LONG_DOUBLE_COMPLEX	-

Table 4.3: Aliases of predefined MPI datatype names in C

Fortran datatype names	Aliases following the naming rule
MPI_CHARACTER	MPI_F_CHARACTER
MPI_LOGICAL	MPI_F_LOGICAL
MPI_INTEGER	MPI_F_INTEGER
MPI_REAL	MPI_F_REAL
MPI_DOUBLE_PRECISION	MPI_F_DOUBLE_PRECISION
MPI_COMPLEX	MPI_F_COMPLEX
MPI_DOUBLE_COMPLEX	MPI_F_DOUBLE_COMPLEX
(the followings are optional predefined datatypes)	
MPI_INTEGER1	MPI_F_INTEGER__1
MPI_INTEGER2	MPI_F_INTEGER__2
MPI_INTEGER4	MPI_F_INTEGER__4
MPI_INTEGER8	MPI_F_INTEGER__8
MPI_INTEGER16	MPI_F_INTEGER__16
MPI_REAL2	MPI_F_REAL__2
MPI_REAL4	MPI_F_REAL__4
MPI_REAL8	MPI_F_REAL__8
MPI_REAL16	MPI_F_REAL__16
MPI_COMPLEX4	MPI_F_COMPLEX__4
MPI_COMPLEX8	MPI_F_COMPLEX__8
MPI_COMPLEX16	MPI_F_COMPLEX__16
MPI_COMPLEX32	MPI_F_COMPLEX__32

Table 4.4: Aliases of predefined MPI datatype names in Fortran

# Current Status on FP16

- Architectures

- CPU (GPU)

- AArch64

FP16 is supported

- X86

FP16 conversion is supported

- Power9

FP16 conversion is supported

- Nvidia

FP16 is supported

- Network

- Mellanox

FP16 (and others) is (are) supported

- Language Standards

- C

`_Float16` (ISO/IEC JTC 1/SC 22/WG 14 N1896)

- C++

`_Half_float` (ISO/IEC JTC1 SC22 WG14 N2017)

- Compilers

- GCC

`__fp16` (ARM only)

- LLVM

`_Float16`

- Intel

Intrinsics to convert FP16 type

# MPI Implementations

- MPICH  
`MPIX_C_FLOAT16`
- Open MPI
  - Fujitsu/Mellanox `MPIX_SHORT_FLOAT`
  - Mellanox/Uber are running Horovod (TensorFlow) over patched Open MPI
    - CPU is not involved in reductions (GPU & Network)
- others ?

# Upcoming Datatype

- The **bfloat16** format is utilized in upcoming Intel AI processors, such as Nervana NNP-L1000, Xeon processors, and Intel FPGAs, Google Cloud TPUs, and TensorFlow.  
— <https://en.wikipedia.org/wiki/Minifloat>

# Summary of Current Status

- User demands are growing
  - MPI is considered as a communication library for AI/DL frameworks
  - not only AI/DL but also traditional HPC apps
- C language standard does not define FP16 yet
  - C working group proposal
- Not all major C compilers support FP16
- Some MPI implementations (will) have FP16 independent from the MPI standard
  - MPICH, Open MPI
- Yet another FP16 datatype (bfloat16) is on its way



# Major Arguments So Far

- #65 16-bit floating-point support for C/C++
  - We should wait until C standardize FP16
- #74 Datatype Naming Rule
  - We MAY lose portability

# My Objections

- **Support FP16 *as soon as possible***
  - It can be a good message to MPI users that MPI supports what users want when users want
  - Some MPI implementations already support FP16 out of the MPI standard
  - We force users to have another programming effort to utilize FP16 unless MPI supports FP16
- **Naming rule**
  - As long as major MPI implementations have good manners, no portability issue arises

# My Points

- **MPI supports FP16 ASAP**
- **Without having C standard, MPI datatype names cannot be defined**
- **Give implementors the freedom of introducing new datatypes following the datatype naming rule**

# Straw Votes

- ~~MPI should support FP16 as soon as possible~~
  - ~~Y[ ] N[ ] A[ ] O( )~~
- MPI should not define FP16 datatype until C standard defines it
  - Y[15] N[3] A[1] O( )
- ~~MPI should have FP16 datatype according to the C working group proposal - MPI\_FLOAT16 (MPI\_FLOAT32 and MPI\_FLOAT64)~~
  - ~~Y[ ] N[ ] A[ ] O( )~~
- ~~MPI should have the datatype naming rule~~
  - ~~Y[ ] N[ ] A[ ] O( )~~