

D R A F T

Document for a Standard Message-Passing Interface

Message Passing Interface Forum

February 5, 2015

This work was supported in part by NSF and ARPA under NSF contract CDA-9115428 and Esprit under project HPC Standards (21111).

This is the result of a LaTeX run of a draft of a single chapter of the MPIF Final Report document.

Chapter 19

Language Bindings Summary

In this section we summarize the specific bindings for C and Fortran. First we present the constants, type definitions, info values and keys. Then we present the routine prototypes separately for each binding. Listings are alphabetical within chapter.

19.1 Defined Values and Handles

19.1.1 Defined Constants

The C and Fortran names are listed below. Constants with the type `const int` may also be implemented as literal integer constants substituted by the preprocessor.

Error classes
C type: <code>const int</code> (or unnamed <code>enum</code>)
Fortran type: <code>INTEGER</code>
MPI_SUCCESS
MPI_ERR_BUFFER
MPI_ERR_COUNT
MPI_ERR_TYPE
MPI_ERR_TAG
MPI_ERR_COMM
MPI_ERR_RANK
MPI_ERR_REQUEST
MPI_ERR_ROOT
MPI_ERR_GROUP
MPI_ERR_OP
MPI_ERR_TOPOLOGY
MPI_ERR_DIMS
MPI_ERR_ARG
MPI_ERR_UNKNOWN
MPI_ERR_TRUNCATE
MPI_ERR_OTHER
MPI_ERR_INTERN
MPI_ERR_PENDING

(Continued on next page)

Error classes (continued)

C type: <code>const int</code> (or unnamed <code>enum</code>)
Fortran type: <code>INTEGER</code>

<code>MPI_T_ERR_CANNOT_INIT</code>
<code>MPI_T_ERR_NOT_INITIALIZED</code>
<code>MPI_T_ERR_MEMORY</code>
<code>MPI_T_ERR_INVALID</code>
<code>MPI_T_ERR_INVALID_INDEX</code>
<code>MPI_T_ERR_INVALID_ITEM</code>
<code>MPI_T_ERR_INVALID_SESSION</code>
<code>MPI_T_ERR_INVALID_HANDLE</code>
<code>MPI_T_ERR_INVALID_NAME</code>
<code>MPI_T_ERR_OUT_OF_HANDLES</code>
<code>MPI_T_ERR_OUT_OF_SESSIONS</code>
<code>MPI_T_ERR_CVAR_SET_NOT_NOW</code>
<code>MPI_T_ERR_CVAR_SET_NEVER</code>
<code>MPI_T_ERR_PVAR_NO_WRITE</code>
<code>MPI_T_ERR_PVAR_NO_STARTSTOP</code>
<code>MPI_T_ERR_PVAR_NO_ATOMIC</code>
<code>MPI_ERR_LASTCODE</code>

1
2
3
4
5
6
7 #400
8
9
10
11
12 #377

Buffer Address Constants

C type: <code>void * const</code>
Fortran type: (predefined memory location) ¹

<code>MPI_BOTTOM</code>
<code>MPI_IN_PLACE</code>

¹ Note that in Fortran these constants are not usable for initialization expressions or assignment. See Section 2.5.4.

Assorted Constants

C type: <code>const int</code> (or unnamed <code>enum</code>)
Fortran type: <code>INTEGER</code>

<code>MPI_PROC_NULL</code>
<code>MPI_ANY_SOURCE</code>
<code>MPI_ANY_TAG</code>
<code>MPI_UNDEFINED</code>
<code>MPI_BSEND_OVERHEAD</code>
<code>MPI_KEYVAL_INVALID</code>
<code>MPI_LOCK_EXCLUSIVE</code>
<code>MPI_LOCK_SHARED</code>
<code>MPI_ROOT</code>

No Process Message Handle

C type: <code>MPI_Message</code>
Fortran type: <code>INTEGER</code> or <code>TYPE(MPI_Message)</code>

<code>MPI_MESSAGE_NO_PROC</code>

47
48

C Constants Specifying Ignored Input (no Fortran)

C type: MPI_Fint*	equivalent to Fortran
MPI_F_STATUSES_IGNORE	MPI_STATUSES_IGNORE in mpi / mpif.h
MPI_F_STATUS_IGNORE	MPI_STATUS_IGNORE in mpi / mpif.h
C type: MPI_F08_status*	equivalent to Fortran
MPI_F08_STATUSES_IGNORE	MPI_STATUSES_IGNORE in mpi_f08
MPI_F08_STATUS_IGNORE	MPI_STATUS_IGNORE in mpi_f08

C preprocessor Constants and Fortran Parameters

C type: C-preprocessor macro that expands to an int value

Fortran type: INTEGER

MPI_SUBVERSION

MPI_VERSION

Null handles used in the MPI tool information interface

MPI_T_ENUM_NULL

MPI_T_enum

MPI_T_CVAR_HANDLE_NULL

MPI_T_cvar_handle

MPI_T_PVAR_HANDLE_NULL

MPI_T_pvar_handle

MPI_T_PVAR_SESSION_NULL

MPI_T_pvar_session

Verbosity Levels in the MPI tool information interface

C type: const int (or unnamed enum)

← No Fortran

MPI_T_VERBOSITY_USER_BASIC

MPI_T_VERBOSITY_USER_DETAIL

MPI_T_VERBOSITY_USER_ALL

MPI_T_VERBOSITY_TUNER_BASIC

MPI_T_VERBOSITY_TUNER_DETAIL

MPI_T_VERBOSITY_TUNER_ALL


MPI_T_VERBOSITY_MPIDEV_BASIC

MPI_T_VERBOSITY_MPIDEV_DETAIL

MPI_T_VERBOSITY_MPIDEV_ALL

#354


**Constants to identify associations of variables
in the MPI tool information interface**

C type: `const int` (or unnamed `enum`) 

#354

`MPI_T_BIND_NO_OBJECT`
`MPI_T_BIND_MPI_COMM`
`MPI_T_BIND_MPI_DATATYPE`
`MPI_T_BIND_MPI_ERRHANDLER`
`MPI_T_BIND_MPI_FILE`
`MPI_T_BIND_MPI_GROUP`
`MPI_T_BIND_MPI_OP`
`MPI_T_BIND_MPI_REQUEST`
`MPI_T_BIND_MPI_WIN`
`MPI_T_BIND_MPI_MESSAGE`
`MPI_T_BIND_MPI_INFO`

**Constants describing the scope of a control variable
in the MPI tool information interface**

C type: `const int` (or unnamed `enum`) 


#354

`MPI_T_SCOPE_CONSTANT`
`MPI_T_SCOPE_READONLY`
`MPI_T_SCOPE_LOCAL`
`MPI_T_SCOPE_GROUP`
`MPI_T_SCOPE_GROUP_EQ`
`MPI_T_SCOPE_ALL`
`MPI_T_SCOPE_ALL_EQ`

**Additional constants used
by the MPI tool information interface**

C type: `MPI_T_pvar_handle`
`MPI_T_PVAR_ALL_HANDLES`

**Performance variables classes used by the
MPI tool information interface**

C type: `const int` (or unnamed `enum`) 

#354

`MPI_T_PVAR_CLASS_STATE`
`MPI_T_PVAR_CLASS_LEVEL`
`MPI_T_PVAR_CLASS_SIZE`
`MPI_T_PVAR_CLASS_PERCENTAGE`
`MPI_T_PVAR_CLASS_HIGHWATERMARK`
`MPI_T_PVAR_CLASS_LOWWATERMARK`
`MPI_T_PVAR_CLASS_COUNTER`
`MPI_T_PVAR_CLASS_AGGREGATE`
`MPI_T_PVAR_CLASS_TIMER`
`MPI_T_PVAR_CLASS_GENERIC`

19.1.2 Types

The following are defined C type definitions, included in the file `mpi.h`.

```

/* C opaque types */
MPI_Aint
MPI_Count
MPI_Fint
MPI_Offset
MPI_Status
MPI_F08_status

/* C handles to assorted structures */
MPI_Comm
MPI_Datatype
MPI_Errhandler
MPI_File
MPI_Group
MPI_Info
MPI_Message
MPI_Op
MPI_Request
MPI_Win

/* Types for the MPI_T interface */
MPI_T_enum
MPI_T_cvar_handle
MPI_T_pvar_handle
MPI_T_pvar_session

The following are defined Fortran type definitions, included in the mpi_f08 and mpi
modules.

! Fortran opaque types in the mpi_f08 and mpi modules
TYPE(MPI_Status)

! Fortran handles in the mpi_f08 and mpi modules
TYPE(MPI_Comm)
TYPE(MPI_Datatype)
TYPE(MPI_Errhandler)
TYPE(MPI_File)
TYPE(MPI_Group)
TYPE(MPI_Info)
TYPE(MPI_Message)
TYPE(MPI_Op)
TYPE(MPI_Request)
TYPE(MPI_Win)

```

#345

19.1.3 Prototype Definitions

C Bindings

The following are defined C typedefs for user-defined functions, also included in the file `mpi.h`.

```

/* prototypes for user-defined functions */
typedef void MPI_User_function(void *invec, void *inoutvec, int *len,
                               MPI_Datatype *datatype);

typedef int MPI_Comm_copy_attr_function(MPI_Comm oldcomm,
                                         int comm_keyval, void *extra_state, void *attribute_val_in,
                                         void *attribute_val_out, int *flag);
typedef int MPI_Comm_delete_attr_function(MPI_Comm comm,
                                         int comm_keyval, void *attribute_val, void *extra_state);

typedef int MPI_Win_copy_attr_function(MPI_Win oldwin, int win_keyval,
                                         void *extra_state, void *attribute_val_in,
                                         void *attribute_val_out, int *flag);
typedef int MPI_Win_delete_attr_function(MPI_Win win, int win_keyval,
                                         void *attribute_val, void *extra_state);

typedef int MPI_Type_copy_attr_function(MPI_Datatype oldtype,
                                         int type_keyval, void *extra_state,
                                         void *attribute_val_in, void *attribute_val_out, int *flag);
typedef int MPI_Type_delete_attr_function(MPI_Datatype datatype,
                                         int type_keyval, void *attribute_val, void *extra_state);

typedef void MPI_Comm_errhandler_function(MPI_Comm *, int *, ...);
typedef void MPI_Win_errhandler_function(MPI_Win *, int *, ...);
typedef void MPI_File_errhandler_function(MPI_File *, int *, ...);

typedef int MPI_Grequest_query_function(void *extra_state,
                                         MPI_Status *status);
typedef int MPI_Grequest_free_function(void *extra_state);
typedef int MPI_Grequest_cancel_function(void *extra_state, int complete);

typedef int MPI_Datarep_extent_function(MPI_Datatype datatype,
                                         MPI_Aint *file_extent, void *extra_state);
typedef int MPI_Datarep_conversion_function(void *userbuf,
                                         MPI_Datatype datatype, int count, void *filebuf,
                                         MPI_Offset position, void *extra_state);

```

Fortran 2008 Bindings with the `mpi_f08` Module

The callback prototypes when using the Fortran `mpi_f08` module are shown below:

The user-function argument to `MPI_Op_create` should be declared according to:

```
ABSTRACT INTERFACE
```



```

SUBROUTINE MPI_User_function(invec, inoutvec, len, datatype)
  USE, INTRINSIC :: ISO_C_BINDING, ONLY : C_PTR
  TYPE(C_PTR), VALUE :: invec, inoutvec
  INTEGER :: len
  TYPE(MPI_Datatype) :: datatype

```

← BIND(C) removed
in all ABSTRACT
INTERFACE
definitions

The copy and delete function arguments to MPI_Comm_create_keyval should be declared according to:

ABSTRACT INTERFACE

```

SUBROUTINE MPI_Comm_copy_attr_function(oldcomm, comm_keyval, extra_state,
  attribute_val_in, attribute_val_out, flag, ierror)
  TYPE(MPI_Comm) :: oldcomm
  INTEGER :: comm_keyval, ierror
  INTEGER(KIND=MPI_ADDRESS_KIND) :: extra_state, attribute_val_in,
  attribute_val_out
  LOGICAL :: flag

```

ABSTRACT INTERFACE

```

SUBROUTINE MPI_Comm_delete_attr_function(comm, comm_keyval,
  attribute_val, extra_state, ierror)
  TYPE(MPI_Comm) :: comm
  INTEGER :: comm_keyval, ierror
  INTEGER(KIND=MPI_ADDRESS_KIND) :: attribute_val, extra_state

```

The copy and delete function arguments to MPI_Win_create_keyval should be declared according to:

ABSTRACT INTERFACE

```

SUBROUTINE MPI_Win_copy_attr_function(oldwin, win_keyval, extra_state,
  attribute_val_in, attribute_val_out, flag, ierror)
  TYPE(MPI_Win) :: oldwin
  INTEGER :: win_keyval, ierror
  INTEGER(KIND=MPI_ADDRESS_KIND) :: extra_state, attribute_val_in,
  attribute_val_out
  LOGICAL :: flag

```

ABSTRACT INTERFACE

```

SUBROUTINE MPI_Win_delete_attr_function(win, win_keyval, attribute_val,
  extra_state, ierror)
  TYPE(MPI_Win) :: win
  INTEGER :: win_keyval, ierror
  INTEGER(KIND=MPI_ADDRESS_KIND) :: attribute_val, extra_state

```

The copy and delete function arguments to MPI_Type_create_keyval should be declared according to:

ABSTRACT INTERFACE

```

SUBROUTINE MPI_Type_copy_attr_function(oldtype, type_keyval, extra_state,
  attribute_val_in, attribute_val_out, flag, ierror)
  TYPE(MPI_Datatype) :: oldtype
  INTEGER :: type_keyval, ierror
  INTEGER(KIND=MPI_ADDRESS_KIND) :: extra_state, attribute_val_in,

```

```

1      attribute_val_out
2      LOGICAL :: flag
3
4  ABSTRACT INTERFACE
5      SUBROUTINE MPI_Type_delete_attr_function(datatype, type_keyval,
6      attribute_val, extra_state, ierror)
7          TYPE(MPI_Datatype) :: datatype
8          INTEGER :: type_keyval, ierror
9          INTEGER(KIND=MPI_ADDRESS_KIND) :: attribute_val, extra_state

```

The handler-function argument to MPI_Comm_create_errhandler should be declared like this:

```

12 ABSTRACT INTERFACE
13     SUBROUTINE MPI_Comm_errhandler_function(comm, error_code)
14         TYPE(MPI_Comm) :: comm
15         INTEGER :: error_code

```

The handler-function argument to MPI_Win_create_errhandler should be declared like this:

```

18 ABSTRACT INTERFACE
19     SUBROUTINE MPI_Win_errhandler_function(win, error_code)
20         TYPE(MPI_Win) :: win
21         INTEGER :: error_code

```

The handler-function argument to MPI_File_create_errhandler should be declared like this:

```

25 ABSTRACT INTERFACE
26     SUBROUTINE MPI_File_errhandler_function(file, error_code)
27         TYPE(MPI_File) :: file
28         INTEGER :: error_code

```

The query, free, and cancel function arguments to MPI_Grequest_start should be declared according to:

```

32 ABSTRACT INTERFACE
33     SUBROUTINE MPI_Grequest_query_function(extra_state, status, ierror)
34         TYPE(MPI_Status) :: status
35         INTEGER :: ierror
36         INTEGER(KIND=MPI_ADDRESS_KIND) :: extra_state

```

```

37 ABSTRACT INTERFACE
38     SUBROUTINE MPI_Grequest_free_function(extra_state, ierror)
39         INTEGER :: ierror
40         INTEGER(KIND=MPI_ADDRESS_KIND) :: extra_state

```

```

42 ABSTRACT INTERFACE
43     SUBROUTINE MPI_Grequest_cancel_function(extra_state, complete, ierror)
44         INTEGER :: ierror
45         INTEGER(KIND=MPI_ADDRESS_KIND) :: extra_state
46         LOGICAL :: complete

```

The extent and conversion function arguments to MPI_Register_datarep should be de-

clared according to:

ABSTRACT INTERFACE

```
SUBROUTINE MPI_Datarep_extent_function(datatype, extent, extra_state,
    ierror)
```

```
    TYPE(MPI_Datatype) :: datatype
    INTEGER(KIND=MPI_ADDRESS_KIND) :: extent, extra_state
    INTEGER :: ierror
```

ABSTRACT INTERFACE

```
SUBROUTINE MPI_Datarep_conversion_function(userbuf, datatype, count,
    filebuf, position, extra_state, ierror)
```

```
    USE, INTRINSIC :: ISO_C_BINDING, ONLY : C_PTR
    TYPE(C_PTR), VALUE :: userbuf, filebuf
    TYPE(MPI_Datatype) :: datatype
    INTEGER :: count, ierror
    INTEGER(KIND=MPI_OFFSET_KIND) :: position
    INTEGER(KIND=MPI_ADDRESS_KIND) :: extra_state
```

Fortran Bindings with mpif.h or the mpi Module

With the Fortran `mpi` module or `mpif.h`, here are examples of how each of the user-defined subroutines should be declared.

The user-function argument to `MPI_OP_CREATE` should be declared like this:

```
SUBROUTINE USER_FUNCTION(INVEC, INOUTVEC, LEN, DATATYPE)
```

```
    <type> INVEC(LEN), INOUTVEC(LEN)
    INTEGER LEN, DATATYPE
```

The copy and delete function arguments to `MPI_COMM_CREATE_KEYVAL` should be declared like these:

```
SUBROUTINE COMM_COPY_ATTR_FUNCTION(OLDCOMM, COMM_KEYVAL, EXTRA_STATE,
```

```
    ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERROR)
```

```
    INTEGER OLDCOMM, COMM_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) EXTRA_STATE, ATTRIBUTE_VAL_IN,
        ATTRIBUTE_VAL_OUT
    LOGICAL FLAG
```

```
SUBROUTINE COMM_DELETE_ATTR_FUNCTION(COMM, COMM_KEYVAL, ATTRIBUTE_VAL,
    EXTRA_STATE, IERROR)
```

```
    INTEGER COMM, COMM_KEYVAL, IERROR
    INTEGER(KIND=MPI_ADDRESS_KIND) ATTRIBUTE_VAL, EXTRA_STATE
```

The copy and delete function arguments to `MPI_WIN_CREATE_KEYVAL` should be declared like these:

```
SUBROUTINE WIN_COPY_ATTR_FUNCTION(OLDWIN, WIN_KEYVAL, EXTRA_STATE,
```

```
    ATTRIBUTE_VAL_IN, ATTRIBUTE_VAL_OUT, FLAG, IERROR)
```

```
    INTEGER OLDWIN, WIN_KEYVAL, IERROR
```

#388

19.1.5 Info Keys

The following info keys are reserved. They are strings.

```

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48

```

red
bar:
#347

access_style
accumulate_ops
accumulate_ordering
alloc_shared_noncontig
appnum
arch
cb_block_size
cb_buffer_size
cb_nodes
chunked_item
chunked_size
chunked
collective_buffering
file_perm
filename
file
host
io_node_list
ip_address
ip_port
nb_proc
no_locks
num_io_nodes
path
same_disp_unit
same_size
soft
striping_factor
striping_unit
wdir

#369
#347

19.1.6 Info Values

The following info values are reserved. They are strings.

```

1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48

```

false
random
rar
raw
read_mostly
read_once
reverse_sequential
same_op
same_op_no_op
sequential

#347	true	1
	war	2
	waw	3
	write_mostly	4
	write_once	5
		6
		7
		8
		9
		10
		11
		12
		13
		14
		15
		16
		17
		18
		19
		20
		21
		22
		23
		24
		25
		26
		27
		28
		29
		30
		31
		32
		33
		34
		35
		36
		37
		38
		39
		40
		41
		42
		43
		44
		45
		46
		47
		48