

Index: io-2.tex

```
=====
--- io-2.tex      (revision 1903)
+++ io-2.tex      (working copy)
@@ -1090,12 +1090,14 @@
\emph{offsets} &
& \mpifunc{MPI\_FILE\_WRITE\_AT} & \mpifunc{MPI\_FILE\_WRITE\_AT
\_ALL} \\
\cline{2-4}
-& \emph{nonblocking} \&
- & \mpifunc{MPI\_FILE\_IREAD\_AT} & \mpifunc{MPI\_FILE\_READ\_AT
\_ALL\_BEGIN} \\
-& \emph{split collective}
- &
& \mpifunc{MPI\_FILE\_READ\_AT\_ALL
\_END} \\
-& & \mpifunc{MPI\_FILE\_IWRITE\_AT} & \mpifunc{MPI\_FILE\_WRITE\_AT
\_ALL\_BEGIN} \\
-& &
& \mpifunc{MPI\_FILE\_WRITE\_AT\_ALL
\_END} \\
+& \emph{nonblocking}
+ & \mpifunc{MPI\_FILE\_IREAD\_AT} & \mpifunc{MPI\_FILE\_IREAD\_AT
\_ALL} \\
+& & \mpifunc{MPI\_FILE\_IWRITE\_AT} & \mpifunc{MPI\_FILE\_IWRITE\_AT
\_ALL} \\
+\cline{2-4}
+& \emph{split collective} & {N/A} & \mpifunc{MPI\_FILE\_READ\_AT\_ALL
\_BEGIN} \\
+& & & \mpifunc{MPI\_FILE\_READ\_AT\_ALL\_END} \\
+& &
& \mpifunc{MPI\_FILE\_WRITE\_AT
\_ALL\_BEGIN} \\
+& &
& \mpifunc{MPI\_FILE\_WRITE\_AT
\_ALL\_END} \\
\hline %-----
\emph{individual} & \emph{blocking}
& \mpifunc{MPI\_FILE\_READ} & \mpifunc{MPI\_FILE\_READ\_ALL} \\
@@ -1102,12 +1104,14 @@
\emph{file pointers} &
& \mpifunc{MPI\_FILE\_WRITE} & \mpifunc{MPI\_FILE\_WRITE\_ALL} \\
\cline{2-4}
-& \emph{nonblocking} \&
- & \mpifunc{MPI\_FILE\_IREAD} & \mpifunc{MPI\_FILE\_READ\_ALL
\_BEGIN} \\
-& \emph{split collective}
- &
& \mpifunc{MPI\_FILE\_READ\_ALL\_END} \\
-& & \mpifunc{MPI\_FILE\_IWRITE} & \mpifunc{MPI\_FILE\_WRITE\_ALL
\_BEGIN} \\
-& &
& \mpifunc{MPI\_FILE\_WRITE\_ALL\_END} \\
+& \emph{nonblocking}
+ & \mpifunc{MPI\_FILE\_IREAD} & \mpifunc{MPI\_FILE\_IREAD\_ALL} \\
+& & \mpifunc{MPI\_FILE\_IWRITE} & \mpifunc{MPI\_FILE\_IWRITE\_ALL}
```

```

+\cline{2-4}
+& \emph{split collective} & {N/A} & \mpifunc{MPI\_FILE\_READ\_ALL
\_BEGIN} \\
+& & \mpifunc{MPI\_FILE\_READ\_ALL\_END} \\
+& & & \mpifunc{MPI\_FILE\_WRITE\_ALL
\_BEGIN} \\
+& & & \mpifunc{MPI\_FILE\_WRITE\_ALL
\_END} \\
\hline %-----
\emph{shared} & \emph{blocking}
& \mpifunc{MPI\_FILE\_READ\_SHARED} & \mpifunc{MPI\_FILE\_READ
\_ORDERED} \\
@@ -1114,12 +1118,14 @@
\emph{file pointer} &
& \mpifunc{MPI\_FILE\_WRITE\_SHARED} & \mpifunc{MPI\_FILE\_WRITE
\_ORDERED} \\
\cline{2-4}
-& \emph{nonblocking} \&
- & \mpifunc{MPI\_FILE\_IREAD\_SHARED} & \mpifunc{MPI\_FILE\_READ
\_ORDERED\_BEGIN} \\
-& \emph{split collective}
- & & \mpifunc{MPI\_FILE\_READ\_ORDERED
\_END} \\
-& & \mpifunc{MPI\_FILE\_IWRITE\_SHARED} & \mpifunc{MPI\_FILE\_WRITE
\_ORDERED\_BEGIN} \\
-& & & \mpifunc{MPI\_FILE\_WRITE\_ORDERED
\_END} \\
+& \emph{nonblocking}
+ & \mpifunc{MPI\_FILE\_IREAD\_SHARED} & {N/A} \\
+& & \mpifunc{MPI\_FILE\_IWRITE\_SHARED} & \\
+\cline{2-4}
+& \it split collective & {N/A} & \mpifunc{MPI\_FILE\_READ\_ORDERED
\_BEGIN} \\
+& & \mpifunc{MPI\_FILE\_READ\_ORDERED\_END} \\
+& & & \mpifunc{MPI\_FILE\_WRITE
\_ORDERED\_BEGIN} \\
+& & & \mpifunc{MPI\_FILE\_WRITE
\_ORDERED\_END} \\
\hline
\end{tabular}
\end{center}
@@ -1458,6 +1464,28 @@
\mpifunc{MPI\_FILE\_IREAD\_AT} is a nonblocking version
of the \mpifunc{MPI\_FILE\_READ\_AT} interface.

+\begin{funcdef}{MPI\_FILE\_IREAD\_AT\_ALL(fh, offset, buf, count,
datatype, request)}
+\funcarg{\IN}{fh}{file handle (handle)}
+\funcarg{\IN}{offset}{file offset (integer)}
+\funcarg{\OUT}{buf}{initial address of buffer (choice)}

```

```

+\funcarg{\IN}{count}{number of elements in buffer (integer)}
+\funcarg{\IN}{datatype}{datatype of each buffer element (handle)}
+\funcarg{\OUT}{request}{request object (handle)}
+\end{funcdef}
+
+\cdeclindex{MPI\_Request}%
+\cdeclindex{MPI\_File}%
+\cdeclindex{MPI\_Offset}%
+\mpibind{MPI\_File\_iread\_at\_all(MPI\_File~fh, MPI\_Offset~offset,
void~*buf, int~count, MPI\_Datatype~datatype, MPI\_Request~*request)}
+\mpifnewbind{MPI\_File\_iread\_at\_all(fh, offset, buf, count,
datatype, request, ierror) BIND(C) \fargs TYPE(MPI\_File),
INTENT(IN) :: fh \\ INTEGER(KIND=MPI\_OFFSET\_KIND), INTENT(IN) ::
offset \\ TYPE(*), DIMENSION(..), ASYNCHRONOUS :: buf \\ INTEGER,
INTENT(IN) :: count \\ TYPE(MPI\_Datatype), INTENT(IN) :: datatype \\
TYPE(MPI\_Request), INTENT(OUT) :: request \\ INTEGER, OPTIONAL,
INTENT(OUT) :: ierror}
+\mpifbind{MPI\_FILE\_IREAD\_AT\_ALL(FH, OFFSET, BUF, COUNT, DATATYPE,
REQUEST, IERROR)\fargs <type> BUF(*) \\ INTEGER FH, COUNT, DATATYPE,
REQUEST, IERROR \\ INTEGER(KIND=MPI\_OFFSET\_KIND) OFFSET}
+
+\mpifunc{MPI\_FILE\_IREAD\_AT\_ALL} is a nonblocking version of
+\mpifunc{MPI\_FILE\_READ\_AT\_ALL}. See
+Section~\ref{sec:io-semantics-nb-collective},
+page~\pageref{sec:io-semantics-nb-collective} for semantics of
nonblocking
+collective file operations.
+
\begin{funcdef}{MPI\_FILE\_IWRITE\_AT(fh, offset, buf, count,
datatype, request)}
\funcarg{\INOUT}{fh}{file handle (handle)}
\funcarg{\IN}{offset}{file offset (integer)}
@@ -1478,6 +1506,25 @@
\mpifunc{MPI\_FILE\_IWRITE\_AT} is a nonblocking version
of the \mpifunc{MPI\_FILE\_WRITE\_AT} interface.

+\begin{funcdef}{MPI\_FILE\_IWRITE\_AT\_ALL(fh, offset, buf, count,
datatype, request)}
+\funcarg{\INOUT}{fh}{file handle (handle)}
+\funcarg{\IN}{offset}{file offset (integer)}
+\funcarg{\IN}{buf}{initial address of buffer (choice)}
+\funcarg{\IN}{count}{number of elements in buffer (integer)}
+\funcarg{\IN}{datatype}{datatype of each buffer element (handle)}
+\funcarg{\OUT}{request}{request object (handle)}
+\end{funcdef}
+
+\cdeclindex{MPI\_Request}%
+\cdeclindex{MPI\_File}%
+\cdeclindex{MPI\_Offset}%
+\mpibind{MPI\_File\_iwrite\_at\_all(MPI\_File~fh, MPI\_Offset~offset,

```

```

const void~*buf, int~count, MPI\_Datatype~datatype, MPI
\_Request~*request)}
+ \mpifnewbind{MPI\_File\_iwrite\_at\_all(fh, offset, buf, count,
datatype, request, ierror) BIND(C) \fargs TYPE(MPI\_File),
INTENT(IN) :: fh \\ INTEGER(KIND=MPI\_OFFSET\_KIND), INTENT(IN) ::
offset \\ TYPE(*), DIMENSION(..), INTENT(IN), ASYNCHRONOUS :: buf \\
INTEGER, INTENT(IN) :: count \\ TYPE(MPI\_Datatype), INTENT(IN) ::
datatype \\ TYPE(MPI\_Request), INTENT(OUT) :: request \\ INTEGER,
OPTIONAL, INTENT(OUT) :: ierror}
+ \mpifbind{MPI\_FILE\_IWRITE\_AT\_ALL(FH, OFFSET, BUF, COUNT,
DATATYPE, REQUEST, IERROR) \fargs <type> BUF(*) \\ INTEGER FH, COUNT,
DATATYPE, REQUEST, IERROR \\ INTEGER(KIND=MPI\_OFFSET\_KIND) OFFSET}
+
+ \mpifunc{MPI\_FILE\_IWRITE\_AT\_ALL} is a nonblocking version of
+ \mpifunc{MPI\_FILE\_WRITE\_AT\_ALL}.
+
\subsection{Data Access with Individual File Pointers}
%-----
\label{sec:io-indiv-ptr}
@@ -1691,6 +1738,23 @@
\end{verbatim}
\end{example}

+ \begin{funcdef}{MPI\_FILE\_IREAD\_ALL(fh, buf, count, datatype,
request)}
+ \funcarg{\INOUT}{fh}{file handle (handle)}
+ \funcarg{\OUT}{buf}{initial address of buffer (choice)}
+ \funcarg{\IN}{count}{number of elements in buffer (integer)}
+ \funcarg{\IN}{datatype}{datatype of each buffer element (handle)}
+ \funcarg{\OUT}{request}{request object (handle)}
+ \end{funcdef}
+
+ \cdeclindex{MPI\_Request}%
+ \cdeclindex{MPI\_File}%
+ \mpibind{MPI\_File\_iread\_all(MPI\_File~fh, void~*buf, int~count,
MPI\_Datatype~datatype, MPI\_Request~*request)}
+ \mpifnewbind{MPI\_File\_iread\_all(fh, buf, count, datatype, request,
ierror) BIND(C) \fargs TYPE(MPI\_File), INTENT(IN) :: fh \\ TYPE(*),
DIMENSION(..), ASYNCHRONOUS :: buf \\ INTEGER, INTENT(IN) :: count \\
TYPE(MPI\_Datatype), INTENT(IN) :: datatype \\ TYPE(MPI\_Request),
INTENT(OUT) :: request \\ INTEGER, OPTIONAL, INTENT(OUT) :: ierror}
+ \mpifbind{MPI\_FILE\_IREAD\_ALL(FH, BUF, COUNT, DATATYPE, REQUEST,
IERROR) \fargs <type> BUF(*) \\ INTEGER FH, COUNT, DATATYPE, REQUEST,
IERROR}
+
+ \mpifunc{MPI\_FILE\_IREAD\_ALL} is a nonblocking version
+ of \mpifunc{MPI\_FILE\_READ\_ALL}.
+
\begin{funcdef}{MPI\_FILE\_IWRITE(fh, buf, count, datatype, request)}
\funcarg{\INOUT}{fh}{file handle (handle)}

```

```
\funcarg{\IN}{buf}{initial address of buffer (choice)}  
@@ -1708,6 +1772,23 @@
```

\mpifunc{MPI_FILE_IWRITE} is a nonblocking version of the
\mpifunc{MPI_FILE_WRITE} interface.

```
+ \begin{funcdef}{MPI\_FILE\_IWRITE\_ALL(fh, buf, count, datatype,  
request)}  
+ \funcarg{\INOUT}{fh}{file handle (handle)}  
+ \funcarg{\IN}{buf}{initial address of buffer (choice)}  
+ \funcarg{\IN}{count}{number of elements in buffer (integer)}  
+ \funcarg{\IN}{datatype}{datatype of each buffer element (handle)}  
+ \funcarg{\OUT}{request}{request object (handle)}  
+ \end{funcdef}  
+  
+ \cdeclindex{MPI\_Request}%  
+ \cdeclindex{MPI\_File}%  
+ \mpibind{MPI\_File\_iwrite\_all(MPI\_File~fh, const void~*buf,  
int~count, MPI\_Datatype~datatype, MPI\_Request~*request)}  
+ \mpifnewbind{MPI\_File\_iwrite\_all(fh, buf, count, datatype,  
request, ierror) BIND(C) \fargs TYPE(MPI\_File), INTENT(IN) :: fh \\  
TYPE(*), DIMENSION(..), INTENT(IN), ASYNCHRONOUS :: buf \\  
INTEGER, INTENT(IN) :: count \\  
TYPE(MPI\_Datatype), INTENT(IN) :: datatype \\  
TYPE(MPI\_Request), INTENT(OUT) :: request \\  
INTEGER, OPTIONAL,  
INTENT(OUT) :: ierror}  
+ \mpifbind{MPI\_FILE\_IWRITE\_ALL(FH, BUF, COUNT, DATATYPE, REQUEST,  
IERROR)\fargs <type> BUF(*) \\  
INTEGER FH, COUNT, DATATYPE, REQUEST,  
IERROR}  
+  
+ \mpifunc{MPI\_FILE\_IWRITE\_ALL} is a nonblocking version  
+ of \mpifunc{MPI\_FILE\_WRITE\_ALL}.  
+  
+ \begin{funcdef}{MPI\_FILE\_SEEK(fh, offset, whence)}  
+ \funcarg{\INOUT}{fh}{file handle (handle)}  
+ \funcarg{\IN}{offset}{file offset (integer)}  
@@ -3339,6 +3420,37 @@
```

Different processes can pass different values for other arguments
of a collective routine unless specified otherwise.

```
+ \subsection{Nonblocking Collective File Operations}  
+ %-----  
+ \label{sec:io-semantics-nb-collective}  
+ Nonblocking collective file operations are defined only for data  
access  
+ routines with explicit offsets and individual file pointers but not  
with  
+ shared file pointers.  
+  
+ Nonblocking collective file operations are subject to the same  
restrictions as
```

+blocking collective I/O operations. All processes belonging to the group of

+the communicator that was used to open the file must call collective I/O

+operations (blocking and nonblocking) in the same order. This is consistent

+with the ordering rules for collective operations in threaded environments.

+For a complete discussion, please refer to the semantics set forth in

+Section~\ref{coll:correct} on page~\pageref{coll:correct}.

+

+Nonblocking collective I/O operations do not match with blocking collective

+I/O operations. Multiple nonblocking collective I/O operations can be

+outstanding on a single file handle. High quality MPI implementations should

+be able to support a large number of pending nonblocking I/O operations.

+

+All nonblocking collective I/O calls are local and return immediately,

+irrespective of the status of other processes. The call initiates the

+operation which may progress independently of any communication, computation,

+or I/O. The call returns a request handle, which must be passed to a

+completion call. Input buffers should not be modified and output buffers

+should not be accessed before the completion call returns. The same progress rules

+described for nonblocking collective operations apply for nonblocking

+collective I/O operations. For a complete discussion, please refer to the

+semantics set forth in Section~\ref{sec:nbcoll} on

+page~\pageref{sec:nbcoll}.

+

\subsection{Type Matching}

%-----

@@ -3750,6 +3862,21 @@

%%SKIP

%%ENDHEADER

\begin{verbatim}

+MPI_File_iread_all(fh,...) ;

+MPI_File_iread_all(fh,...) ;

+MPI_Waitall(...) ;

+\end{verbatim}

+

+In addition, as mentioned in Section~\ref{sec:io-semantics-nb-collective} on

+page~\pageref{sec:io-semantics-nb-collective}, nonblocking collective

I/O

+operations have to be called in the same order on the file handle by all

+processes.

+

+Similar considerations apply to conflicting accesses of the form:

+%%HEADER

+%%SKIP

+%%ENDHEADER

+\begin{verbatim}

 MPI_File_write_all_begin(fh,...) ;

 MPI_File_iread(fh,...) ;

 MPI_Wait(fh,...) ;