



MPI for High-Level Languages

Douglas Gregor: did all the work

Jeff Squyres: encouraged Doug

Torsten Hoefler: got suckered in after the fact

Andrew Lumsdaine: paid for it

October 2008



MPI for High-Level Languages

Premise

- MPI Forum should only define basic C and Fortran bindings
 - Let other language communities devise their own bindings / libraries
- But MPI needs a few more hooks
 - Otherwise, high-level languages will not have what they need to create higher-level libraries
- If all this occurs, deprecate (and eventually remove) the current C++ bindings

October 2008



MPI for High-Level Languages

Rationale

- C++ bindings aren't used much
 - Boost.mpi is much "better"
 - Designed/maintained by C++ community
- Same will hold true for other languages
 - Including advanced Fortran usage
- MPI Forum doesn't have time/resources to debate + maintain other language bindings
 - Java, Python, C#, Ruby, ...etc.

October 2008



MPI for High-Level Languages

What's Missing

- **Assumption:** threading is going to become more important over time
- Transmission of objects
 - Variable-length [serialized] data
- User-defined operation support
 - Associate state with MPI_Op's
- Memory management / garbage collection
 - Send buffer access / multiple heaps

October 2008



MPI for High-Level Languages



Transmission of Objects

October 2008



MPI for High-Level Languages

Transmission of Objects

- We already have pack/unpack
 - Works fine
- But variable-length data is the bigger issue
 - Currently requires 2 sends, posting really large receives (not always possible), or probe (not thread safe)
 - The “Mprobe” proposal fixes this issue by divorcing matching from receiving

October 2008



MPI for High-Level Languages

Variable Length Pt2Pt Data

```
// do the match
MPI_Message msg;
MPI_Mprobe(..., &msg, &status)
// extract the size
MPI_Get_count(&status, MPI_BYTE, &count)
// malloc for that size
char *buffer = malloc(count);
// do the receive
MPI_Mrecv(buffer, count, MPI_BYTE, &msg, ...);
// deserialize the message
```

October 2008



MPI for High-Level Languages

Variable Length Coll. Data

- Currently no provisions for this
 - Reductions are problematic
 - Example: MPI_ALLREDUCE on strings to concatenate into a big string
- Problems:
 - Collectives take a fixed count
 - Can bcast the count first, but may not know final size until reduction is performed
 - MPI_User_function takes a fixed count and cannot realloc inoutvec

October 2008



MPI for High-Level Languages

Variable Length 1Sided Data

- No provisions or proposals at this time
- Follow the one-sided committee and see what happens

October 2008



MPI for High-Level Languages

Variable Length Solution: Blobs

```
struct MPI_Blob {  
    MPI_Aint length; // or some new type  
    MPI_Aint address;  
};
```

- Can use blobs for pt2pt, collectives
 - Per previous, wait and see for one-sided

October 2008



MPI for High-Level Languages

Blobs and Reductions

```
char *string = ...;
MPI_Blob out, in;

// setup "out" blob
out.length = strlen(string);
out.address = string;

// Do the reduce, get result in a new blob
MPI_Allreduce(&out, &in, 1, MPI_BLOB,
concatOp, MPI_COMM_WORLD);
```

October 2008



MPI for High-Level Languages

Blob Restrictions

- Cannot be used with `MPI_TYPE_CREATE_STRUCT`
- Cannot be used with `MPI_IN_PLACE` (?)
- Where does the memory come from?
 - Implementation has to alloc blob memory that is given back to the user ("in" blob in previous example)
 - Require `MPI_ALLOC/FREE_MEM`?

October 2008



MPI for High-Level Languages



User-Defined Operations

October 2008



MPI for High-Level Languages

User-Defined Operations

- Problematic for multi-threaded
 - No way to associate state with an MPI_Op (e.g, an object)
- Do not want to change all reduction collective function signatures
- Simple solution
 - Allow associating a (void*) to new MPI_Op creation function

October 2008



MPI for High-Level Languages

User-Defined Operations

```
typedef void MPI_User_function_data(void
    *invec, void *inoutvec, int *len,
    MPI_Datatype *datatype, void *data);

int MPI_Op_create_data(MPI_User_function_data
    *function, int commute, MPI_Op *op, void
    *data);
```

- Assume making MPI_Op's is fast / cheap
 - Can do it frequently for multiple different (void *data)'s

October 2008



MPI for High-Level Languages

One-Sided MPI_Op's

- Not possible with MPI-2
MPI_ACCUMULATE
- Follow prior recommendation:
 - See what happens in the MPI-3 RMA WG

October 2008



MPI for High-Level Languages



Memory Management / Garbage Collection

October 2008



MPI for High-Level Languages

Multi-Send Buffer Access

- Garbage collectors may scan memory that is being used in non-blocking sends
- Microsoft's proposal addresses this issue
 - Trac ticket #45
 - Looks like this ticket will be included in 2.2

October 2008



MPI for High-Level Languages

Multiple Heaps

- Some languages have 2 heaps:
 - Garbage collected heap
 - System heap (i.e., what the underlying C MPI implementation uses)
- If language provides no way to “pin” memory in GC heap, serialized (blobs) may need to copy between the heaps

October 2008



MPI for High-Level Languages

Multiple Heaps

- If language cannot read from system heap, MPI may need to do an additional copy (and free from system heap)
 - E.g., de-serializing a received BLOB
- Random note: memory allocations on one heap can affect the other

October 2008



MPI for High-Level Languages

So... What's the Problem?

- Some MPI extensions may be needed to help eliminate the need for extraneous copies of serialized data between system and GC heaps.
- May also be useful with some networks that require registered memory...?
- Don't know exactly what is needed yet
 - Still working on it...
 - Feedback would be appreciated

October 2008



MPI for High-Level Languages



Moving Forward

October 2008



MPI for High-Level Languages

Moving Forward

- Some points in this proposal are good
 - Others need [a lot] more discussion
 - Need language specialists for these issues!
- Probably need a group where to discuss these issues
 - New WG?
 - Rename Fortran WG → Language WG?
 - Still have some Fortran-specific issues to discuss

October 2008



MPI for High-Level Languages



Additional Issues

October 2008



MPI for High-Level Languages

Datatypes

- What do we do about C++ datatypes?
 - Complex is not necessarily the same as complex is not necessarily the same as complex...
- Is there any [other] C++ functionality that we lose if we ditch the C++ bindings?
 - `THROW_EXCEPTIONS` can still be implemented on top of the C bindings