

MPI-3 One Sided

- 😊 some simple changes 😊 -
 - half-baked -
-

Torsten Hoefler

Open Systems Lab
Indiana University
Bloomington, IN, USA

Introduction to MPI-2 One Sided

- Two concepts and three calls
 - Central concepts: Window + Epochs
 - Data manipulation calls:
 - `MPI_Put(src_addr, src_cnt, src_type, tgt_rank, tgt_disp, tgt_cnt, tgt_type, win)`
 - `MPI_Get(src_addr, src_cnt, src_type, tgt_rank, tgt_disp, tgt_cnt, tgt_type, win)`
 - `MPI_Accumulate(src_addr, src_cnt, src_type, tgt_rank, tgt_disp, tgt_cnt, tgt_type, op, win)`
- Looks very similar to Active Messages
 - Possible implementation:
 - Put and accumulate are special handlers
 - Get triggers a put message on remote end
 - Interface needs to support hardware
 - RDMA and Shared Memory



Do we need Memory Windows?

□ Windows enable:

- Protection
- Process-local offsets (heterogeneity)
- Group contexts (cf. Communicators)
- Explicit memory exposure helps hw support

□ Windows require:

- $\Omega(P)$ offset translation tables (likely if RDMA is used)
- dense collective semantics
 - dense access is assumed
 - lazy translation entry fetching?
 - scalability problems



What do we really need?

- ❑ Do we want to (can we) give this up?
 - ❑ Shared memory needs mapping for direct access
 - `shm_open()`, `mmap()`
 - ❑ RDMA interconnects are not first class citizens (yet)
 - Need memory registration (pinning, explicit addr. transl.)
 - Future developments might use IOMMUs (page req. interf.)
- ❑ We need memory exposure operation
 - ❑ often referred to as “memory registration”
 - ❑ GASNet+ARMCI also require special memory
- ❑ We don't need this collectively though
 - ❑ Some apps might benefit from collective semantics?



Proposal (P2P Exposure)

- MPI-2 Memory windows expose collectively
 - P2P exposure can increase scalability -> local windows
- MPI_Win_create_local(base, size, displ, info, comm, win)
 - Local call to create a local window object
 - We might want to add a test for the type of win (not included yet)
 - MPI_Win_free() semantics change with type of win (coll. or not)
- Memory registration vs. registered allocation
 - Window creation takes existing memory
 - Some underlying APIs have strange requirements
 - posix shmattr wants page-aligned addresses or rounds addr. up
 - MPI RMA is not transparent → CPU can translate addresses
 - however, addr-translation can be simplified with special allocators



Proposal (AM Emulation)

- Allow user-defined ops in MPI_Accumulate()
 - requires **either** sending of MPI_Ops **or**
 - collective MPI_Op_create()
 - remember: ops *may have side effects*
- Differences to “real” AMs
 - can not trigger new messages (one-level)
 - have fixed-size vector input and output
 - access to more complex structures, e.g., lists or queues, needs to be “emulated” (hacked via side effects ;-))
 - MPI guarantees atomicity (this is good and bad)
 - no guaranteed asynchronous progr. (epochs)



Proposal (MPI_Get_accumulate)

- ❑ To implement lock-free remote atomics
 - (test&set, test&accumulate)
- ❑ Avoids unnecessary synchronization
- ❑ Allows concurrent accesses to same mem
 - unlike MPI_Put(), MPI_Get() which have undefined outcome in the current proposal
- ❑ Can play tricks with MPI_REPLACE or a new op MPI_OP_NULL (? -- simple get)
- ❑ Can easily be optimized for predefined ops!
 - Many interconnects (IB) do so already



What are Epochs (good for)?

- ❑ Epochs enable:
 - ❑ BSP-like bulk synchronicity
 - ❑ implementation on top of Message Passing (S/R)
 - ❑ batch nonblocking accesses without requests
- ❑ Synchronization modes:
 - ❑ define consistency and scope of epochs
 - ❑ they seem rather complex in MPI-2
- ❑ Three modes:
 - ❑ fence – BSP-like global synch (dynamic patterns)
 - ❑ start-complete/post-wait – p2p synch (fixed patterns)
 - ❑ lock-unlock – target-specific epoch between lock/unlock (passive target)



Bonachea's and Duell's criticism

- Do not need collective semantics
 - they chose passive target mode (lock/unlock)
- 1. Window creation is collective
 - hinders efficient exposure for local objects
 - no “sparse” communication
 - adding local windows
- 2. Exposed memory must be MPI_Alloc_mem()'d
 - no exposure of static memory or stack-variables
 - alloc_mem might be limited by the implementation
 - force better guarantees on MPI_Alloc_mem?
 - static mem and stack variables remain problematic (is ok?)



Bonachea's and Duell's criticism ff

3. Forbids conflicting get/put (or local load/store) accesses to same memory
 - really hard to track for compilers (halting problem?)
 - easy source of bugs in user codes
 - conflicting accesses are undefined not erroneous (no atomicity!)
4. Window's memory may not be updated by remote gets and local stores concurrently
 - simplifies MPI implementation significantly
 - seems very artificial and suboptimal from user's perspective
 - conflicting accesses are undefined not erroneous (no atomicity!)



Bonachea's and Duell's criticism ff


5. Overlapping memory regions of multiple windows can be created but not be used
 - “concurrent communications may lead to erroneous results”
 - conflicting accesses are undefined not erroneous (no atomicity!)

6. Passive target RMA ops only lock a single process during an epoch
 - ops from one source to different targets are serialized
 - one window for each target to enable concurrent access?
 - scalability limitation
 - use one local memory window for each process



What about the access calls?

- Put() and Get() are simple/sufficient
 - Lots of arguments though
 - Put()/Get(): 8; Acc(): 9; Get_acc(): 9 or 12
 - test: call a function 10^5 times with n pointer-args



MacOS results are reproducible!

System	Time
BlueGene/P	3.61 ns
Jaguar (XT-4)	0.78 ns
Odin (2 GHz x86-64)	1.15 ns
Laptop (1 GHz x86)	0.89 ns

Table: Costs per Argument



How to be as fast as SHMEM/ARMCI?

□ The Put Calls:

- `shmem_put64(tgt, addr, size, dst)`
- `ARMCI_Put(src_addr, dst_addr, count, tgt)` **but**
- `ARMCI_PutS(src_addr, src_stride[], dst_addr, dst_stride[], count, stride_levels, tgt)`
- `MPI_Put(src_addr, src_cnt, src_type, tgt_rank, tgt_disp, tgt_cnt, tgt_type, win)`

□ MPI “Overheads” (what makes MPI MPI)

- datatypes (arbitrary patterns, heterogeneity)
- `tgt_displs` (heterogeneity)
- `win` (collective **or** local window)

□ If the Forum endorses a fast-path (advice to impl/users)

- e.g., if `(src_type==tgt_type==MPI_BYTE)` then bypass datatype processing (heterogeneity and access patterns)
- overheads: window lookup (might mean cache miss); `tgt_displs` might be expensive (could mean cache miss or remote action); make all `displs==0` a special case!
- if no cache miss: only two branches more!



Synchronization and Consistency?

□ ARMCI:

- remote fence blocks until all ops completed at dest
- collective version: AllFence or Barrier

□ GASNet:

- AM synchronization by user (poll on local variable)
- RMA functions block until memory is consistent
 - Nonblocking versions available (with and without explicit handles)!

□ SHMEM:

- might require shmem_udcflush() on target (Cray T90)
- synch with shmem_wait(), shmem_barrier() or polling
 - not 100% defined (publicly) as it seems



Synchronization Models in MPI

- Collective windows
 - `MPI_Win_fence()` works well
 - `post/start/complete/wait` is nice but hard to use
 - passive target mode has consistency problems
- Local windows:
 - `MPI_Win_fence()` would need to be called by both processes (doesn't seem useful)
 - `post/start/complete/wait` would be funny
 - passive target mode could be useful on systems with strong memory consistency
 - how do we know the memory consistency?



Querying memory consistency

- `MPI_RMA_query(optype, model)`
 - `optype` – selects put, get, acc, or get_acc
 - `model`:
 - `MPI_RMA_SEPARATE` – not cache coherent/weak semantics
 - `MPI_RMA_ONE` – cache coherent/strict semantics
- Memory consistency explained:
 - `SEPARATE`: behaves like MPI-2, needs target needs to call MPI to make progress
 - `ONE`: behaves like InfiniBand, ARMCI, SHMEM, target memory is updated asynchronously (eventually)



What's different with strong consistency?

- Active target mode:
 - fence: pretty much nothing (the user can assume better overlap behavior ☺)
 - post/start/complete/wait: post-start and complete-wait dependencies persists
- Passive target mode:
 - becomes similar to ARMCI/SHMEM
 - user might choose to ignore lock/unlock
 - access epochs vanish
 - synchronization via polling (ugs)
 - we might want to add MPI_Target_wait() (?)



Optimizing the Interface (possibilities!)

- ❑ RDMA hardware (OFED, DCMF) and SM:
 - MPI_RMA_xxx() will just return success ☺
 - Win_fence() is simply an MPI_Barrier()
 - Ops might need extra message (see below)
- ❑ HW-supported AM (LAPI, DCMF):
 - send memory+op+local flag to target
 - Win_fence() could use termination detection algs.
 - AM handler on target triggers message to set local flag
 - or counter mechanism (cf. LAPI)
- ❑ Message Passing (Send/Recv):
 - emulate active messaging (send reply message after op is finished)



Questions?

