



# MPI-3 Backwards Compatibility

“All the fun of ABI, but with half the calories”



# Suffixes

- ▶ Use suffixes for new MPI functions:
  - Tickets current use: MPI\_File\_write\_l (the letter L)
  - ...but that suffix can be changed easily
- ▶ **QUESTION:** Is “\_l” a good suffix?
  - Other possibilities:
    - MPI\_File\_write\_L
    - MPI\_File\_writel ← this seems problematic
    - MPI\_File\_writeL
    - ...?




# Ticket #117

- ▶ <https://svn.mpi-forum.org/trac/mpi-forum-web/ticket/117>
- ▶ Defines an MPI\_Count datatype
  - Text is currently relative to MPI-3 trunk HEAD
- ▶ Added:
  - MPI\_Count items can be used in reductions (like MPI\_Offset)
    - MPI\_COUNT datatype (and MPI\_COUNT\_KIND)
  - MPI\_GET\_ELEMENTS\_L
  - MPI\_GET\_COUNT\_L
  - MPI\_STATUS\_SET\_ELEMENTS\_L



# Ticket #117


- ▶ Not deprecating anything
  - Can't deprecate MPI\_GET\_ELEMENTS, etc.
- ▶ Not implemented yet; polling Forum for intent first



# Ticket #117


## ► QUESTIONS:

- Ticket changes orientation of text to discuss the “\_l” versions of the functions as the primary. Ok?
- Do we need MPI\_UNDEFINED\_L?
  - For MPI\_GET\_COUNT\_L
  - We *think* so...
- What happens if you send 4B elements and call MPI\_GET\_ELEMENTS?
  - I.e., what is the error case?



# Ticket #220


- ▶ <https://svn.mpi-forum.org/trac/mpi-forum-web/ticket/220>
- ▶ Create MPI-IO functions that use MPI\_Count
  - Add MPI\_File\_<foo>\_l functions
  - Also add MPI\_CONVERSION\_FN\_NULL\_L (etc.)
    - But not MPI\_DATAREP\_EXTENT\_FUNCTION (!)
  - Updated examples to use MPI\_Count
- ▶ Not implemented yet – polling Forum for intent first



# Ticket #220


## ► QUESTIONS:

- Ticket changes orientation of text to discuss the “\_l” versions of the functions as the primary. Ok?
- Deprecate the non-\_l functions?
  - No text has been added yet about deprecation
  - Will need to list them all in the deprecated chapter if we actually deprecate them
- MPI\_SET\_STATUS\_ELEMENTS\_L: needed if we assume file operations will use generalized requests



# More Questions


- ▶ Assuming that we do nothing for C++ bindings for MPI\_Count (i.e., no MPI::Count)
- ▶ What to do with all new MPI-3 functions?
  - Do they all use MPI\_Count?





# More Questions

- ▶ What about the blocking collectives?
  1. MPI\_Bcast, MPI\_Ibcast\_L(...MPI\_Count...)
    - The lack of symmetry seems painful
  2. MPI\_Bcast, MPI\_Ibcast(...int...)
    - Not using MPI\_Count for a new routine seems like a step backward
  3. MPI\_Bcast\_L(...MPI\_Count...), MPI\_Ibcast\_L(...MPI\_Count...)
    - If we introduce \_L versions of all the collectives, then why not just do it for all routines...?
  4. MPI\_Bcast, MPI\_Ibcast(...MPI\_Count...)
    - Only using \_L extensions when there is a name conflict makes it unclear when a routine takes int vs. MPI\_Count.
- All new MPI-3 functions with a count argument are \_L
- Have all 4 variants: Bcast, Bcast\_L, Ibcast, Ibcast\_L



- ▶ 1. bcast and ibcast both with int
  - ibcast with integer for fortran safety and storage scalability (e.g., gatherv)
- ▶ 2. new ticket for bcast\_l and ibcast\_l
- ▶ 3. new ticket(s) for other things with \_l
- ▶ Point made:
  - Really only need 64 bit queries: get\_elements\_l and get\_count\_l
  - So maybe we don't need MPI\_Count versions of MPI\_File\_<foo>\_l ...etc....?
  - It could be useful to have Datatype constructors that take MPI\_Count



# More Questions

- ▶ If we add \_L versions of the collectives...
  - What about MPI\_Accumulate?
- ▶ If we do MPI\_ACCUMULATE\_L...
  - What about MPI\_GET and MPI\_PUT?
- ▶ Also note: there's ~15 collectives
  - There's ~15 flavors of send, too
  - So why not do the pt2pt functions?
- ▶ **WARNING: SLIPPERY SLOPE**

