

## MPI-3 RMA Draft



## RMA Goals

---

- Extend the applicability of MPI RMA to additional (but not necessarily *all*) applications
- Address documented shortcomings in using MPI RMA
- Maintain look and feel of MPI; maintain relevance to future platforms



## New Features at a Glance

- Additional ways to create MPI\_Win
  - ♦ Allocate memory for window (more scalable)
  - ♦ Dynamically attach/detach memory to an existing window
- Additional remote operations
  - ♦ Including read-modify-write, absent in MPI-2
- Extensions to passive target RMA
  - ♦ Win\_lock\_all locks all processes in MPI\_Win
  - ♦ Local and remote completion
- RMA operations with requests



3

## New Features Continued

- Many operations previously erroneous now undefined
  - ♦ Enables different programming models
- Ordering:
  - ♦ Any combination of read/write ordering
  - ♦ rr, rw, wr, ww as info key value string
- Extended info keys
  - ♦ Relax ordering (see above)
  - ♦ Info query mechanism suggested (for misc WG, will be discussed later)



4

## New Routines at a Glance

- MPI\_Win\_allocate
- MPI\_Win\_create\_dynamic
  - ♦ MPI\_Win\_register
  - ♦ MPI\_Win\_deregister
- MPI\_Get\_accumulate
  - ♦ Generic, similar to MPI\_Accumulate
- MPI\_Fetch\_and\_op
  - ♦ Single datatype, single item specialized (fast)
- MPI\_Compare\_and\_swap



5

## New Routines Continued

- MPI\_Win\_lock\_all
  - ♦ MPI\_Win\_unlock\_all
- MPI\_Win\_flush
  - ♦ MPI\_Win\_flush\_all
- MPI\_Win\_flush\_local
  - ♦ MPI\_Win\_flush\_local\_all
- MPI\_Win\_sync
- MPI\_Win\_query



6

## New Routines Continued

- MPI\_Rget, MPI\_Rget\_acumulate, MPI\_Rput, MPI\_Raccumulate



7

## Win Allocate

- Collective Window Allocation
  - ◆ Enables symmetric allocation for simple offset calculation
- MPI\_WIN\_ALLOCATE(size, disp\_unit, info, comm, base, win)
  - ◆ int MPI\_Win\_allocate(MPI\_Aint size, int disp\_unit, MPI\_Info info, MPI\_Comm comm, void \*base, MPI\_Win \*win)



8

## Win Create Dynamic

- Creates a window that allows local registration
  - ◆ Memory needs to be registered before being accessed remotely
- `int MPI_Win_create_dynamic(MPI_Info info, MPI_Comm comm, MPI_Win *win)`
  - ◆ `int MPI_Win_register(MPI_Win win, void *base, MPI_Aint size)`
  - ◆ `int MPI_Win_deregister(MPI_Win win, void *base)`
  - ◆ Overlapping registrations are erroneous!



9

## Get Accumulate

- Similar to fetch&add
- `int MPI_Get_accumulate(void *origin_addr, int *origin_count, MPI_Datatype origin_datatype, void *result_addr, int *result_count, MPI_Datatype result_datatype, int target_rank, MPI_Aint target_disp, int *target_count, MPI_Datatype target_datatype, MPI_Op op, MPI_Win win)`
  - ◆ Supports `MPI_NO_OP`



10

## Fetch and Op

- Specialized Get Accumulate for performance
- MPI\_Fetch\_and\_op(void \*origin\_addr, void \*result\_addr, MPI\_Datatype datatype, int target\_rank, MPI\_Aint target\_disp, MPI\_Op op, MPI\_Win win)
  - ◆ No count (count = 1)
  - ◆ Only one datatype



11

## Compare and Swap

- int MPI\_Compare\_and\_swap(void \*origin\_addr, void \*compare\_addr, void \*result\_addr, MPI\_Datatype datatype, int target\_rank, MPI\_Aint target\_disp, MPI\_Win win)
  - ◆ Datatypes are limited to some predefined datatypes
  - ◆ C integer, Fortran integer, Logical, Complex, Byte, MPI\_AINT, MPI\_OFFSET



12

## Win Lock All

- Locks all target processes in a window
  - ♦ Obviously a shared lock ☺
  - ♦ Optimization for a simple loop
- `int MPI_Win_lock_all(int assert, MPI_Win win)`
  - ♦ `int MPI_Win_unlock_all(MPI_Win win)`
- Did not include `MPI_Win_lock_group`
  - ♦ No use-case, can implement as loop over `Win_lock`



13

## Win Flush

- Blocks until all operations completed remotely
- `int MPI_Win_flush(int rank, MPI_Win win)`
  - ♦ Flush a specific rank
- `int MPI_Win_flush_all(MPI_Win win)`
  - ♦ Flush all ranks
- Applies only to passive target RMA



14

## Win Flush Local

- Blocks until all operations completed locally
  - ♦ Local buffers can be re-used
  - ♦ Might imply remote completion (e.g., Get)
- `int MPI_Win_flush_local(int rank, MPI_Win win)`
  - ♦ Operations targeted to a specific rank
- `int MPI_Win_flush_local_all(MPI_Win win)`
  - ♦ All previous operations on window win



15

## RMA Query

- Query remote memory consistency
- `int MPI_RMA_query(MPI_Win win, int *model)`
  - ♦ Returns either `MPI_RMA_SEPARATE` (MPI-2) or `MPI_RMA_UNIFIED` (public window = private window)
  - ♦ `MPI_RMA_UNIFIED` changes semantic rules 5+6



16



## RMA With Request

- Allows process to wait for local completion
  - ◆ Permits reuse of buffers on Rput, Raccumulate
  - ◆ Permits use of data on Rget, Rget\_accumulate
  - ◆ Uses an MPI\_Request (can use MPI\_Wait/Test)



20

## Bonachea's and Duell's criticism

- Window creation is collective
  - ◆ hinders efficient exposure for local objects
  - ◆ no "sparse" communication
- MPI\_Win\_create\_dynamic
  - ◆ Well suited for automatic memory management
  - ◆ Good as compilation target



21

## Bonachea's and Duell's criticism

- Exposed memory must be MPI\_Alloc\_mem()'d
  - ◆ no exposure of static memory or stack-variables
  - ◆ alloc\_mem might be limited by the implementation
- Also addressed in MPI\_Win\_create\_dynamic
  - ◆ Can register stack
  - ◆ Dangerous though! Undefined results if stack is out of scope.



22

## Bonachea's and Duell's criticism

- Forbids conflicting get/put (or local load/store) accesses to same memory
  - ◆ really hard to track for compilers (halting problem?)
  - ◆ Easy source of bugs in user codes
- Outcome is undefined (does not change the behavior of *correct* MPI-2 programs)



23

## Bonachea's and Duell's criticism

- Window's memory may not be updated by remote gets and local stores concurrently
  - ◆ simplifies MPI implementation significantly
  - ◆ seems very artificial and suboptimal from user's perspective
- Outcome is undefined
  - ◆ MPI\_Win\_sync provides way to make outcome defined and as user expects
  - ◆ MPI\_Win\_query with result MPI\_RMA\_UNIFIED also makes outcome defined and as user expects



24

## Bonachea's and Duell's criticism

- Overlapping memory regions of multiple windows can be created but not be used
  - ◆ "concurrent communications may lead to erroneous results"
- Now also "undefined"
  - ◆ [is this much better than erroneous?]
  - ◆ But these applications are likely to use a single dynamic window, so problem avoided



25

## Bonachea's and Duell's criticism

- Passive target RMA ops only lock a single process during an epoch
  - ♦ ops from one source to different targets are serialized
  - ♦ one window for each target to enable concurrent access?
    - scalability limitation
- MPI\_Win\_lockall, may use multiple MPI\_Win\_lock



26

## Comparison to ARMCI

- Similar to MPI\_RMA\_UNIFIED (if available)
  - ♦ ARMCI does not support non-CC systems without an additional activity (thread/process/callback)
- Local completion is different
  - ♦ Either blocking, implicit handles or Test/Wait(all)
  - ♦ Similar to "implicit handles"
- Remote completion similar
  - ♦ (all)Fence == Flush(\_all)
- Ordering with collectives
  - ♦ ARMCI\_Barrier combines barrier + allfence



27

## Comparison to ARMCI cont.

- ARMCI\_Malloc for exposed memory
  - ◆ Collective and local; no register (in an undocumented version)
- ARMCI\_RMW
  - ◆ Similar to Get\_accumulate
  - ◆ No compare-and-swap
- ARMCI\_Lock/Unlock
  - ◆ Lock special synch. objects
  - ◆ MPI RMA still doesn't include user-defined locks
- Limited set of collectives
  - ◆ Supports MPI collectives



28

## Co-Array Fortran

- RMA through co-arrays
  - ◆ Explicit collective allocation
- Execution divided in segments (delimited by "image control statements")
  - ◆ (non-volatile) operations in a segment are unordered
  - ◆ Exceptions are values with "atom" or volatile argument
- Image control statements:
  - ◆ Sync all, sync images, sync memory, lock/unlock
- Can (almost?) be implemented with MPI-3 RMA



29

## Unified Parallel C

- Explicit collective allocation
- Relaxed and strict access
  - ♦ Relaxed accesses are unordered
    - Exception are conflicting accesses to the same memory which appear in program order
  - ♦ Strict accesses are ordered (program order)
- Sync operations:
  - ♦ upc\_fence, upc\_(notify,wait), upc\_barrier, upc\_(un)lock
  - ♦ flush, (nonblocking) barrier, barrier, <no user locks>



30

## Comparison Summary

- MPI-3 RMA covers most features from existing RMA systems
  - ♦ No user-defined locks
    - Can be emulated with compare and swap
- MPI-3 RMA can efficiently simulate many other RMA systems
  - ♦ Note more RMA programming models do not include an active message interface
  - ♦ Likely due to the complexity in defining the detailed semantics of AM for users
  - ♦ Some RMA features in other systems may require emulation, e.g., with MPI pt-2-pt



31

## Executive Summary

---

- Fixes known deficits of MPI-2 RMA and can be implemented on non-CC architectures
  - Part of the complexity is pushed to the user
- ◆ Enables the use of MPI-RMA for new important application domains
- ◆ Does not handle everything but that was not a goal (e.g., no active messages)

