

MPI-3 Hybrid Working Group Status

MPI interoperability with Shared Memory

- Motivation: sharing data between processes on a node without using threads
- Sandia applications motivation:
 - Dump data into a “common” memory region
 - Synchronize
 - All processes just use this memory region in read-only mode
 - And they want this in a portable manner

MPI interoperability with Shared Memory: Plan A

- The original plan was to provide
 - Routines to allocate/deallocate shared memory
 - Routines to synchronize operations, but not to operate on memory (similar to MPI_Win_sync)
 - Operations on shared memory are done using load/store operations (unlike RMA)
 - Synchronization would be done with something similar to MPI_Win_sync
 - There was a suggestion to provide operations to operate on the data as well, but there was no consensus on that in the working group
 - A routine to create a communicator on which shared memory can be created
- The Forum's feedback that this was not possible to do in MPI unless it knows the compilers capabilities

Process 0	Process 1
store(X)	
MPI_Shmem_sync()	
MPI_Barrier()	MPI_Barrier()
	MPI_Shmem_sync()
	load(X)

MPI interoperability with Shared Memory: Plan B

- The second plan was to remove shared memory synchronization routines; we still have:
 - Routines to allocate/deallocate shared memory
 - A routine to create a communicator on which shared memory can be created
- The Forum's feedback was that allocation/deallocation might not be useful without operations to synchronize data
 - The use-case for only creating shared memory and expect the user to handle memory barriers was fairly minimal
 - Also, another feedback was to do this as an external library

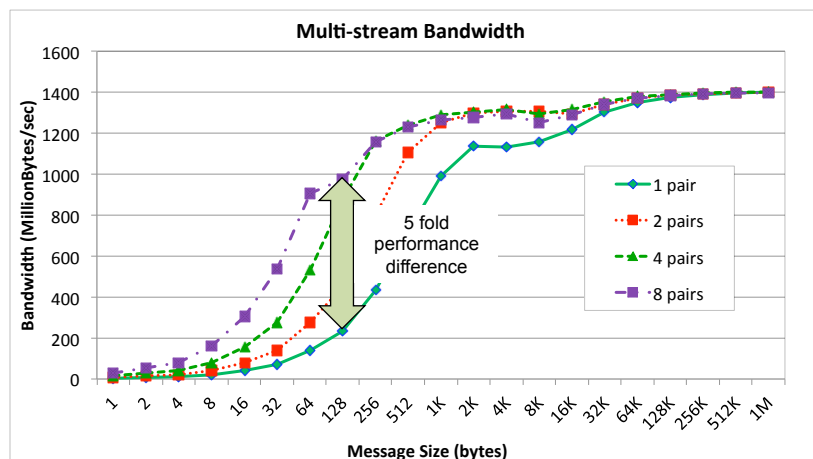
MPI interoperability with Shared Memory: Plan C

- The third plan is to remove shared memory allocation routines; we still have:
 - A routine to create a communicator on which shared memory can be created
- `MPI_Shmem_comm_create(old_comm, info, &new_comm)`
 - The info argument can provide implementation-specific information such as, within the socket, shared by a subset of processes, whatever else
 - No predefined info keys
 - There has been a suggestion to generalize this to provide a communicator creator function that provides “locality” information
 - Will create an array of communicators, where the lower index communicators “might contain” closer processes (best effort from the MPI implementation)
 - An attribute on the communicator would tell if shared memory can be created on it


MPI Interoperability with Thread Teams

- Motivation: Allow coordination with the application to get access to threads, instead of the application maintaining a separate pool of threads than the MPI implementation
- Proposal in a nutshell
 - Thread teams are created
 - The user can provide the threads in the team to the MPI implementation to help out the MPI implementation
 - A user-defined team allows the user to have control on locality and blocking semantics for threads (i.e., which threads block waiting to help)

Point-to-point Communication on InfiniBand



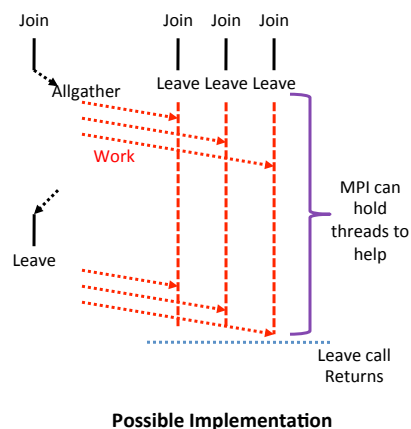
Allreduce on BG/P



Thread Teams Models


- Two models proposed
 - Synchronous model: all threads in the team synchronously join and leave the team – more restrictive for the application but has more optimization opportunity for MPI
 - Asynchronous model: threads join the team asynchronously; processes can leave synchronously or asynchronously break out
- Both models are essentially performance hints
 - Optimized MPI implementations can use this information

Thread Teams: Synchronous Model



- In the synchronous model, the Join and Leave calls can be synchronizing between the threads
 - The MPI implementation can assume that all threads will be available to help
 - The programmer should keep the threads synchronized for good performance
- The MPI implementation knows how many threads are going to help, so it can statically partition the available work, for example
- This model does not allow for threads to “escape” without doing a synchronous leave

Thread Teams: Asynchronous Model



- In the asynchronous model, the Join call is not synchronizing between the threads
- The leave call is still synchronizing, but threads are allowed to either help using “leave” or “break out” without helping
- The MPI implementation does not know how many threads are going to help
- This model allows for threads to “break out” without doing a synchronous leave

Thread Teams Proposed API

- Team creation/freeing
 - `MPI_Team_create(team_size, info, &team)`
 - One predefined info key for “synchronous”; default is “asynchronous”
 - Similar to the MPI RMA chapter in that the info arguments are true assertions; if the user says “synchronous” and tries to break out, that’s an erroneous program
 - `MPI_Team_free(team)`
- Team join/leave functionality
 - `MPI_Team_join(team)`
 - A thread can only join one team at a time
 - `MPI_Team_leave(team)`
 - `MPI_Team_break(team)`

Hartree-Fock Example

```

do {
    One_Electron_Contrib(Density, Fock)
    while (task = next_task()) {
        {i, j, k} = task.dims
        X = Get(Density, {i,j,k} .. {i+C,j+C,k+C})
    #pragma omp parallel {
        Y = Work({i,j,k}, X)                ; <----- compute intensive
        omp_barrier()
        MPI_Team_join(team)
        if (omp_master) {
            Accumulate(SUM, Y, Fock, {i,j,k}, {i+C,j+C,k+C}); <----- communication intensive
        }; OMP master
        MPI_Team_leave(team)
    }; OMP - end parallel block
    }
    #pragma omp parallel {
        Update_Density(Density, Fock)        ; <----- communication intensive
        my_energy = omp_gather()
        MPI_Team_join(team)
        if (omp_master) {
            energy = MPI_Allgather(my_energy)    ; <----- moderately communication intensive
        }; OMP master
        MPI_Team_leave(team)
    }; OMP - end parallel block
    } while (abs(new_energy - energy) > tolerance)

```