

Keyval and callbacks

Rules and Behaviors

Background

Keyvals are used for caching attributes on {COMM,TYPE,WIN} objects

Discussion will focus on COMM functions for convenience

APIs:

- MPI_Comm_create_keyval
- MPI_Comm_free_keyval
- MPI_Comm_set_attr
- MPI_Comm_get_attr

Callbacks

- MPI_Comm_copy_attr_function
- MPI_Comm_delete_attr_function

Behavior

Create a key (global) that can subsequently be used to set/clear attributes on objects.

Keys are represented by their key value (keyval) (page 286, line 7)

- Keyval is set to MPI_KEYVAL_INVALID by MPI_Comm_free_keyval

1 - * relationship between keys and attributes

Attributes on objects reference the key, even if the key is freed (and its keyval set to MPI_KEYVAL_INVALID).

Callbacks

Copy callback is invoked when an object is duplicated

- E.g. `MPI_Comm_dup`

Delete callback is invoked when an object is freed

- E.g. `MPI_Comm_free`

Callback takes 'comm' and 'keyval' as input parameter

- 'oldcomm' in copy, but for simplicity following will use 'comm' for both

Delete callback harder to handle than copy callback

- Following slides focus (mostly) on delete callbacks

Ordering of delete callbacks

MPI_Comm_free indicates delete callback called in arbitrary order

- Page 248, Line 27

MPI_Finalize indicates that attributes on MPI_COMM_SELF are deleted in reverse order from creation, despite behaving as-if MPI_Comm_free was called for MPI_COMM_SELF.

- Page 363, Line 17

Recommend:

- Make standard consistent
- Define ordering as reverse from creation
 - Backward compatible

Are Keyvals Handles?

They behave like handles

- MPI_Comm_free_keyval sets the keyval to MPI_KEYVAL_INVALID
- Similar to MPI_Comm_free sets comm to MPI_COMM_NULL

They don't have an associated handle type

- Standard defines it as 'integer' type
- MPI_KEYVAL_INVALID vs. MPI_COMM_NULL

Can copy of keyval be used after MPI_Comm_free_keyval?

- Generally erroneous to use a handle after it has been freed
 - Exception is MPI_Grequest_complete, which may be called (using a copy of the original handle value) after MPI_Request_free.
 - Didn't find any other exception
 - Standard doesn't clearly state anywhere that using stale handles is 'erroneous'
 - Duh!

The Real Problem

Standard does not specify what can (or cannot) be done in a callback

- i.e. nothing is forbidden, thus everything is valid
- Must tread lightly here, as we could break back-compat

Standard does not specify if the 'keyval' param can be temporary

- i.e. for the duration/scope of the callback only
- Again, potential back-compat issues

Standard does not specify if 'keyval' can be reused before 'key' is freed.

- Not forbidden, thus valid (despite introducing issues)

Scenario

`MPI_Comm_create_keyval(..., &keyval, ...)`

- `Keyval = 3`

`MPI_Comm_set_attr(MPI_COMM_WORLD, keyval, attr_val)`

`MPI_Comm_free_keyval(&keyval)`

- `Keyval = MPI_KEYVAL_INVALID`

`MPI_Comm_dup(MPI_COMM_WORLD, &dup)`

- Copy callback is called – what is keyval param?
 - `MPI_KEYVAL_INVALID`?
 - 3?
 - Something else?

`MPI_Comm_free(&dup)`

- Delete callback is called – what is keyval param?
 - `MPI_KEYVAL_INVALID`?
 - 3?
 - Something else?

Calls from (delete) callbacks

Function	MPI_KEYVAL_INVALID	Original keyval	Something Else
blocking comms/IO	Y	Y	Y
non-blocking comms/IO	GTFO?	GTFO?	GTFO?
MPI_Comm_free_keyval	N	? <already freed!!!>	Y
MPI_Comm_get_attr	N	? <already freed!!!>	Y
MPI_Comm_set_attr	N	? <already freed!!!>	Y
MPI_Comm_delete_attr	N	Diff comm or diff keyval?	Y (except same comm??)

Communication or I/O

Don't need the keyval, so independent of whether keyval was freed

Blocking probably OK, a bit strange though

Non-blocking potentially problematic unless waiting in the callback

- Nasty if initiated from an attribute delete callback on MPI_COMM_SELF during MPI_Finalize

MPI_Comm_free_keyval

OK for different comm and/or keyval (retrieved via extra_state or global?)

- E.g. Free keyval B from a callback for key A

Can get messy if keyval was already freed

- Double free?
- Implementation must have a way to track whether keyval was freed
 - Proper reference counting on key when callback unwinds
- Mitigated by using “something else” than original keyval
 - keyval encodes that MPI_Comm_free_keyval is noop somehow

MPI_Comm_get_attr

OK for different comm and/or keyval

Unnecessary for same comm and keyval

- Attribute value is passed as input parameter
- Probably harmless

MPI_Comm_set_attr

OK for copy callback

Likely always wrong in delete callback

- Throw in MPI_THREAD_MULTIPLE for extra fun!

Definitely wrong in delete callback for same comm

- Seriously messes with callback ordering requirements
- communicator is being freed, what's the point?

MPI_Comm_delete_attr

Probably OK for different comm and/or keyval

Just plain wrong from a delete callback for same comm and keyval

Unwind Semantics

If delete_fn fails, the associated MPI*_FREE call fails

Potentially nasty implications to ordering requirement for attribute callbacks if attributes can be set from delete callback

- If delete_fn succeeds, ordering would require newly added attributes to be deleted next
- If delete_fn fails, ordering would require newly added attributes to precede the failed attribute

Summary

It's like picking a scab...

It's pretty horrid with many corner cases

- Even if the keyval wasn't freed
- Worse yet if the keyval was freed

Fixing it can break existing apps

Solution should minimize potential for back-compat

- Except egregious behavior

Recommendation

Option 1 (Cleanest):

- MPI_Comm_free_keyval sets keyval to MPI_KEYVAL_INVALID, which will also be the value that is passed to the callbacks.
- Likely back-compat issues for lazy apps (e.g. PETc)

Option 2 (Best back-compat):

- Document that keyval parameter to callbacks may not be the same as the keyval created by MPI_Comm_create_keyval if that keyval was freed with MPI_comm_free_keyval.
- Enable encoding 'behavior' in keyval, e.g. MPI_Comm_free_keyval is noop
- Applications can use 'extra_state' for context if needed
- Makes behavior if original keyval was freed rational
- Create either white list or black list of acceptable calls during callbacks
 - Laborious
 - Black list likely more reasonable given back-compat

All I/O (comm or file) initiated during callback must be completed in callback

- Allows non-blocking, but with restrictions

Suggested Black List

Only needed for delete callbacks?

Function	Same Comm	Same Keyval	Both Same	Neither Same
MPI_Comm_set_attr	No	No	No	No
MPI_Comm_delete_attr	Yes	Yes	No	Yes
MPI_Comm_get_attr	Yes	Yes	Yes (but silly)	Yes