

# Endpoints Plenary

James Dinan

Hybrid Working Group

December 10, 2013

# Status of Endpoints

1. Proposal text #380 ready
2. Explored interoperability story [EuroMPI '13]
3. Exploring performance story [IJHPCA in prep]
4. Working on implementation in MPICH
  - Will be open source
5. Target: Formal reading in March

# Motivation for Endpoints


## 1. Interoperability argument

- On-node programming model
- Multi-node models that use threads

## 2. Performance argument


- Increase communication concurrency
  - Preserve shared memory/node-level programming
  - Make number of VA spaces free parameter
- Reduce synchronization penalties
  - Privatize thread communication state and resources

# Achievable Network Performance (Dramatization)



- Network endpoint design evolving to support many cores
- ***Not real data, represents my personal views***
- Gathering real data for paper, will present at next meeting

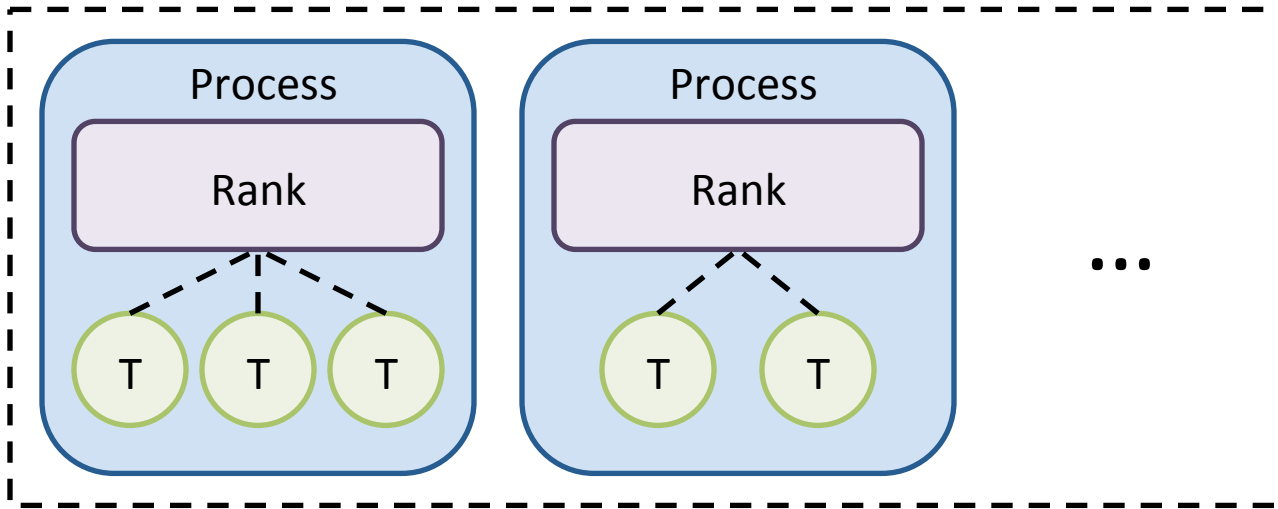
# Impact of Queue Depth on Message Rate



- Brian Barrett, et al. [EuroMPI '13]
- Threads sharing a rank increase posted receive queue depth (x-axis)

# Mapping of Ranks to Processes

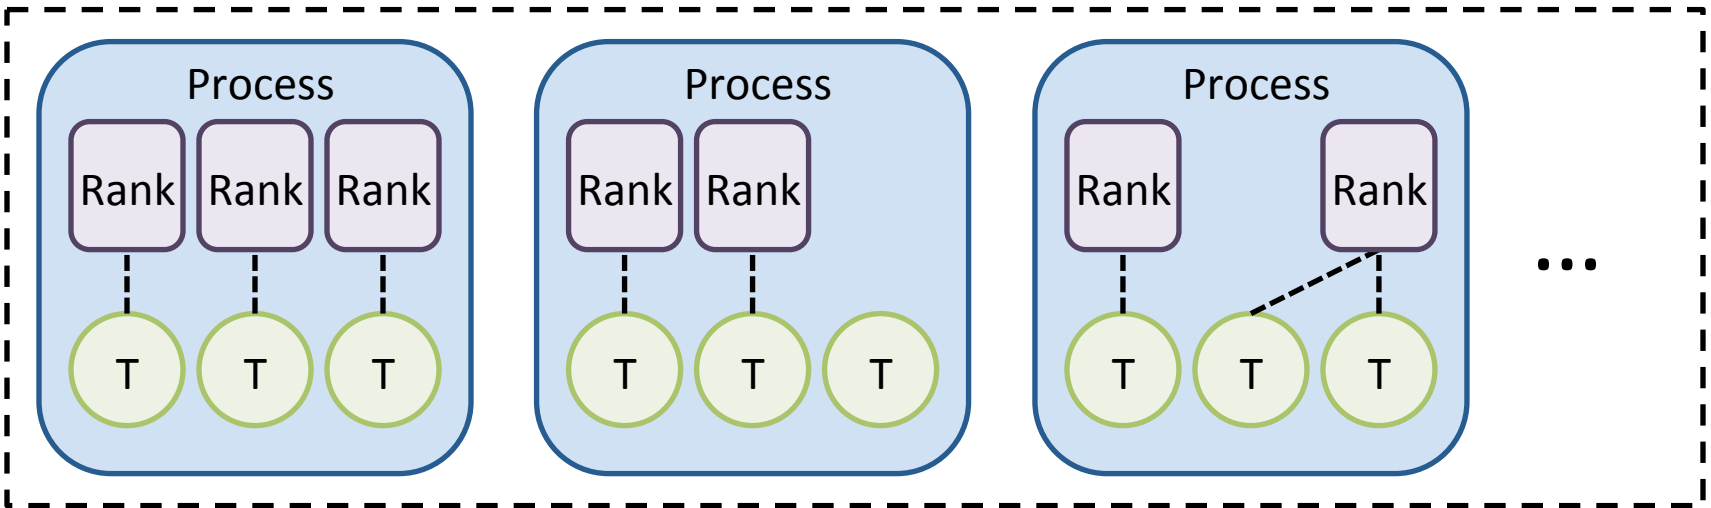
Conventional Communicator



- MPI provides a 1-to-1 mapping of ranks to processes
- This was good in the past
- Usage models and systems have evolved
  - Hybrid MPI+Threads programming
  - Ratio of core to network endpoint performance decreasing

# Endpoints Model

## Endpoints Communicator



- Many-to-one mapping of ranks to processes
  - Threads act as first-class participants in MPI operations
  - Improve programmability of MPI + X
  - Threads drive independent network endpoints
- Endpoint: Set of resources that supports the independent execution of MPI communications
  - Endpoints have process semantics


# Current THREAD\_MULTIPLE Usage

- MPI message matching space: <communicator, sender, tag>
- Two approaches to using THREAD\_MULTIPLE
  1. Match specific thread using the tag:
    - Partition the tag space to address individual threads
    - Limitations:
      - Collectives – Multiple threads at a process can't participate concurrently
      - Wildcards – Multiple threads concurrently requires care
  2. Match specific thread using the communicator:
    - Split threads across different communicators (e.g. Dup and assign)
    - Can use wildcards and collectives
    - However, limits connectivity of threads with each other

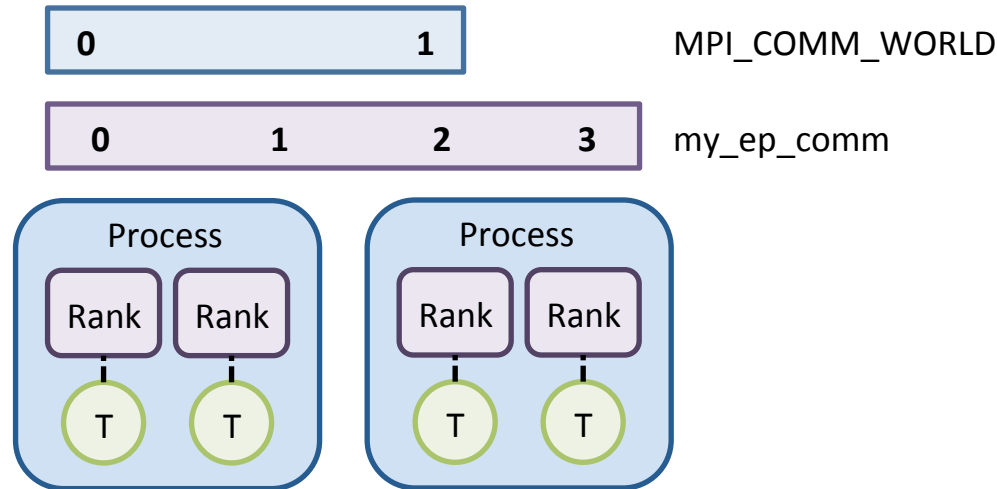


# Implementation of Endpoints

- Two implementation strategies
  1. Each rank is a network endpoint
  2. Ranks are multiplexed on endpoints
    - Effectively adds destination rank to matching
  3. Combination of the above
- Potential to reduce threading overheads
  - Separate resources per thread
    - Rank can represent distinct network resources
    - Increase HFI/NIC concurrency
  - Separate software state per thread
    - Per-endpoint message queues/matching
    - Enable per-communicator threading levels
- FG-MPI implements “static” endpoints
  - A little different, still demonstrates implementation and performance benefits



# Endpoints Interface



```
int MPI_Comm_create_endpoints(  
    MPI_Comm parent_comm, int my_num_ep,  
    MPI_Info info, MPI_Comm *out_comm_hdls[])
```

- Out handle array takes TLS out of the implementation and off the critical path
- Each process requests an independent number of endpoints
- `MPI_ERR_ENDPOINTS` – Endpoints could not be created

# Endpoints Proposal

<https://svn.mpi-forum.org/trac/mpi-forum-web/ticket/380>