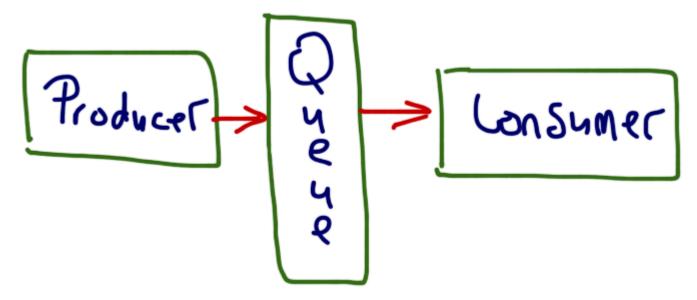


```
while(true) {
  produce()
  MPI_Send(prod, 1, INT, 1, 99, comm)
}
```

```
while(true) {
   MPI_Recv(prod, 1, INT, 1, 99, comm, stat)
   consume()
}
```



Active synchronization: fence

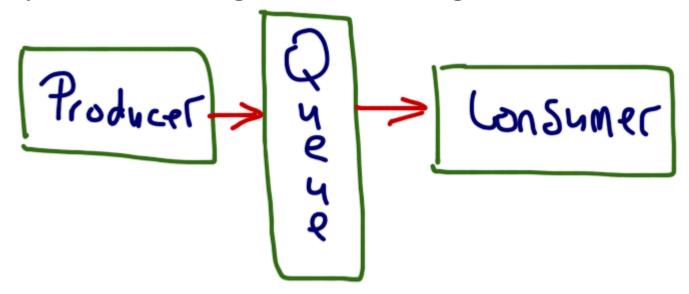


```
while(true) {
  produce()
  MPI_Put(prod, 1, INT, 1, 0, 1, INT, win)
  MPI_Win_fence(win);
}
```

```
while(true) {
   MPI_Win_fence(win)
   consume()
}
```



Active synchronization: general active target

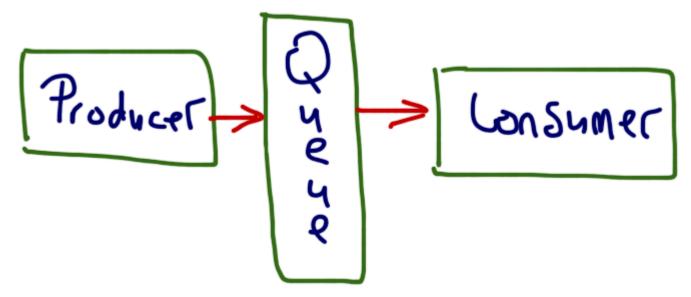


```
while(true) {
  produce()
  MPI_Win_start(grp, 0, win)
  MPI_Put(prod, 1, INT, 1, 0, 1, INT, win)
  MPI_Win_complete(win)
}
```

```
while(true) {
  MPI_Win_post(grp, 0, win) consume()
  MPI_Win_wait(win)
  consume()
}
```



Passive synchronization: lock/unlock with barrier

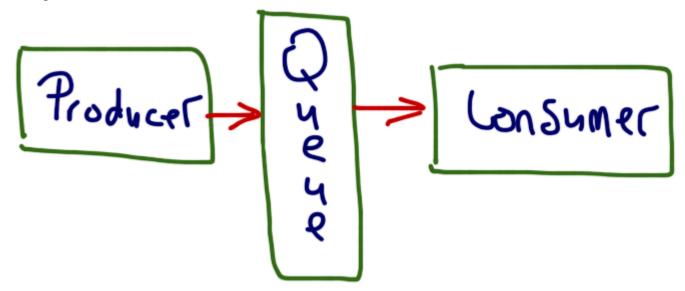


```
while(true) {
  produce()
  MPI_Win_lock(excl, 0, 0, win)
  MPI_Put(prod, 1, INT, 1, 0, 1, INT, win)
  MPI_Win_unlock(0,win)
  MPI_Barrier()
}
```

```
while(true) {
   MPI_Barrier()
   MPI_Win_lock(excl, 0, 0, win)
   consume()
   MPI_Win_unlock(0,win)
}
```



Passive synchronization: lock/unlock with send/recv

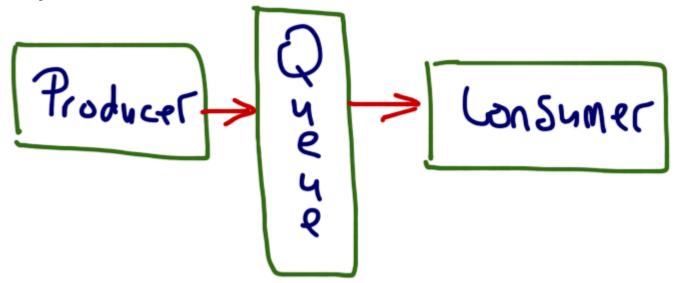


```
while(true) {
  produce()
  MPI_Win_lock(excl, 0, 0, win)
  MPI_Put(prod, 1, INT, 1, 0, 1, INT, win)
  MPI_Win_unlock(0,win)
  MPI_Send(flag, 1, INT, 1, 99, comm)
}
```

```
while(true) {
   MPI_Recv(flag, 1, INT, 1, 99, comm, stat)
   MPI_Win_lock(excl, 0, 0, win)
   consume()
   MPI_Win_unlock(0,win)}
```



Passive synchronization: lock_all with send/recv



```
MPI_Win_lock_all(0, win)
while(true) {
  produce()
  MPI_Put(prod, 1, INT, 1, 0, 1, INT, win)
  MPI_Win_flush(0, win)
  MPI_Send(flag, 1, INT, 1, 99, comm)
}
MPI_Win_unlock_all(win)
```

```
while(true) {
   MPI_Recv(flag, 1, INT, 1, 99, comm, stat)
   MPI_Win_lock(excl, 0, 0, win)
   consume()
   MPI_Win_unlock(0,win)
}
```



What do we want for a Prod/Cons Pattern

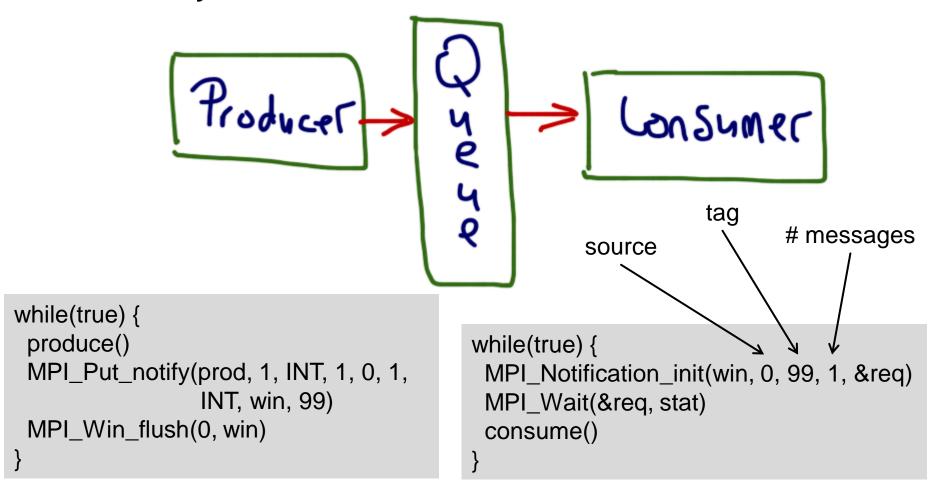
- Something in between active and passive synchronization
 - Not the full exposure epoch concept for active (too heavy)
 - Yet, have remote involvement (notification) at the destination
 - Destination is **not** needed to make progress

Half-passive synchronization: remote notifications

- Re-using MPI requests similar to persistent communication
 MPI_Start(MPI_Request *request)
- One-sided use-case often involves multiple transfers for a logical message We need a counter (well, less MPI-ish) What do we count??
- Counting messages!
 Could also count elements, but that would involve datatypes? Possible!
- Matching like MPI messages (src, tag)-tuple
 Re-using existing concepts



Passive synchronization: lock_all with send/recv



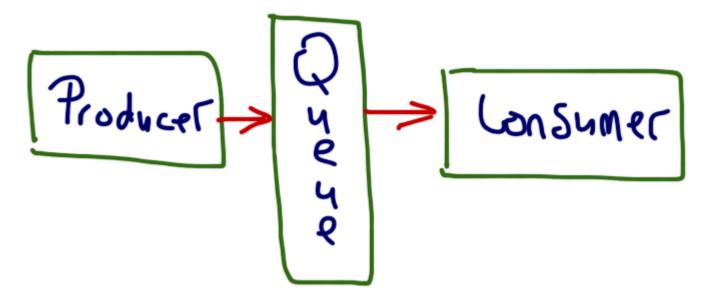


```
while(true) {
   MPI_Notification_init(win, 0, 99, 1, &req)
   MPI_Wait(&req, stat)
   consume()
}
```

- Somewhat odd with the epoch concept
- Option 1:
 - Put_notify() would close and open an access epoch
 - Would be a communication and synchronization operation
 - Has to flush all previous messages as well
- Option 1:
 - Put_notify() is outside the epoch concept
 - Somewhat ugly in the details and for tools



Alternative? Maybe hard to implement? Cf. #439



```
while(true) {
   MPI_Notification_init(win, 0, 99, 1, &req)
   MPI_Wait(&req, stat)
   consume()
}
```



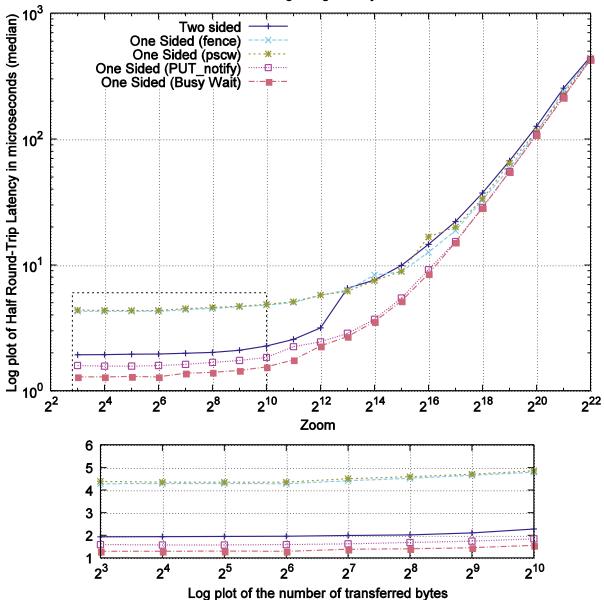
Alternative? Maybe hard to implement? Cf. #439

```
while(true) {
   MPI_Notification_init(win, 0, 99, 1, &req)
   MPI_Wait(&req, stat)
   consume()
}
```

- Slightly better!
- Yet, hard to implement on today's hardware
 - InfiniBand: rely on in-order and put with immediate
 - Cray DMAPP: impossible, uGNI: not 100% sure (maybe)
 - Shared memory: mfence + store

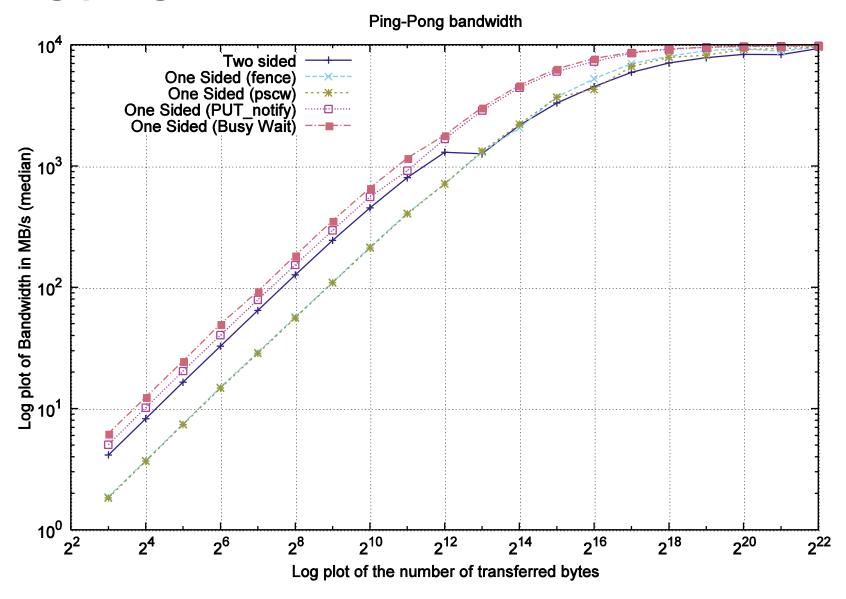


Ping-Pong latency



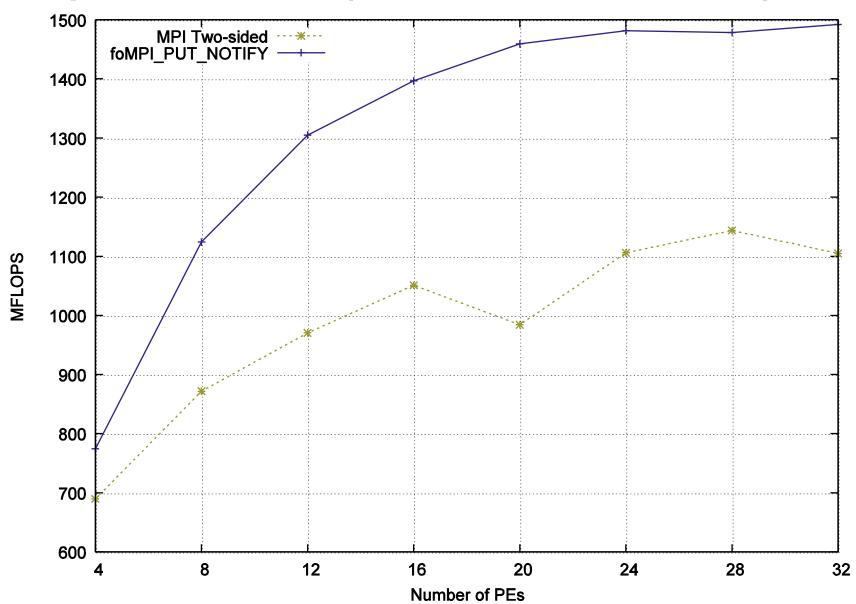


Ping-pong bandwidth





Pipelined Stencil (Intel Microbenchmarks)





Discussion

- Our foMPI [1] implementation will be released soon
 - on top of uGNI
 - using Op_notify
 - will be 100ns slower than polling (but safe without seq. consistency)
 - 100 ns still missing (fighting for it)
- Op_notify or Sync_notify?
 - Implementation tradeoffs (should they hold us back)?
- This is all single-producer-single-consumer
 - What about multiple-producer-multiple-consumer?
 - Needs advanced HW support to be fast (remote memory management)
- Do we want this at all?