

# Active Messages for MPI

- would make things easier! -

Torsten Hoefer and Jeremiah Willcock

Open Systems Lab  
Indiana University  
Bloomington, USA

10<sup>th</sup> MPI Forum Meeting June'09  
Menlo Park, CA, USA  
Jun 8–10th, 2009

# Motivation—What are Active Messages?

## Original Active Message Motivation

- T. von Eicken, et al.: *“Active Messages: a Mechanism for Integrated Communication and Computation”* (1992)
- Overlap communication and computation
- Asynchronous communication with minimal synch.
- Low overhead injection, pipelined transport
- Potentially run in interrupt handler → no buffering

## Active Messages have many Faces

- Similar to but not quite like remote procedure invocation
- Perform small non-blocking functions on message data
- Work is often inserted in a queue to be handled by main thread (cf. first-level interrupt handlers)
- Thread (and signal) safe?
- Handlers must halt for every input in each environment!

# Motivation—State of the Art?

## Many Middlewares Provide Active Messages

- GASNet
- IBM's LAPI and DCMF
- Myrinet's MX
- POOMA/CHEETAH
- Application layers like in PBGL ...

## Some Application Examples

- Parallel graph computations (inserting into remote queues)
- Implementing DSM systems (copying data to remote addresses without full (non-scalable or global) mapping)
- High-level language bindings (manipulating remote data structures)
- Also usable for (irregular) halo exchange (like one-sided)

# Can we implement the functionality using MPI-2.2?

## One-Sided?

- MPI\_Accumulate goes in the direction but is too limited
- cf. D. Bonachea: *“The Inadequacy of the MPI 2.0 One-sided Communication API for Implementing Parallel Global Address-Space Languages”*

## Two-Sided?

- The usual suspect—*asynchronous progress* (invocation)  
⇒ polling vs. thread
- Polling would be impractical (eliminate all AM advantages)
- Thread requires *THREAD\_MULTIPLE* to not limit options
- Performance penalty for threaded execution
- Is workable in practice (we have an implementation)  
⇒ performance penalty is clearly visible

# Should we have it in the Spec?

## Disadvantages (Problems, Architectures)

- A new (but simple) concept in the spec
- Some esoteric architectures could be problematic (learn from GASNet?)
- Same problems as ticket #26 (“Add a callback function if a request completes”)

## Advantages (Performance, Semantics)

- It is universal (one could implement one-sided with AM)
- Could serve well as compilation target (UPC, CAF, ...)
- It solves many high-level language binding issues
- All the original AM advantages (overlap, asynch., ...)

# Option #1

## The Handler Function

- `(*MPI_AM_Handler)(void* userdata, const void* recvbuf, MPI_Status* status)`
- `MPI_AM_Register(int id, MPI_AM_Handler handler, void* userdata, int maxcnt, MPI_Datatype type, MPI_Comm c)`
- `MPI_AM_Deregister(int id, MPI_Comm comm)`

## Sending an AM

- `MPI_AM_Send(const void* sendbuf, int count, MPI_Datatype type, int dest, int id, MPI_Comm comm)`
- ... `Isend`, `Bsend`, `Ssend`?

## Comments

- Only one handler per message (id)
- Special send calls
- Datatype hard-coded in handler function

# Option #2

## The Handler Function

- `(*MPI_AM_Handler)(void* userdata, const void* recvbuf, MPI_Status* status)`
- `MPI_AM_Register(int tag, MPI_AM_Handler handler, void* userdata, int maxcnt, MPI_Datatype type, MPI_Comm c)`
- `MPI_AM_Deregister(int tag, MPI_Comm comm)`

## Sending an Active Message

- Just like normal point-to-point (`MPI_{Is,Rs,Ss,S}end`)

## Comments

- Only one handler per message (tag)
- Same tag namespace and matching logic as for P2P
- Integrates well in MPI

# More Options

## Multiple Handlers per ID/Tag?

- `MPI_AM_Register(MPI_AM_Handler handler, ..., void* userdata, MPI_Comm comm, MPI_AM_Func *f);`
- `MPI_AM_Deregister(MPI_AM_Func *f);`

## Multiple Datatypes per Handler?

- `(*MPI_AM_Handler)(void* userdata, MPI_Status status)`
- `MPI_AM_Register(MPI_AM_Handler handler, int tag, void* userdata, MPI_Comm comm, MPI_AM_Func *f)`
- Handler function has to `MPI_Recv` the message
- $\Rightarrow$  would lead to buffering issues

## Reintroducing needed Synchronization Semantics?

- `MPI_AM_Quiesce(MPI_Comm comm {, int id})`
- Effectively a barrier—drains all AMs



# Even More Options

## Enforce progress?

- `MPI_AM_Flush(MPI_Comm comm, int id)`
- Flushes the local AM queue (if coalescing is used)

## Collective (de)registration?

- (De)register collectively
- Can simplify tag AM tables (is this significant?)
- Initialize hardware support (?)

## Reply Messages?

- cf. `gasnet_AMReply()`
- Can be invoked in handler, would require reply handler
- Can be done on top of MPI but could be more efficient (?)

# And Even More Options (credits go to GASNet Spec)

## Differentiate between Message Sizes

- cf. GASNet short, medium and large AM sends
- Short—only register transfers (cf. inline)
- Medium—short data in temp buffer (cf. eager)
- Long—long data in buffer specified at sender (cf. rendezvous)
- MPI P2P philosophy is different
- Would enable small message optimizations (inlining etc.)

## Separate Progression

- Program logic might depend on progression (no MPI\_Wait)
- Back at the old MPI\_Progress discussion
- See tickets #25 and #154

# Restrictions on Handlers (the hard part)

## Local Computation

- Handlers **must** terminate for any input in any env
- Normal mutex locks vs. special MPI locks (like GASNet) ?
- Need to be thread-safe (also signal-safe?)
- Should finish “quickly” (performance)

## Messaging Operations

- Are not allowed to call all MPI calls (e.g., MPI\_Recv)
- All local MPI calls should be allowed (e.g., MPI\_Get\_count)
- Restricted set of remote operations
  - Only “immediate” (fire & forget) send operations
  - Generate new AMs!
  - Special send operations?
  - Restrict to reply messages?
  - Very hairy topic

# Restrictions on Handlers (the hard part)

## Memory Management

- MPI allocates memory to receive messages
- The handler “uses” the data and the runtime deallocates it
- Data can be stored in fast communication memory
- Should the handler be allowed to “keep” the memory?

- We do **not** intend to change MPI's progress semantics!
- Handlers can either be asynch or synch
- Handlers should be invoked **last** before leaving an MPI call
- Handlers can also be invoked asynchronously (e.g., in threads)
- Handlers that call MPI functions would be tail-recursive calls (no problem to handle them)
- Called MPI functions follow normal progress rule
- Deadlock/race problems not worse than in threaded environments

## Discussion

Thanks for input from Nick Edmonds and Marcin Zalewski!