



MPI Collective Workgroup

Portland, OR, November 2009



Torsten Hoefler

Agenda

- ▶ Topological Collectives
- ▶ Dick's proposal (collectives on Cartesian topologies)
- ▶ NBC proposal (#109)
- ▶ Nonblocking collective helper functions
- ▶ Graph Topology issues (a fancy new interface?)
- ▶ MPI-2.2 Errata Item

Topological Collectives

- ▶ Dynamic optimization -- all has been said, right? 😊
 - ▶ Gather, Alltoall, Reduce
- ▶ Has synchronization issues
 - ▶ thinking about “streaming” version
 - ▶ `MPI_Ineighbor_gather(..., reqs[recvcount]) + MPI_Waitany?`
 - ▶ `MPI_Ineighbor_gather(..., *req) + MPI_Cwaitany(req)?`
 - ▶ Doesn't seem to be very MPI-ish
 - ▶ might reduce usefulness of comm. schedule
- ▶ We need to find a (representative) set of applications
 - ▶ low on resources 😞
 - ▶ anybody interested (I have an implementation)?

Dick's proposal

- ▶ Expose knowledge about collective communication in Cartesian communicators
 - ▶ E.g., Alltoall along lines, Bcast along columns
- ▶ I think it cannot be done with Topologies
 - ▶ How big is the advantage?
 - ▶ Would this influence process reordering?
- ▶ Does anybody (from IBM) have more details?
- ▶ Does anybody want to pick it up?

NBC Ticket #109

- ▶ First, a minimal change
 - ▶ rev 6-2: page 66:41: add “blocking” before collective
 - ▶ clarifies but it not urgent
 - ▶ can be done with the MPI-2.2 changes?
- ▶ Second, a (voted on) nontrivial change
 - ▶ displs, count arrays can not be modified
 - ▶ simple fix:
 - ▶ rev 6-2: page 49:47: add “and displacement and count arrays” after “send buffers”
 - ▶ do we need an extra ticket?



NBC MPI-2.2 changes

- ▶ Will be applied before the Jan. meeting
- ▶ I want (need) formal reviews (responsibilities)!
 - ▶ any volunteers?
- ▶ Do we read/vote on it again? How do we organize this?
 - ▶ I don't really need to know this right now but we should define a methodology (also for future proposals)



Collective helper functions #168

- ▶ Add nonblocking versions
 - ▶ Allow nonblocking transparent libraries
 - ▶ We need for sure:
 - ▶ MPI_Comm_dup (for sure)
 - ▶ Maybe? Questionable imho:
 - ▶ MPI_Comm[_create,_split,_free]
 - ▶ MPI_Intercomm[_create,_merge]
 - ▶ MPI_[Dist_]Graph_create[_adjacent]
 - ▶ MPI_Cart_create
 - ▶ MPI_Comm_spawn[_multiple]
 - ▶ MPI_Comm[_accept,_connect,_disconnect]
 - ▶ MPI_Win[_create,_free,_fence]
 - ▶ +some I/O functions (but I/O is perfect ☺)



MPI_Cart_create (#194/#195)

- ▶ See tickets!

MPI-2 Errata

- ▶ See ticket #183
 - ▶ Typo MPI_op_free () vs. MPI_Op_free()
 - ▶ we can as well not put it to the errata and just fix it for MPI-3?