



Performance Impact of MPI Options on RMA

Keith D. Underwood, Brian Barrett

Point of this Presentation

- There are many complaints about MPI one-sided operations
 - Semantics
 - Lack of progress
 - Synchronization
 - Overhead of current synchronization
 - Passive target's write completion semantics
 - Atomic operations
 - Compare and swap
 - Fetch and increment
 - **Performance**
- ***This presentation only discusses the last one***
 - And only a narrow subset of it at that
 - *What is the overhead associated with MPI_Put?*

What is the overhead for MPI_Put?

- We attempt to provide some measurement of the overhead associated with MPI_Put()
 - We focus on what is possible, ***not what is done***
 - We ignore semantic issues and the impact of those
- As a reminder:

```
int MPI_Put(void *origin_addr, int origin_count,  
            MPI_Datatype  origin_datatype, int target_rank,  
            MPI_Aint target_disp, int target_count,  
            MPI_Datatype target_datatype, MPI_Win win)
```
- Several aspects that can impact performance
 - Long argument list
 - Dual datatype decoding
 - Windows (pointer indirection)
 - Target displacement (array lookup)


Experiment

- Compare instruction count and message rate
- Memory-to-memory copy for data move (based on SMARTMAP)
- Comparisons:
 - Complex Datatype: Transfer based on a “complex” datatype which requires decoding
 - Simple Datatype: Transfer based on simple datatype, with information encoded in datatype
 - Small args: Assume origin and target datatypes are the same (and simple). Eliminates origin and target datatypes from the call
 - SHMEM: Cray SHMEM implementation over SMARTMAP

Results

Test	Instructions	Message Rate (million messages/sec)
Complex Datatype	272	10.16
Simple Datatype	84	43.07
Small args	58	53.55
Cray SHMEM	48	59.38

Instructions



Message Rate

