



Shared Memory Extensions for MPI



Proposed API

- **MPI_COMM_ALLOC_MEM(comm, size, info, baseptr)**
 - IN comm input communicator (handle)
 - IN size size of memory segment in bytes (non-negative int)
 - IN info info argument (handle)
 - OUT baseptr pointer to beginning of memory segment allocated (choice)

- **MPI_COMM_FREE_MEM(comm, base)**
 - IN comm input communicator (handle)
 - IN base initial address of memory segment allocated by MPI_COMM_ALLOC_MEM (choice)



Proposed Semantics

- **MPI_COMM_ALLOC_MEM()**
 - Collective call
 - Allocates region of shared memory accessible by ranks in input communicator
 - No guarantee of identical baseptr across ranks
 - Otherwise, semantics are same as MPI_ALLOC_MEM()
 - Returns MPI_ERR_COMM if no shared memory is possible
 - Return MPI_ERR_NO_MEM if memory is exhausted
- **MPI_COMM_FREE_MEM()**
 - Collective call
 - Same semantics as MPI_FREE_MEM()



Why shared memory?

- **Performance**
 - **Direct load/store access between processes is more efficient than any MPI communication method**
- **Ease of use**
 - **Supports structured programming**
 - **Data is private until explicitly shared**
 - **Easier to use than threads**
 - **Where everything is shared and must be explicitly made private**
- **Reduce replicated state across processes**
- **Available on all systems (that I know of)**



Why do this in MPI? (1/2)

- **Performance**

- Integrating into MPI offers opportunity for optimization
 - POSIX shared memory allocation is not collective
 - Making it collective offers opportunity to optimize for layout and access
 - Also can make message passing more efficient
 - Affinity for multi-rail transfers
 - Potentially useful for integrating accelerators
 - May optimize checkpointing/resiliency
 - No need to replicate shared memory for all ranks
- Opportunity for using non-POSIX shared memory portably



Why do this in MPI? (2/2)

- **Integration with MPI run-time system**
 - **Simplifies shared memory allocation**
 - An MPI application would want run-time system information to allocate shared memory anyway
 - **Simplifies shared memory cleanup**
 - Leftover state on node ends up being MPI's fault anyway 😊
- **Allows integration with MPI tools**
 - Debuggers, performance debuggers, etc.
- **Ease of programming**
 - Incremental approach for existing MPI applications
 - POSIX shared memory is not easy to use
- **Ease of implementation**
 - MPI implementations already use shared memory



Hybrid MPI/Multi-Threaded Programming in Scientific Computing

Workshop to Explore the Introduction of
Threads into SNL-ASC codes

August 30, 2010

***Michael Wolf*, Mike Heroux, Erik Boman**
Scalable Algorithms Department (1416)
Extreme-scale Algorithms and Software Institute
(EASI)



Sandia National Laboratories is a multi-program laboratory managed and operated by Sandia Corporation, a wholly owned subsidiary of Lockheed Martin Corporation, for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-AC04-94AL85000.


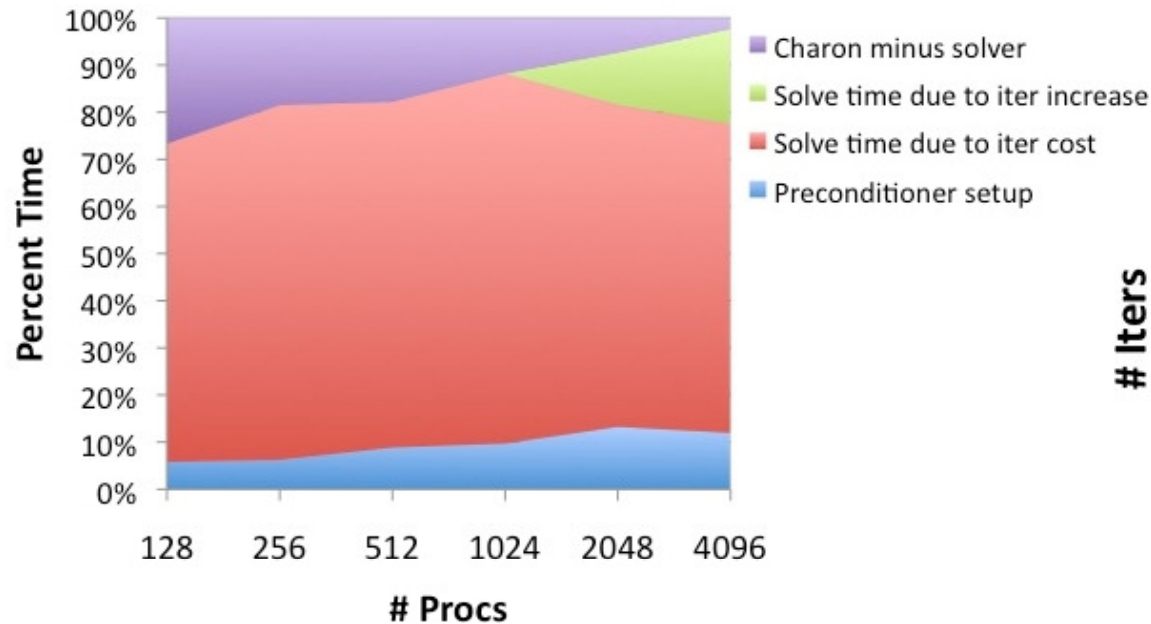




Outline

- Bimodal MPI-only/MPI + X programming
 - Integrating hybrid kernels into MPI-only applications in painless manner
 - MPI extensions for shared memory allocation
 - Simple example
 - Work in progress: Hybrid MPI/multithreaded PCG

Motivation

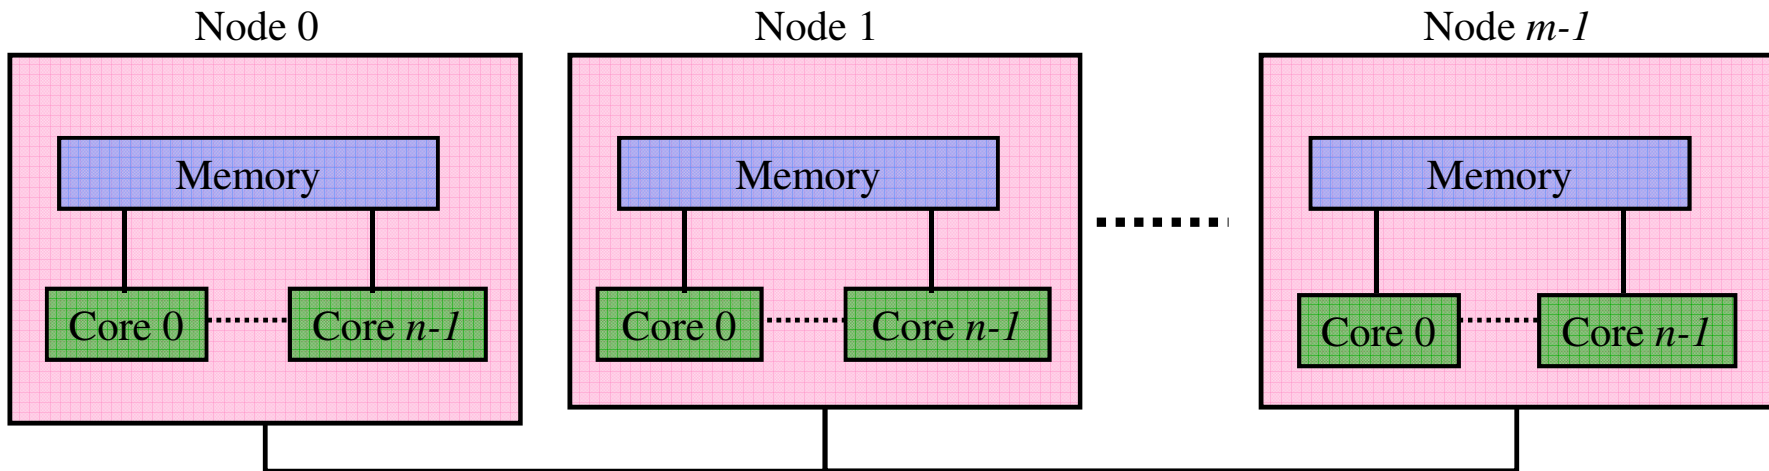


Strong scaling of Charon on TLCC (P. Lin, J. Shadid 2009)

- Domain decomposition preconditioning with incomplete factorizations
- Inflation in iteration count due to number of subdomains
- With scalable threaded triangular solves
 - Solve triangular system on larger subdomains
 - Reduce number of subdomains (MPI tasks)



MPI + Hybrid MPI/Multithreaded Programming



- Parallel machine with $p = m * n$ processors:
 - m = number of nodes
 - n = number of shared memory cores per node
- Two typical ways to program
 - Way 1: p MPI processes (flat MPI-only)
 - Way 2: m MPI processes with n threads per MPI process
- Third way (bimodal approach)
 - “Way 1” in some parts of the execution (the app)
 - “Way 2” in others (the solver)



MPI Shared Memory Allocation

Idea:

- Shared memory alloc/free functions:
 - MPI_Comm_alloc_mem
 - MPI_Comm_free_mem
- Status:
 - Available in current development branch of OpenMPI
 - Demonstrated usage with threaded triangular solve

Collaborators: B. Barrett, R. Brightwell - SNL; Vallee, Koenig - ORNL



Simple MPI Program

```
double *x = new double[4];  
double *y = new double[4];
```

```
MPIkernel1(x,y);  
MPIkernel2(x,y);
```

```
delete [] x;  
delete [] y;
```

- Simple MPI application
 - Two distributed memory/MPI kernels
- Want to replace an MPI kernel with more efficient hybrid MPI/threaded
 - Threading on multicore node



Simple MPI + Hybrid Program

```
double *x = new double[4];  
double *y = new double[4];
```

```
MPIkernel1(x,y);  
MPIkernel2(x,y);
```

```
delete [] x;  
delete [] y;
```

```
MPI_Comm_size(MPI_COMM_NODE, &nodeSize);  
MPI_Comm_rank(MPI_COMM_NODE, &nodeRank);
```

```
double *x, *y;
```

```
MPI_Comm_alloc_mem(MPI_COMM_NODE, n*nodeSize*sizeof(double),  
                  . MPI_INFO_NULL, &x);  
MPI_Comm_alloc_mem(MPI_COMM_NODE, n*nodeSize*sizeof(double),  
                  . MPI_INFO_NULL, &y);
```

```
MPIkernel1(&(x[nodeRank * n]), &(y[nodeRank * n]));
```

```
if(nodeRank==0)  
{  
    hybridKernel2(x,y);  
}
```

```
MPI_Comm_free_mem(MPI_COMM_NODE, &x);  
MPI_Comm_free_mem(MPI_COMM_NODE, &y);
```

n=4

- Very minor changes to code
 - MPIKernel1 does not change
- Hybrid MPI/Threaded kernel runs on rank 0 of each node
 - Threading on multicore node



Iterative Approach to Hybrid Parallelism

- Many sections of parallel applications scale extremely well using MPI-only model.
 - Don't change these sections much
- Approach allows introduction of multithreaded kernels in iterative fashion
 - “Tune” how multithreaded an application is
- Can focus on parts of application that don't scale with MPI-only programming
- Approach requires few changes to MPI-only sections



Iterative Approach to Hybrid Parallelism

```
MPLComm_size(MPI_COMM_NODE, &nodeSize);
MPLComm_rank(MPI_COMM_NODE, &nodeRank);

double *x, *y;

MPLComm_alloc_mem(MPI_COMM_NODE, n*nodeSize*sizeof(double),
    .               MPI_INFO_NULL, &x);
MPLComm_alloc_mem(MPI_COMM_NODE, n*nodeSize*sizeof(double),
    .               MPI_INFO_NULL, &y);

MPIkernel1(&(x[nodeRank * n]), &(y[nodeRank * n]));

if(nodeRank--0)
{
    .   hybridKernel2(x,y);
}

MPLComm_free_mem(MPI_COMM_NODE, &x);
MPLComm_free_mem(MPI_COMM_NODE, &y);
```

- Can use 1 hybrid kernel



Iterative Approach to Hybrid Parallelism

```
MPI_Comm_size(MPI_COMM_NODE, &nodeSize);
MPI_Comm_rank(MPI_COMM_NODE, &nodeRank);

double *x, *y;

MPI_Comm_alloc_mem(MPI_COMM_NODE, n*nodeSize*sizeof(double),
.                  MPI_INFO_NULL, &x);
MPI_Comm_alloc_mem(MPI_COMM_NODE, n*nodeSize*sizeof(double),
.                  MPI_INFO_NULL, &y);

if(nodeRank==0)
{
.  hybridKernel1(x,y);
.  hybridKernel2(x,y);
}

MPI_Comm_free_mem(MPI_COMM_NODE, &x);
MPI_Comm_free_mem(MPI_COMM_NODE, &y);
```

- Or use 2 hybrid kernels



Work in Progress

Bimodal MPI-only/Multithreaded PCG



PCG Algorithm

$$r_0 = b - Ax_0$$

$$z_0 = M^{-1}r_0$$

$$p_0 = z_0$$

for ($k = 0$; $k < \text{maxit}$, $\|r_k\| < \text{tol}$)

{

$$\begin{aligned} \cdot \quad & \alpha_k = \frac{r_k^T z_k}{p_k^T A p_k} \\ \cdot \quad & x_{k+1} = x_k + \alpha_k p_k \\ \cdot \quad & r_{k+1} = r_k - \alpha_k A p_k \\ \cdot \quad & z_{k+1} = M^{-1} r_{k+1} \\ \cdot \quad & \beta_k = \frac{r_{k+1}^T z_{k+1}}{r_k^T z_k} \\ \cdot \quad & p_{k+1} = z_{k+1} + \beta_k p_k \\ \cdot \quad & \end{aligned}$$

}

Used symmetric Gauss-Seidel as preconditioner (2 triangular solves)



PCG Algorithm

$$r_0 = b - Ax_0$$

$$\boxed{z_0} = \boxed{M}^{-1} \boxed{r_0}$$

$$p_0 = z_0$$

for ($k = 0$; $k < \text{maxit}$, $\|r_k\| < \text{tol}$)

{

$$\cdot \quad \alpha_k = \frac{r_k^T z_k}{p_k^T A p_k}$$

$$\cdot \quad x_{k+1} = x_k + \alpha_k p_k$$

$$\cdot \quad r_{k+1} = r_k - \alpha_k A p_k$$

$$\cdot \quad \boxed{z_{k+1}} = \boxed{M}^{-1} \boxed{r_{k+1}}$$

$$\cdot \quad \beta_k = \frac{r_{k+1}^T z_{k+1}}{r_k^T z_k}$$

$$\cdot \quad p_{k+1} = z_{k+1} + \beta_k p_k$$

}

Shared
memory
variables



PCG Algorithm – MPI part

$$r_0 = b - Ax_0$$

$$z_0 = M^{-1}r_0$$

$$p_0 = z_0$$

for ($k = 0$; $k < \text{maxit}$, $\|r_k\| < \text{tol}$)

{

$$\cdot \quad \alpha_k = \frac{r_k^T z_k}{p_k^T A p_k}$$

$$\cdot \quad x_{k+1} = x_k + \alpha_k p_k$$

$$\cdot \quad r_{k+1} = r_k - \alpha_k A p_k$$

$$\cdot \quad z_{k+1} = M^{-1} r_{k+1}$$

$$\cdot \quad \beta_k = \frac{r_{k+1}^T z_{k+1}}{r_k^T z_k}$$

$$\cdot \quad p_{k+1} = z_{k+1} + \beta_k p_k$$

}

MPI-only operations



PCG Algorithm – Threaded Part

$$r_0 = b - Ax_0$$

$$z_0 = M^{-1}r_0$$

$$p_0 = z_0$$

for ($k = 0$; $k < \text{maxit}$, $\|r_k\| < \text{tol}$)

{

$$\cdot \quad \alpha_k = \frac{r_k^T z_k}{p_k^T A p_k}$$

$$\cdot \quad x_{k+1} = x_k + \alpha_k p_k$$

$$\cdot \quad r_{k+1} = r_k - \alpha_k A p_k$$

$$\cdot \quad z_{k+1} = M^{-1}r_{k+1}$$

$$\cdot \quad \beta_k = \frac{r_{k+1}^T z_{k+1}}{r_k^T z_k}$$



$$\cdot \quad p_{k+1} = z_{k+1} + \beta_k p_k$$

}

Multithreaded block
preconditioning to reduce
number of subdomains



Preliminary PCG Results



Runtime relative to flat MPI PCG



Summary: Bimodal MPI-only / MPI + X Programming

- Interface traditional MPI-only applications with efficient MPI + X kernels
 - Only change parts of applications that don't scale
- MPI shared memory allocation useful
 - Allows seamless combination of traditional MPI programming with MPI+X kernels
- Iterative approach to multithreading
- Implemented PCG using MPI shared memory extensions and level set method
 - Effective in reducing iterations
 - Runtime did not scale (work in progress)
 - Better triangular solver algorithms needed