

MPI-3.0 Fortran Tickets

Rolf Rabenseifner, Jeff Squyres, Craig Rasmussen



Major ideas

- Solving the argument checking problems
 - Allowing also checking of wrong handle types
- Being backward compatible
- Allowing triplet array subscripts a(lb:ub:incr) everywhere, i.e., also in nonblocking routines
 - No need of MPI_Type_create_subarray for buffer descriptions.
 - No local copying
 - Prohibited on compiler level
 - Not required inside the MPI library
 - Although a quick MPI implementation can do such copying, as done currently by the compilers
- Additional features:
 - optional ierror,
 - “status ignore” through overloading
 - INTENT (IN, OUT, INOUT)

Major ideas, continued

- Fixing the Fortran-MPI-incompatibilities
 - At least with advices to users how to use Fortran in combination with MPI
- New **mpi_f08** module with all features
 - Fully Fortran 2008 compatible definition of all MPI routines
- Existing **mpi** module with several enhancements
- Different buffer handling in C-wrappers for
 - Fortran 2008 compilers
 - Older Fortran levels
- The use of `mpif.h` is strongly discouraged in the future
- In the future, i.e. all Fortran compilers support “assumed-type&rank” and `BIND(C) CHARACTER*(*)`, it is possible to have only one set of C-wrappers for all three: `mpi_f08` & `mpi` module, and `mpif.h`
- We should install a common source code infrastructure for new `mpi_f08`

Problems in discussion with Fortran committees:

- The Fortran 2008 method assumed-type & assumed-rank is currently too weak
 - `TYPE(*)`, `DIMENSION(..)` buf
- Requirement:
 - Existing interface
 - `CALL MPI_SEND(any possible actual argument)`
 - `SUBROUTINE MPI_SEND` defined as implicit interface
 - New explicit interface must allow all possible existing calls to `MPI_SEND`, i.e., all possible actual buffer arguments
- Problems:
 - `SUBROUTINE MPI_SEND(...)` `BIND(C)` with `TYPE(*)`. `DIMENSION(..)` buf does not allow assumed-size arrays as actual arguments
 - `BIND(C)` does not allow `CHARACTER*(*)` string dummy arguments
 - Hardest Problem: `MPI_(UN)PACK_EXTERNAL` has 2 choice buffers and also a string argument
- Fortran committees WG5 and J3 work on solutions → to be included into TR
 - Expectation: Solution only valid if `MPI_....` defined as `BIND(C)`

#229-A – Overview on all related Fortran Tickets

- Ticket #229-A gives an overview on all related Fortran tickets
- Don't worry that we have about 25 tickets
- They reflect separate decisions
- Before final voting,
we can combine them to one single ticket
(as done with nonblocking-collectives tickets)

-
- With all Fortran tickets together,
we can solve all problems reported in MPI-2.2

- <https://svn.mpi-forum.org/trac/mpi-forum-web/ticket/229>

#230-B - New module "USE mpi_f08"

- New module mpi_f08
 - With all new features:
 - Full compile-time argument checking
- New wording (instead of Basic & Extended Fortran Support)
 - 3 Fortran support methods:
 - include 'mpif.h'
 - use mpi
 - use mpi_f08

Specifically for the MPI Forum, the Fortran 2008 standardization committee developed new syntax TYPE(*), DIMENSION(..) to define choice buffers in a standardized way

Overview on all 3 Methods: **Include file mpif.h**

- Define all named MPI constants & declare MPI functions that return a value.
 - For each MPI routine, an implementation can choose to use an implicit or explicit interface.
 - The handles are defined as INTEGER
 - mpif.h must be valid and equivalent for both fixed- and free- source form.
 - **Advice to users: Instead of using mpif.h, the use of the mpi or mpi_f08 module is strongly encouraged. See ... Reasons ...**
-
- Almost the same as MPI-2.2
 - Only some small additions to MPI-2.2 (based on these tickets)
 - Of course, all new stuff from other MPI-3.0 tickets and chapters will be added

Overview on all 3 Methods: **Module mpi**

- Define all named MPI constants & declare MPI functions that return a value.
 - **Provide explicit interfaces for all MPI routines → compile-time argument checking.**
 - The handles are defined as INTEGER
 - Same values as in `mpif.h`
- An alternative will be in new `mpi_f08` module
- An MPI implementation may provide in the `mpi` module other features that enhance the usability of MPI while maintaining adherence to the standard.
For example, it may provide argument attributes `INTENT(IN,OUT,INOUT)` in these interface blocks.
-
- Backward compatible to MPI-2.2 for applications with bug-free syntax
 - Details later

Overview on all 3 Methods: **Module mpi_f08**

- Define all named MPI constants & declare MPI functions that return a value.
- **Provide explicit interfaces for all MPI routines → compile-time argument checking.**
- **All handles are defined with named types.**
- **If the Fortran compiler provides *assumed type* and *assumed rank*:**
 - All choice buffers are declared with **TYPE(*), DIMENSION(..)**
 - **MPI_SUBARRAYS = MPI_SUBARRAYS_SUPPORTED**
 - **non-contiguous sub-arrays are also valid in nonblocking routines.**

Only if the target compiler does **not** support *assumed type* and *assumed rank* or an equivalent non-standard alternative:

- Same requirements as for mpi module
- **MPI_SUBARRAYS = MPI_SUBARRAYS_NOT_SUPPORTED**
- non-contiguous sub-arrays are **not** valid also in nonblocking routines

```
Advice to users.  
IF (MPI_SUBARRAYS ==  
    MPI_SUBARRAYS_SUPPORTED)  
...  
ELSE IF (MPI_SUBARRAYS ==  
    MPI_SUBARRAYS_NOT_SUPPORTED)  
...  
ELSE  
    PRINT *, 'value not yet standardized'  
ENDIF
```

- With this module, **new Fortran 2008 definitions are added for each MPI routine**, except for routines that are deprecated in MPI-2.2.
- **Each argument has an INTENT=IN, OUT, or INOUT attribute** if appropriate.
- **All ierror output arguments are declared as optional**, except for user-defined callback functions (e.g., comm_copy_attr_fn) and their predefined callbacks (e.g., MPI_NULL_COPY_FN).

Available with all 3 methods

- Each application routine can freely choose one of the 3 methods.
- **Compile-time constant `MPI_F08_SUBARRAYS` is `.TRUE.`**
if all choice buffer arguments are implemented with `TYPE(*)`, `DIMENSION(..)`
 - Can be set different within all 3 methods
- `MPI_SIZEOF`, `MPI_TYPE_MATCH_SIZE`, `MPI_TYPE_CREATE_F90_INTEGER`, `MPI_TYPE_CREATE_F90_REAL` and `MPI_TYPE_CREATE_F90_COMPLEX`.
- **New `MPI_SYNC_REG` as one of several methods to solve register optimization problems.**
- Handle values are equivalent in *mpif.h* & *USE mpi*
- **Handles easily convertible to/from *USE mpi_f08***
- All new MPI-3.0 features are implemented in all 3 Fortran support methods

#231-C - Fortran argument checking with individual handles

- **This is the first major topic of new mpi_f08 module.**
- In mpi_f08, all handles use named types:

```
TYPE MPI_Comm  
SEQUENCE  
INTEGER :: MPI_VAL  
END TYPE MPI_Comm
```

- Trivial conversion between old-style (module mpi) and new handles:
comm_new%MPI_VAL = comm_old
- Same C-binding: one Fortran integer
 - The new handles do not require any changes for the wrappers written in C

#232-D: Existing module "USE mpi" with argument checking

- **Compile-time argument checking will now be mandatory for mpi module**
 - Can be done because most relevant Fortran compilers know directives like
 - !DEC\$ ATTRIBUTES NO_ARG_CHECK :: BUF
 - !\$PRAGMA IGNORE_TKR BUF
 - REAL, DIMENSION(*) :: BUF
 - To be checked: Those directives imply same argument handling as with implicit interfaces in mpif.h for choice buffers
 - May need a Fortran compiler-specific flag!
 - (such additional compiler flags are okay for compiling the module, but not a good idea for compiling the user application)
- Only if the compiler has no such method, then implicit interface is still okay
- No checking of MPI handle types (all are INTEGER)

Straw	Yes	No	Abstain
Votes:	7	-	11

#233-E: The use of 'mpif.h' is strongly discouraged

- **Advice to users: Instead of using mpif.h, the use of the mpi or mpi_f08 module is strongly encouraged. See ... Reasons ...**
 - Most have no compile-time argument checking
 - Too many bugs are in MPI applications, e.g., due to
 - Missing IERROR as last additional argument in Fortran
 - Status only as INTEGER, instead of INTEGER, DIMENSION(MPI_STATUS_SIZE)
 - Passing the wrong MPI handle types
 - ...
 - Easy migration to USE mpi
 - For syntax-correct programs
- We do not deprecate mpif.h because
 - All new features should be still added to mpif.h
 - Because there is only **one** interface per routine for mpif.h and USE mpi

#234-F: Choice buffers through "TYPE(*), DIMENSION(..)"

- This is the second major topic of new `mpi_f08` module.
- All choice buffers are declared with new Fortran 2008
 - **TYPE(*), DIMENSION(..) = assumed type & assumed rank**
- Internally, a dope-vector (i.e., a descriptor) is passed to the Fortran wrapper
(The dope-vector is generated by the compiler and describes the buffer.)
- Based on this dope-vector, a virtual or real copying from non-contiguous to contiguous scratch arrays can be done.
- This contiguous scratch buffer is under control of MPI
- With non-blocking routines, it must not be released before `MPI_Wait`
- Implication:
 - All Fortran sub-arrays can be used also in non-blocking MPI routines
 - Handling of sub-arrays can be done through Fortran syntax instead of MPI derived datatypes
- Exception: If compiler without `TYPE(*), DIMENSION(..)` → `MPI_F08_SUBARRAYS=.FALSE.`

Many thanks
to the
Fortran 2008
standardization
committee

How stable
is a TR?
Is it official?

Straw	Yes	No	Abstain
Votes:	13	-	2

Provided that `TYPE(*), DIMENSION(..)` will
have Fortran Standard Quality

#235-G: Corrections to "Problems with Fortran Bindings" (p. 481) and "Problems Due to Strong Typing" (p.482)

- No decisions, only new wording that is more correct.
- Corrections specific to the 3 methods (mpi_f08, mpi, mpif.h)

#236-H: Corrections to "Problems Due to Data Copying and Sequence Association" (MPI-2.2 page 482)

- No decision, mainly new wording to reflect the new methods and the constant MPI_F08_SUBARRAYS

```
USE mpi_f08
REAL :: buf(100)
CALL MPI_Irecv(buf(1:100:7), 15, MPI_REAL, ..., req)
...
CALL MPI_Wait(req, ...)
```

buf(1), buf(8), buf(15), buf(22), ... buf(99)

Triplet-subscripts can now be used in non-blocking routines if compile-time constant MPI_F08_SUBARRAYS is .TRUE.

- New unsolved minor problem:
Vector-subscripts still not usable in nonblocking calls, e.g. buf([1,2,24,25,33,34])
 - It is not planned to solve this problem, i.e, it is okay that only triplets are solved

#237-I: Corrections to "Problems Due to Fortran 90 Derived Types" (MPI-2.2 page 484)

- No decision, only new wording that correct this wrong section.
- Currently, MPI-2.2 says
 - “MPI does not explicitly support passing Fortran 90 derived types to dummy arguments. ...
Use of the SEQUENCE attribute may help here, somewhat.
- Reality is that MPI datatypes
 - work correctly with Fortran sequence derived types, and
 - are not guaranteed to work for Fortran non-sequence derived types.
- The section must be therefore corrected.

#238-J: Corrections to "Registers and Compiler Optimizations" (p. 371) and "A Problem with Register Optimization" (page 485)

- Additional advice about already known "Problems with MPI and Fortran optimization" (or any future method)
 - New advice with the Fortran **TARGET** attribute
 - Solves problems with: nonblocking calls, MPI_BOTTOM, and 1sided
 - New advice with the Fortran **ASYNCHRONOUS** attribute
 - Solves problems with: nonblocking calls
 - Additional helper routine **MPI_F_SYNC_REG** to substitute the user-written DD (buf)
 - Module data and common blocks also work for all three problems
- New problem with "Temporary Memory Modifications"
 - Solved with **ASYNCHRONOUS**, but not with TARGET attribute, DD, MPI_SYNC_REG, module data, or common blocks
- **This is the third major topic of new mpi_f08 module: It is about correctness of MPI nonblocking, 1-sided, and MPI_BOTTOM based calls in Fortran.**

Additional Problem:
buf1 + buf2 in one datatype
MPI_Send(buf1,1,datatype,...)

Straw	Yes	No	Abstain
Votes:	6	-	9

#239-K: IERROR optional

- In the current MPI Fortran interface, the IERROR dummy argument is mandatory.
 - In the MPI C interface, the MPI routines can be called as
 - a function (i.e., the ierror value is returned), or
 - as a procedure (i.e., ignoring the ierror value),and therefore the ierror is optional.
 - **With this ticket, the Fortran IERROR dummy argument is declared as optional** in all MPI routines that provide an IERROR.
 - Exception: For user-defined callback functions (e.g., `comm_copy_attr_fn`) and their predefined callbacks (e.g., `MPI_NULL_COPY_FN`), ierror should not be optional
 - An MPI implementation can also choose to use function overloading instead of implementing IERROR as optional
 - Advantage: Compile-time decision
 - Drawback: Doubling number of wrappers
-
- Implementation: We have to check that wrappers in C can test for this **optional** Fortran IERROR argument !!!

#240-L: New syntax used in the description
of Fortran general interfaces

- New Fortran 95 style, e.g., `INTEGER :: MPI_VERSION`
- New wordings

Straw	Yes	No	Abstain
Votes:	10	-	-

.....

#241-M: Not including old deprecated routines

- Not to include deprecated routines into the new Fortran 2008 bindings.

#242-N: Arguments with INTENT=IN, OUT, INOUT

- **Use of INTENT=IN, OUT, INOUT attributes in all new Fortran 2008 bindings.**
- The Fortran attribute INTENT(IN) is used for all arguments that are IN arguments in the language-independent notation.
- For OUT or INOUT arguments in the language-independent notation, the Fortran attributes INTENT(OUT) or INTENT(INOUT) are used, with following exceptions:
 - If there exists a constant that can be provided as actual argument, then an INTENT attribute is not specified. Examples:
 - MPI_BOTTOM and MPI_IN_PLACE for buffer arguments;
 - MPI_UNWEIGHTED in sourceweights and destweights in MPI_Dist_graph_neighbors.
 - If the argument is a handle type argument and is implemented in C with call-by-value, then INTENT(IN) is specified. Examples:
 - All file-handles in MPI_Write routines;
 - the request in MPI_Grequest_complete.
 - Exception: MPI_Cancel with INTENT(IN) request
 - Buffers without INTENT
- An MPI implementation is allowed to use more restrictive INTENT
 - E.g., INTENT(IN) for send buffers

Straw	Yes	No	Abstain
Votes:	11	-	3

#243-O: MPI_Status as a Fortran derived type

- **TYPE(MPI_Status) status**

Reasons:

- The existing status(MPI_STATUS_SIZE) array is awkward
 - status%MPI_SOURCE versus status(MPI_SOURCE)
- Wrong accesses through integer indexes cannot be detected at compile-time
 - status(1) instead of status(MPI_SOURCE)
- Wrong arrays size is also not detected at compile-time
 - INTEGER status(1) instead of INTEGER status(MPI_STATUS_SIZE)

Advice to implementors.

With the SEQUENCE attribute, one can implement this Fortran derived type with the same memory layout as the Fortran status (MPI_STATUS_SIZE) array and the C structure MPI_Status.

Straw	Yes	No	Abstain
Votes:	9	-	5

#244-P: MPI_STATUS(ES)_IGNORE with function overloading

With USE mpi_f08, the user can freely choose

- CALL MPI_Recv(buf,cnt,datatype,src,tag,comm,status,ierror)
 - CALL MPI_Recv(buf,cnt,datatype,src,tag,comm, ierror)
 - CALL MPI_Recv(buf,cnt,datatype,src,tag,comm,status)
 - CALL MPI_Recv(buf,cnt,datatype,src,tag,comm)
 - Some routines are often in the critical path:
 - Function overloading is at compile-time
 - no conditional branch at run-time
 - Function overloading is more efficient
 - Only 36 routines with status output argument
 - Same API cannot be done with OPTIONAL status argument, i.e., with OPTIONAL status, users must write
 - CALL MPI_File_write(fh,buf,count,datatype, IERROR=ierror)
- instead of
- CALL MPI_File_write(fh,buf,count,datatype, ierror)

Note that here, ierror may be needed, because in all I/O routines, ERRORS_RETURN is the default!

#245-Q: Upper and lower case letters in new Fortran bindings

- The new interfaces in mpi_f08 look like:

```
SUBROUTINE MPI_Recv(buf, count, datatype, source, tag, comm, status, ierror)
  TYPE(*), DIMENSION(..) :: buf
  TYPE(MPI_Datatype), INTENT(IN) :: datatype
  TYPE(MPI_Comm), INTENT(IN) :: comm
  INTEGER, INTENT(IN) :: count, source, tag
  TYPE(MPI_Status), INTENT(OUT) :: status
  INTEGER, OPTIONAL, INTENT(OUT) :: ierror
END
```

- Bill and Adam “hate” this → We will discuss the reasons

#246-R: MPI_ALLOC_MEM and Fortran

- How to use MPI_ALLOC_MEM together with C-Pointers in Fortran.
(instead of non-standard Cray-Pointers)

```
SUBROUTINE MPI_Alloc_mem(size, info, baseptr, ierror)
USE, INTRINSIC :: ISO_C_BINDING
INTEGER(KIND=MPI_ADDRESS_KIND), INTENT(IN) :: size
TYPE(MPI_Info), INTENT(IN) :: info
TYPE(C_PTR), INTENT(OUT) :: baseptr
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
END
```

```
SUBROUTINE MPI_Free_mem(base, ierror)
TYPE(C_PTR), INTENT(IN) :: base
INTEGER, OPTIONAL, INTENT(OUT) :: ierror
END
```

- New interface that can be used together with ALLOCATABLE arrays

Example: To be done

#247-S: All new Fortran 2008 bindings – Part 1

- This ticket shows the principles and some special details

#248-T: All new Fortran 2008 bindings – Part 2

A.4 Fortran 2008 Bindings with module mpi_f08

A.4.1 Point-to-Point Communication Fortran Bindings

Half-automatically generated
from the existing Fortran
bindings!

SUBROUTINE MPI_Bsend(buf, count, datatype, dest, tag, comm, ierror) BIND(C)

TYPE(*), DIMENSION(..) :: buf

INTEGER, INTENT(IN) :: count, dest, tag

TYPE(MPI_Datatype), INTENT(IN) :: datatype

TYPE(MPI_Comm), INTENT(IN) :: comm

INTEGER, OPTIONAL, INTENT(OUT) :: ierror

END

SUBROUTINE MPI_Bsend_init(buf, count, datatype, dest, tag, comm, request, ierror) BIND(C)

TYPE(*), DIMENSION(..) :: buf

INTEGER, INTENT(IN) :: count, dest, tag

TYPE(MPI_Datatype), INTENT(IN) :: datatype

TYPE(MPI_Comm), INTENT(IN) :: comm

TYPE(MPI_Request), INTENT(OUT) :: request

INTEGER, OPTIONAL, INTENT(OUT) :: ierror

END

...

#250-V: Minor Corrections in Fortran Interfaces

- Typo in the existing Fortran Interface of MPI_INTERCOMM_MERGE:
 - INTRACOMM → NEWINTRACOMM
 - Remove double definition of request in the Fortran binding type declaration part of MPI_SEND_INIT and MPI_BSEND_INIT
-
- Regular errata – in all three Fortran support methods

#252-W: Substituting dummy argument name "type" by "datatype" or "oldtype"

- To minimize conflicts with language keywords (TYPE in Fortran), the dummy argument name "type" is substituted by "datatype" or "oldtype".

-
- Note, with explicit interfaces, the user can freely choose between
 - Positional argument lists
 - CALL MPI_Send(buf,cnt,datatype,src,tag,comm,ierr)
 - Keyword-based argument lists and mixed lists
 - CALL MPI_Send(buf,cnt,datatype,source=src, &
& tag=13,comm=MPI_COMM_WORLD,ierror=ierr)

→ Dummy argument names *should be done correctly & should make sense*



Conclusion

- Although separated into tickets, it is one big packet.
- We needed 3 years to detect most of the MPI-Fortran-problems (1994-1997)
- We needed additional 13 year to hopefully solve them all (1997-2010) !
- And we still detect new ones 😊 😞 😊 😞 😊
- Thanks to the Fortran Standardization Committees WG5 and J3 for their working together to solve the MPI-Fortran incompatibility problems.

Appendix

- List of all Fortran tickets

Tickets related to new MPI-3.0 Fortran Interface

- #229-A - Overview over all related tickets
- #230-B - New module "USE mpi_f08"
- #231-C - Fortran argument checking with individual handles
- #232-D - Existing module "USE mpi" with argument checking
- #233-E - The use of 'mpif.h' is strongly discouraged
- #234-F - Choice buffers through TYPE(*) DIMENSION(..) declarations
- #235-G - Corrections to "Problems with Fortran Bindings" (p.481) & "Strong Typing" (482)
- #236-H - Corrections to "Problems Due to Data Copying and Sequence Association" (482)
- #237-I - Corrections to problems due to "Fortran 90 Derived Types" (MPI-2.2 page 484)
- #238-J - Corrections to "A Problem with Register Optimization" (pages 371 and 485)
- #239-K - IERROR optional
- #240-L - New syntax used in all three (mpif.h, mpi, mpi_f08)

<https://svn.mpi-forum.org/trac/mpi-forum-web/ticket/229> .. 253

Tickets related to new MPI-3.0 Fortran Interface

- #241-M - Not including old deprecated routines from MPI-2.0 - MPI-2.2
- #242-N - Arguments with INTENT=IN, OUT, INOUT
- #243-O - MPI_Status as a Fortran derived type
- #244-P - MPI_STATUS(ES)_IGNORE with function overloading
- #245-Q - Upper and lower case letters in new Fortran bindings
- #246-R - MPI_ALLOC_MEM and Fortran
- #247-S - All new Fortran 2008 bindings - Part 1
- #248-T - All new Fortran 2008 bindings - Part 2
- #249-U - Alternative formulation for Section 16.2 Fortran Support
- #250-V - Minor Corrections in Fortran Interfaces
- #252-W - Substituting dummy argument name "type" by "datatype" or "oldtype"
- #253-X - mpi_f08 Interfaces for new MPI-3.0 routines (not yet done)
- #251 is currently a helper ticket: Printable version of all tickets together