

New, mo' bettah attributes

Ticket #304

Squyres and Tillier

Madrid, September 2013

Last edit: v0.8, 11 Sep 2013

Current attribute system sucks

- Complex rules for inter-language interop
 - Including deprecated, truncating Fortran routines
- Only used on some handle types
- Big ambiguities in attribute callback definitions
 - See [Microsoft June 2013 Forum presentation](#)
 - Fixing the current system is... difficult
- Best solution is a new system

Goals of new system

- Cache on any handle type
- Key objects have a proper handle type
- Key objects are global (vs. per handle type)
- Simple language interop between Fortran/C
- Clearly defined callback semantics
- Deprecate old keyval/attribute functions

Design Overview / Goals

- “Key = value” system
 - Keys are ref-counted handles: MPI_Key
 - Values are MPI_Aint (*no translation between C/ Fortran*)
- Values can be cached on *any* MPI handle type
- Callbacks when handles are:
 - Copied
 - Freed
 - Destroyed (*this is new*)

MPI-3 p103:30 Advice to Users: Do we care?

- C users may be tempted to avoid the usage of `MPI_GET_ADDRESS` and rely on the availability of the address operator `&`. Note, however, that `&cast-expression` is a pointer, not an address. ISO C does not require that the value of a pointer (or the pointer cast to `int`) be the absolute address of the object pointed at -- although this is commonly the case. Furthermore, referencing may not have a unique definition on machines with a segmented address space. The use of `MPI_GET_ADDRESS` to “reference” C variables guarantees portability to such machines as well.

Probably not. This is more about pointer math than lossless translation pointer \leftrightarrow `MPI_Aint`. Bill will check the standard to be sure, but is pretty sure.

Overview

- Deprecate all old attribute functions
- New pre-defined key handles:
 - MPI_KEY_NULL
 - MPI_KEY_TAG_UB
 - MPI_KEY_HOST
 - MPI_KEY_IO
 - MPI_KEY_WTIME_IS_GLOBAL

Key functions

- MPI_Key_create
 - MPI_Key_free
 - MPI_Key_c2f
 - MPI_Key_f2c
-
- ...just like many other MPI handle types

Key functions: create

- `MPI_Key_create(value_copy_fn, value_free_fn, value_destroy_fn, MPI_Aint context, MPI_Key *key)`
 - Function pointers for when the handle is copied, freed, and destroyed (prototypes discussed later)
 - Context to be passed to callbacks
- Once created, this key can be used with *any* handle type
 - `MPI_Comm`, `MPI_Datatype`, `MPI_Key`, ...etc.

Key functions: free and c2f/f2c

- As you would expect
 - `MPI_Key_free(MPI_Key *key)`
 - `MPI_Key MPI_Key_f2c(MPI_Fint fkey)`
 - `MPI_Fint MPI_Key_c2f(MPI_Key ckey)`

Value functions

- MPI_Value_set
- MPI_Value_get
- MPI_Value_clear
 - For explicitly unsetting a value on a handle

Value functions: set

- `MPI_Value_set(MPI_Key key, int handle_type, MPI_Handle handle, MPI_Aint value)`
 - `handle_type`: indicates type of handle (e.g., `MPI_HANDLE_COMM`, `MPI_HANDLE_DATATYPE`, ...)
 - `handle`: type that allows any handle
 - `MPI_COMM_WORLD`, `MPI_INT`, `my_comm`, etc.
 - `value`: `MPI_Aint` of the value
- If a value was already set on (key, handle)
 - Behaves as if `MPI_VALUE_CLEAR` was first called to clear the current value before new value is set

Value functions: get

- `MPI_Value_get(MPI_Key key, int handle_type, MPI_Handle handle, MPI_Aint *value, int *flag)`
 - key, handle_type, handle: Same as `VALUE_SET`
 - value: pointer to `MPI_Aint`
 - flag: logical, indicates whether a value is set on this handle / was filled into the “value” OUT param

Value functions: clear

- `MPI_Value_clear(key, handle_type, handle)`
 - `key`, `handle_type`, `handle`: Same as `VALUE_SET`
 - Explicitly clear a value on a handle such that `MPI_VALUE_GET` on that `(key, handle)` will return `flag==false`
 - Not an error to clear a `(key, handle)` for which no value is set

Key callbacks

- All return void
- Do not affect the semantics of invoking MPI functions
 - E.g., MPI_COMM_FREE functionality is not impacted by the return of a value callback
 - Because otherwise, you get deferred destruction craziness (see prior Microsoft presentation)

Key callbacks: Copy

- Function params
 - key
 - handle_type: enum indicating type of handle
 - old_handle
 - new_handle
 - key_context: from Key_create
 - old_value: value set on old handle
 - new_value: value to be set on new handle
 - flag (OUT): logical, indicates whether a new value was set

Key callbacks: Free

- Invoked at MPI_*_FREE time
- Function params
 - key
 - handle_type: enum indicating type of handle
 - handle
 - key_context: from Key_create
 - value: value currently set on handle

Key callbacks: Free

- Invoked for every value set on the handle
 - Before the handle is freed
 - It's ok to use the handle in the callback

`MPI_Comm_free(comm)`
→ invokes `my_callback()`

```
my_callback(..., comm, ...) {  
    MPI_Send(..., comm);  
}
```



This is ok!

Key callbacks: Free

- Calling MPI_*_FREE on the handle in the callback is erroneous

MPI_Comm_free(comm)
→ invokes my_callback()

```
my_callback(..., comm, ...) {  
    MPI_Comm_free(comm);  
}
```



This is NOT ok!

Key callbacks: Free

- Can call `MPI_*_FREE` on *another* handle in a free callback
 - Well-defined, but can lead to reentrancy issues due to recursion
 - Advice to Users: this is complicated; don't do it

`MPI_Comm_free(comm)`
→ invokes `my_callback()`

```
my_callback(..., comm, ...) {  
    MPI_Comm_free(othercomm);  
}
```



This is ok... but risky

Key callbacks: Free

- If Key was previously MPI_KEY_FREEd
 - key argument will be MPI_KEY_NULL
 - But the key context will still be valid

Key callbacks: Free

- Can VALUE_GET on this handle in this callback
 - VALUE_GET(this_key, this_handle) → flag == false
 - VALUE_GET(other_key, this_handle) → flag == ?
 - Ordering of free callbacks is non-deterministic
- Cannot VALUE_SET or VALUE_CLEAR any key *on this handle* in this callback
 - This is under debate. Other option is to allow VALUE_SET/VALUE_CLEAR in callback and queue up future free callbacks

Key callbacks: Destroy

- Invoked:
 - When corresponding MPI object is destroyed
 - Before MPI_VALUE_CLEAR returns
 - Before MPI_VALUE_SET replaces an old value
- Value destroy function params
 - key
 - handle_type: enum indicating type of handle
 - handle
 - key_context: from Key_create
 - value: value currently set on handle

Key callbacks: Destroy

- If Key was previously MPI_KEY_FREEd
 - key argument will be MPI_KEY_NULL
 - But the key context will still be valid
- (just like free callback behavior)

Performance implication

- By allowing attributes on any MPI handle type:
 - Can cache values on an MPI_Request
 - MPI_TEST/MPI_WAIT may invoke attribute callbacks
 - This (probably) adds a single extra “if” statement in the TEST/WAIT completion code path
- All other handle types invoke callbacks at MPI_*_FREE time
 - Performance implications largely irrelevant



KTHXBYE