MPI_Stream

Canvassing feedback on initial ideas

What is an “MPI stream”?

- ★ From sender to receiver only
  - Joined by an existing communicator
- ★ Sender and receiver use intrinsic data-types as the smallest unit of transmission
  - Message boundaries are ignored
- ★ Sender can send any amount(s) of data
- ★ Receiver can receive any amount up to what is available
  - Buffer underrun - receiver asks for 33 integers but only 16 available so gets 16 integers rather than blocking, receiver must ask again to get the rest
- ★ Ordered and reliable

MPI stream example

Sender sends 5 items, then 7 items, then 2 items

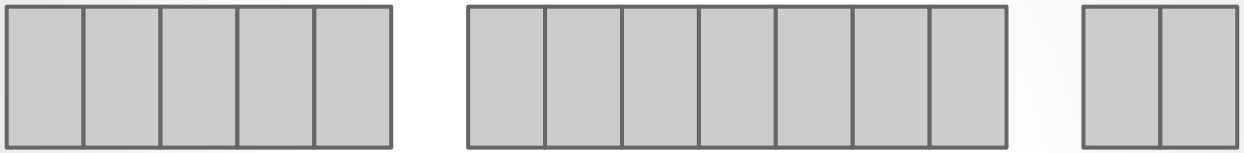

MPI stream example

Sender sends: 5 items, then 7 items, then 2 items

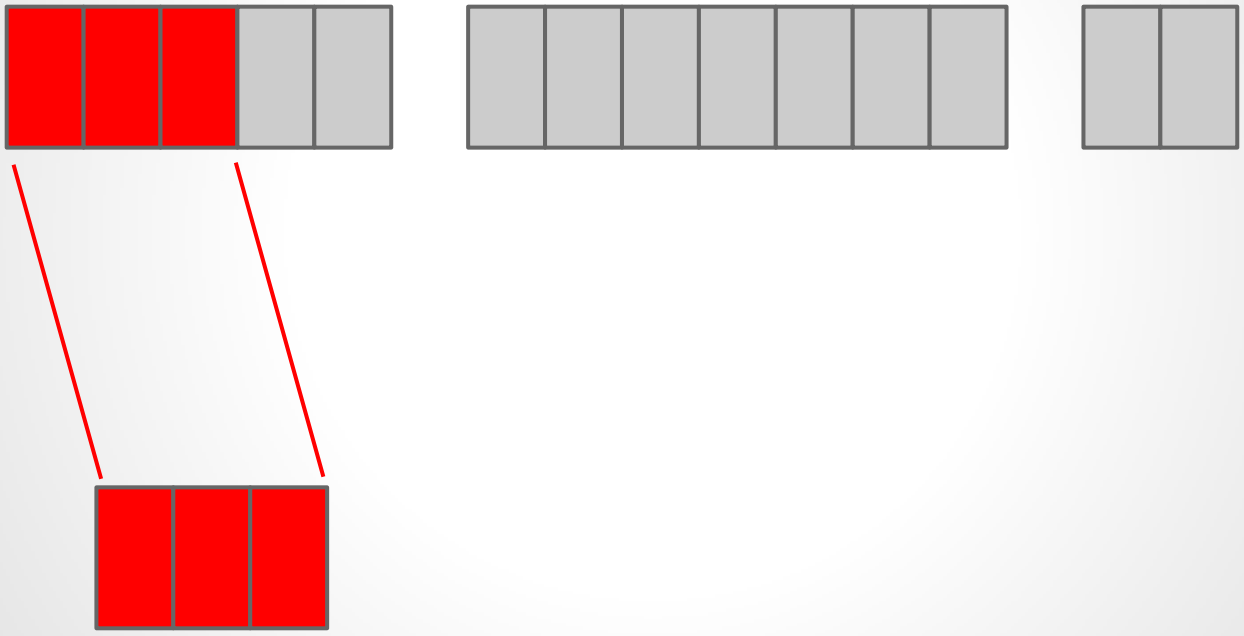

Receiver receives first 3 items

MPI stream example

Sender sends: 5 items, then 7 items, then 2 items

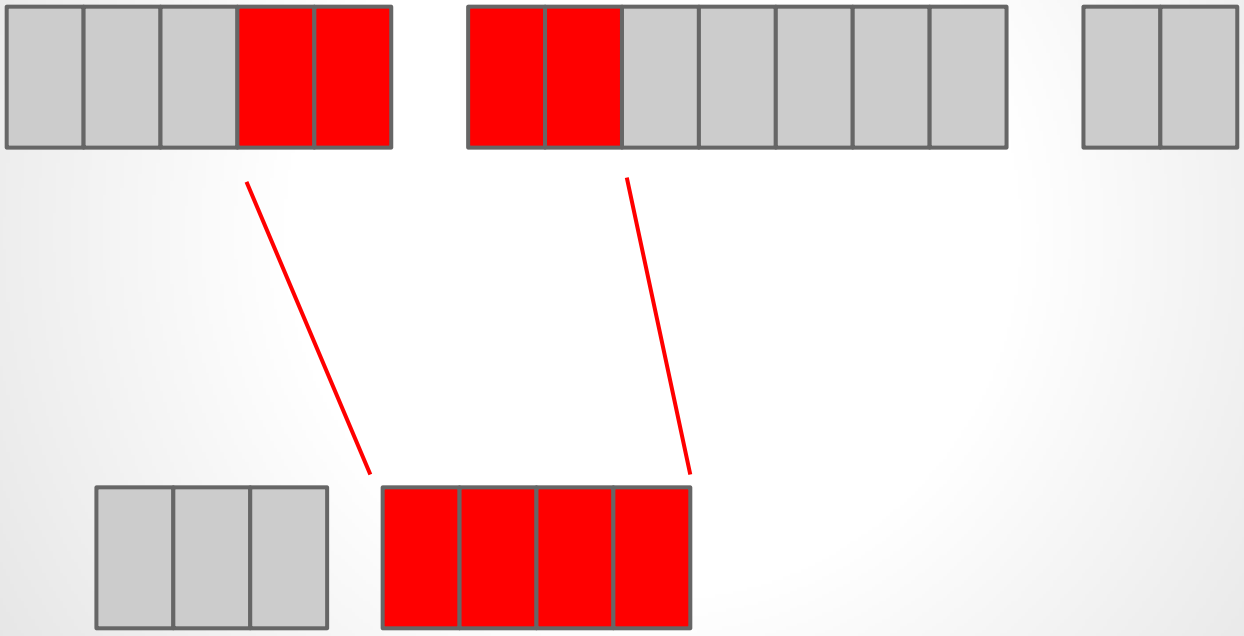

Receiver receives next 4 items

MPI stream example

Sender sends: 5 items, then 7 items, then 2 items

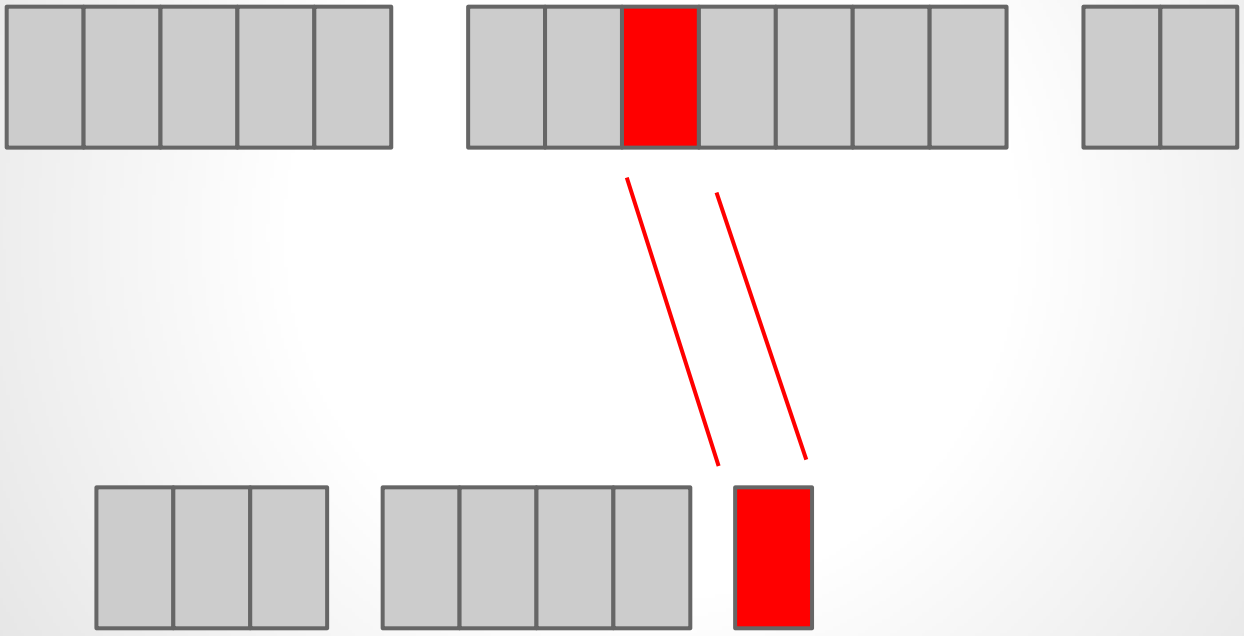

Receiver receives next 1 item

MPI stream example

Sender sends: 5 items, then 7 items, then 2 items

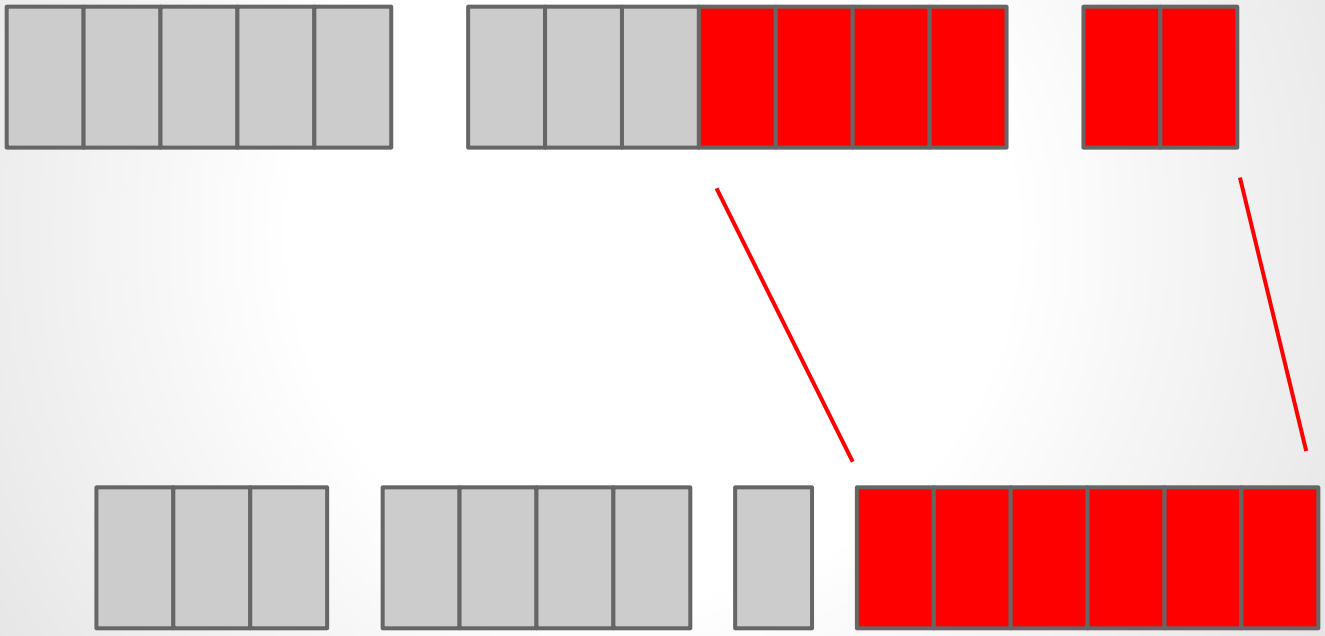

Receiver tries to receive next 20 items, gets 6

Doubts

- ★ Why is this better than TCP?
  - It drops a lot of the stuff from TCP that HPC does not want
  - It adds MPI datatypes
  - Can use high performance networks (without TOE)
- ★ Performance argument
  - Hardware needs to be ready and able to support the proposed semantics
- ★ Is this a solution in search of a problem?

Questions for app developers

★ Would you use this?
★ What would it buy you?