

BigCount Solutions for MPI-4

MPI Forum

Chicago, USA, May 2019

Chattanooga meeting results

A concise summary



Check out these awesome
function pointers for MPI!

But you didn't solve BigCount





...but function pointers are
shiny and awesome!

Chattanooga takeaways

1. Function pointers for C / objects for Fortran for MPI are neat
2. Function pointers would require a major rewrite of apps
 - a. Backwards compatibility will be... tricky
3. *The real issue is to solve BigCount for MPI-4*

Solving BigCount



List of functions affected (133)

MPI_Accumulate	MPI_File_read_at	MPI_lalltoallv	MPI_Neighbor_alltoallv	MPI_Ssend_init
MPI_Allgather	MPI_File_read_at_all	MPI_lalltoallw	MPI_Neighbor_alltoallw	MPI_Startall
MPI_Allgatherv	MPI_File_read_at_all_begin	MPI_lbroadcast	MPI_Pack	MPI_Status_set_elements
MPI_Allreduce	MPI_File_read_ordered	MPI_lbsend	MPI_Pack_external	MPI_Status_set_elements_x
MPI_Alltoall	MPI_File_read_ordered_begin	MPI_lexscan	MPI_Pack_external_size	MPI_T_cvar_handle_alloc
MPI_Alltoallv	MPI_File_read_shared	MPI_lgather	MPI_Pack_size	MPI_T_pvar_handle_alloc
MPI_Alltoallw	MPI_File_write	MPI_lgatherv	MPI_Put	MPI_Testall
MPI_Bcast	MPI_File_write_all	MPI_lmrecv	MPI_Raccumulate	MPI_Testany
MPI_Bsend	MPI_File_write_all_begin	MPI_lneighbor_allgather	MPI_Recv	MPI_Testsome
MPI_Bsend_init	MPI_File_write_at	MPI_lneighbor_allgatherv	MPI_Recv_init	MPI_Type_contiguous
MPI_CONVERSION_FN_NULL	MPI_File_write_at_all	MPI_lneighbor_alltoall	MPI_Reduce	MPI_Type_create_hindexed
MPI_Comm_spawn_multiple	MPI_File_write_at_all_begin	MPI_lneighbor_alltoallv	MPI_Reduce_local	MPI_Type_create_hindexed_block
MPI_Dist_graph_neighbors_count	MPI_File_write_ordered	MPI_lneighbor_alltoallw	MPI_Reduce_scatter	MPI_Type_create_hvector
MPI_Exscan	MPI_File_write_ordered_begin	MPI_lrecv	MPI_Reduce_scatter_block	MPI_Type_create_indexed_block
MPI_File_iread	MPI_File_write_shared	MPI_lreduce	MPI_Rget	MPI_Type_create_struct
MPI_File_iread_all	MPI_Gather	MPI_lreduce_scatter	MPI_Rget_accumulate	MPI_Type_get_extent_x
MPI_File_iread_at	MPI_Gatherv	MPI_lreduce_scatter_block	MPI_Rput	MPI_Type_get_true_extent_x
MPI_File_iread_at_all	MPI_Get	MPI_lrsend	MPI_Rsend	MPI_Type_indexed
MPI_File_iread_shared	MPI_Get_accumulate	MPI_lscan	MPI_Rsend_init	MPI_Type_size_x
MPI_File_iwrite	MPI_Get_count	MPI_lscatter	MPI_Scan	MPI_Type_vector
MPI_File_iwrite_all	MPI_Get_elements	MPI_lscatterv	MPI_Scatter	MPI_Unpack
MPI_File_iwrite_at	MPI_Get_elements_x	MPI_lsend	MPI_Scatterv	MPI_Unpack_external
MPI_File_iwrite_at_all	MPI_Graph_neighbors_count	MPI_lssend	MPI_Send	MPI_Waitall
MPI_File_iwrite_shared	MPI_lallgather	MPI_lmrecv	MPI_Send_init	MPI_Waitany
MPI_File_read	MPI_lallgatherv	MPI_Neighbor_allgather	MPI_Sendrecv	MPI_Waitsome
MPI_File_read_all	MPI_lallreduce	MPI_Neighbor_allgatherv	MPI_Sendrecv_replace	
MPI_File_read_all_begin	MPI_lalltoall	MPI_Neighbor_alltoall	MPI_Ssend	

List of functions affected (133)

MPI_Accumulate	MPI_File_read_at	MPI_lalltoallv	MPI_Neighbor_alltoallv	MPI_Ssend_init
MPI_Allgather	MPI_File_read_at_all	MPI_lalltoallw	MPI_Neighbor_alltoallw	MPI_Startall
MPI_Allgatherv	MPI_File_read_at_all_begin	MPI_lbroadcast	MPI_Pack	MPI_Status_set_elements
MPI_Allreduce	MPI_File_read_ordered	MPI_lbsend	MPI_Pack_external	MPI_Status_set_elements_x
MPI_Alltoall	MPI_File_read_ordered_begin	MPI_lexscan	MPI_Pack_external_size	MPI_T_cvar_handle_alloc
MPI_Alltoallv	MPI_File_read_shared	MPI_lgather	MPI_Pack_size	MPI_T_pvar_handle_alloc
MPI_Alltoallw	MPI_File_write	MPI_lgatherv	MPI_Put	MPI_Testall
MPI_Bcast	MPI_File_write_all	MPI_lmrecv	MPI_Raccumulate	MPI_Testany
MPI_Bsend	MPI_File_write_all_begin	MPI_lneighbor_alltoallv	MPI_Recv	MPI_Testsome
MPI_Bsend_init	MPI_File_write_at		MPI_Recv_init	MPI_Type_contiguous
MPI_CONVERSION_FN_NULL	MPI_File_write_at_all		MPI_Reduce	MPI_Type_create_hindexed
MPI_Comm_spawn_multiple	MPI_File_write_at_all_begin		MPI_Reduce_local	MPI_Type_create_hindexed_block
MPI_Dist_graph_neighbors_count	MPI_File_write_ordered		MPI_Reduce_scatter	MPI_Type_create_hvector
MPI_Exscan	MPI_File_write_ordered_begin		MPI_Reduce_scatter_block	MPI_Type_create_indexed_block
MPI_File_iread	MPI_File_write_shared	MPI_ireduce	MPI_Rget	MPI_Type_create_struct
MPI_File_iread_all	MPI_Gather	MPI_ireduce_scatter	MPI_Rget_accumulate	MPI_Type_get_extent_x
MPI_File_iread_at	MPI_Gatherv	MPI_ireduce_scatter_block	MPI_Rput	MPI_Type_get_true_extent_x
MPI_File_iread_at_all	MPI_Get	MPI_lrsend	MPI_Rsend	MPI_Type_indexed
MPI_File_iread_at_all	MPI_Get_accumulate	MPI_lscan	MPI_Rsend_init	MPI_Type_size_x
MPI_File_iread_shared	MPI_Get_count	MPI_lscatter	MPI_Scan	MPI_Type_vector
MPI_File_iwrite	MPI_Get_elements	MPI_lscatterv	MPI_Scatter	MPI_Unpack
MPI_File_iwrite_all	MPI_Get_elements_x	MPI_lsend	MPI_Scatterv	MPI_Unpack_external
MPI_File_iwrite_at		MPI_lssend	MPI_Send	MPI_Waitall
MPI_File_iwrite_at_all	MPI_Graph_neighbors_count	MPI_lmrecv	MPI_Send_init	MPI_Waitany
MPI_File_iwrite_shared	MPI_lallgather	MPI_Neighbor_allgather	MPI_Sendrecv	MPI_Waitsome
MPI_File_read	MPI_lallgatherv	MPI_Neighbor_allgather	MPI_Sendrecv_replace	
MPI_File_read_all	MPI_lallreduce	MPI_Neighbor_allgatherv		
MPI_File_read_all_begin	MPI_lalltoall	MPI_Neighbor_alltoall		

Constraints / Assumptions

1. All “count” arguments in a given API will be either `int` or `MPI_Count`

```
MPI_Bad(int count1, MPI_Count count2)
```


```
MPI_Good(int count1, int count2)
```

```
MPI_Good(MPI_Count count1, MPI_Count count2)
```

Constraints / Assumptions

2. The only thing that matters is the underlying type map

```
// Use a MPI_Count-sized (big) count  
MPI_Send(..., (4B+4), MPI_INT, ...);
```



Works just fine


```
// Use an int-sized (small) count  
MPI_Type_contiguous((1B+1), MPI_INT, &dtype);  
MPI_Recv(..., 4, dtype, ...)
```

Constraints / Assumptions


3. Must solve for both C and Fortran

- a. Solve it in the “same” way for both languages (from the user perspective)

No:



Yes:





Constraints / Assumptions


4. We need to enable C++ BigCount

- a. Remember: C and C++ are quite different languages at this point
- b. Open question: what does that mean for the C bindings?

No:



Yes:



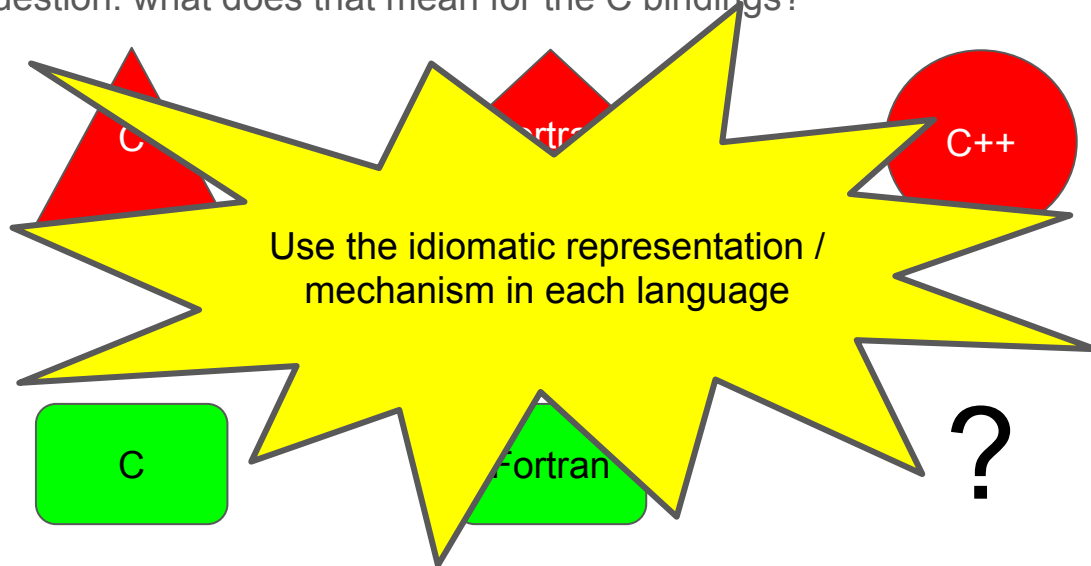
?

Constraints / Assumptions

4. We need to enable C++ BigCount

- a. Remember: C and C++ are quite different languages at this point
- b. Open question: what does that mean for the C bindings?

No:



Yes:

Constraints / Assumptions

5. May require new(ish) compilers

Five main options

	C	C++	Fortran
1.	No change / keep all counts as <code>int</code>		
2.	Change all counts to <code>MPI_Count</code>		
3.	“_X” functions	“_X” functions	“_X” functions
4.	Function pointers	Function pointers (?)	Objects
5.	C11 <code>_Generic</code>	Function overloading	Generic functions

1. No changes / keep all counts `int` in MPI-4

Users be like:



2. Change all counts to `MPI_Count`

Users be like:



3. “_X” functions

```
MPI_Send(..., int count, ...)
```

```
MPI_Send_x(..., MPI_Count count, ...)
```

Discussed in prior Forum meetings:

- PRO: Simple
- PRO: Started down this path in MPI-3.0
- CON: Lots of new functions
- CON: Ugly (evil)
- CON: (Very) Bad precedent



4. Function pointers / objects

```
comm.send(..., int count, ...)  
comm.send(..., MPI_Count count, ...)
```

Discussed in Chattanooga / Jan 2019

- PRO: Opens the door for many interesting possibilities (e.g., QMPI, sessions)
- CON: Problems still to be solved, including:
 - Backwards compatibility
 - Incremental app upgrades
- CON: Potentially a lot of new text for MPI-4



5. “Polymorphism”

Build upon Jeff Hammond’s prior C11 discussions

Goal: user just calls `MPI_Send(...)`
with either `int` or `MPI_Count`

5. “Polymorphism”

Turns into three different things:

C: C11 `_Generic`


C++: Function overloading

Fortran: Generic functions

Goal: user just calls `MPI_Send(...)`
with either `int` or `MPI_Count`

5a. C11 _Generic

- Implementation provides two functions with distinct names
 - Standardize the names for PMPI reasons
- Compiler chooses between them at compile time
- Uses the `C11 _Generic` keyword
 - Does not exist prior to C11
 - Does not exist in C++
- We tell users to still use “main” name (e.g., `MPI_Send`)



5a. C11 _Generic

- Implementation provides two functions with distinct names
 - Standardize the names for PMPI reasons
- Compiler chooses between them at compile time
- Uses the C11 `_Generic` keyword
 - Does not exist prior to C11 → Pre-C11 implementations get `int`
 - Does not exist in C++
- We tell users to still use “main” name (e.g., `MPI_Send`)

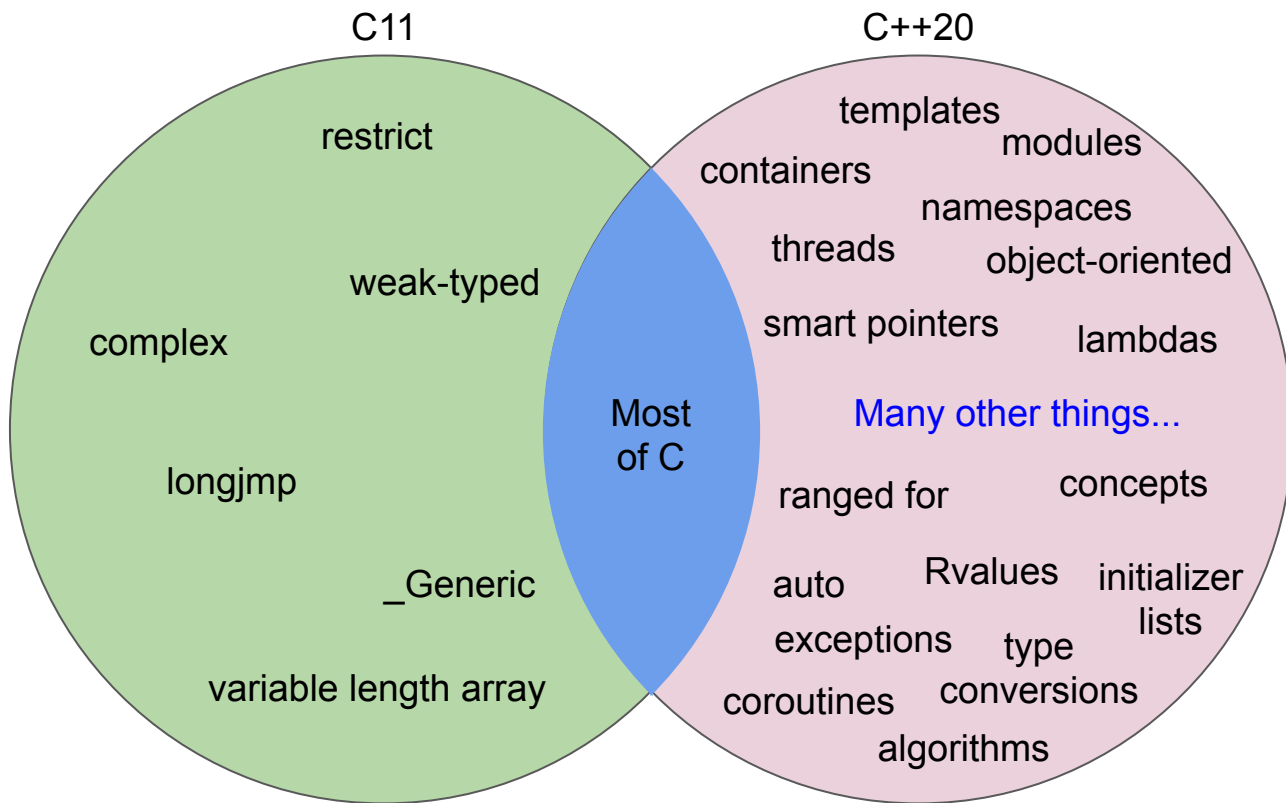


5a. C11 _Generic

- Implementation provides two functions with distinct names
 - Standardize the names for PMPI reasons
- Compiler chooses between them at compile time
- Uses the `C11 _Generic` keyword
 - Does not exist prior to C11
 - Does not exist in C++ → Use a different solution for C++
- We tell users to still use “main” name (e.g., `MPI_Send`)



Sidenote: why is C++ no longer “C with objects”?



Sidenote: why is C++ no longer “C with objects”?

Let's not forget things that are spelled the same in C and C++, but are actually different

- Boolean type
- sizeof(struct types)
- Character literal types
- Implicit extern- vs. file-scoping
- inline functions



Dr. Evil Approved™

5a. C11 _Generic: example

```
// mpi.h
typedef long MPI_Count;

#define MPI_Send(buf, count, dt, rank, tag, comm) \
    _Generic(count, \
        default: MPI_Send_int, \
        int: MPI_Send_int, \
        MPI_Count: MPI_Send_count \
    )(buf, count, dt, rank, tag, comm)

int MPI_Send_int(..., int count, ...)
{
    return MPI_Send_count(...,
        (MPI_Count)count, ...);
}

int MPI_Send_count(..., MPI_Count count, ...)
{
    ...
}
```

5a. C11 _Generic: example

```
// mpi.h
typedef long MPI_Count;

#define MPI_Send(buf, count, dt, rank, tag, comm) \
    _Generic(count, \
        default: MPI_Send_int, \
        int: MPI_Send_int, \
        MPI_Count: MPI_Send_count \
    )(buf, count, dt, rank, tag, comm)

int MPI_Send_int(..., int count, ...)
{
    return MPI_Send_count(...,
        (MPI_Count)count, ...);
}

int MPI_Send_count(..., MPI_Count count, ...)
{
    ...
}
```

```
#include <mpi.h>

int main(int argc, char **argv)
{
    ...

    int i = 32;
    MPI_Send(..., i, ...);
    MPI_Send(..., 32, ...);
    MPI_Send(..., (MPI_Count)32, ...);

    ...

    MPI_Count bigI = 8589934592;
    MPI_Send(..., bigI, ...);
    MPI_Send(..., 8589934592, ...);

    ...
}
```

5b. C++ function overloading

- C++ does not have the `_Generic` keyword
- MPI implementation provides two C++ functions with the same name
- PMPI interfaces will need to be compiled with same compiler
- `mpi.h` must distinguish between C and C++



5b. C++ function overloading

- C++ does not have the `_Generic` keyword
- MPI implementation provides two C++ functions with the same name
- PMPI interfaces will need to be compiled with same compiler
- `mpi.h` must distinguish between C and C++ → ...unless we make `<mpi.hpp>`



5b. C++ function overloading: example

```
// mpi.hpp
```

```
int MPI_Send(..., int count, ...)
{
    return MPI_Send(...,
        static_cast<long>(count), ...);
}
```

```
int MPI_Send(..., MPI_Count count, ...)
{
    ...
}
```


5b. C++ function overloading: example

```
// mpi.hpp

int MPI_Send(..., int count, ...)
{
    return MPI_Send(...,
        static_cast<long>(count), ...);
}

int MPI_Send(..., MPI_Count count, ...)
{
    ...
}
```

```
#include <mpi.hpp>

int main(int argc, char **argv)
{
    ...

    int i = 32;
    MPI_Send(..., i, ...);
    MPI_Send(..., 32, ...);
    MPI_Send(..., (MPI_Count)32, ...);

    ...

    MPI_Count bigI = 8589934592;
    MPI_Send(..., bigI, ...);
    MPI_Send(..., 8589934592, ...);

    ...
}
```

5a+5b. Note that user code is 99.9% the same

```
#include <mpi.h>

int main(int argc, char **argv)
{
    ...

    int i = 32;
    MPI_Send(..., i, ...);
    MPI_Send(..., 32, ...);
    MPI_Send(..., (MPI_Count)32, ...);

    ...

    MPI_Count bigI = 8589934592;
    MPI_Send(..., bigI, ...);
    MPI_Send(..., 8589934592, ...);

    ...
}
```

`#include <mpi.h>` ← The only difference

```
int main(int argc, char **argv)
{
    ...

    int i = 32;
    MPI_Send(..., i, ...);
    MPI_Send(..., 32, ...);
    MPI_Send(..., (MPI_Count)32, ...);

    ...

    MPI_Count bigI = 8589934592;
    MPI_Send(..., bigI, ...);
    MPI_Send(..., 8589934592, ...);

    ...
}
```

5c. Fortran generic functions

- Implementation provides a single interface with two subroutines of distinct names
 - Names standardized for PMPI reasons
- Only do this for `mpi_f08`
 - Not the `mpi` module or `mpif.h`



F08

5c. Fortran generic functions: example

```
module mpi_f08
```

```
  implicit none
```

```
  interface mpi_send ! Generic interface
```

```
    module procedure mpi_send_f08 ! Default API
```

```
    module procedure mpi_send_f08_x ! Big count API
```

```
  end interface mpi_send
```

contains

```
  subroutine mpi_send_f08(..., count, ...)
```

```
    INTEGER, INTENT(IN) :: count
```

```
    ...
```

```
  end subroutine mpi_send_f08
```

```
  subroutine mpi_send_f08_x(..., count, ...)
```

```
    INTEGER(KIND=MPI_COUNT_KIND), INTENT(IN) :: count
```

```
    ...
```

```
  end subroutine mpi_send_f08_x
```

```
end module mpi_f08
```

5c. Fortran generic functions: example

```
module mpi_f08

  implicit none
  interface mpi_send ! Generic interface
    module procedure mpi_send_f08 ! Default API
    module procedure mpi_send_f08_x ! Big count API
  end interface mpi_send

contains

  subroutine mpi_send_f08(..., count, ...)
    INTEGER, INTENT(IN) :: count
    ...
  end subroutine mpi_send_f08

  subroutine mpi_send_f08_x(..., count, ...)
    INTEGER(KIND=MPI_COUNT_KIND), INTENT(IN) :: count
    ...
  end subroutine mpi_send_f08_x

end module mpi_f08
```

```
program main
  use mpi_f08
  implicit none

  integer :: i = 100
  integer(kind=MPI_COUNT_KIND) :: bigI = 9876543210
  integer, parameter :: SMALLINT = 100
  integer(kind=MPI_COUNT_KIND), parameter :: BIGINT = 9876543210

  call mpi_send(..., i, ...)
  call mpi_send(..., SMALLINT, ...)
  call mpi_send(..., 100, ...)

  call mpi_send(..., bigI, ...)
  call mpi_send(..., BIGINT, ...)
  call mpi_send(..., 9876543210, ...)

  call mpi_send(..., INT(100,KIND=MPI_COUNT_KIND), ...)
  call mpi_send(..., INT(9876543210,KIND=MPI_COUNT_KIND), ...)

end program main
```

There is a long-term goal beyond MPI-4




Fixing BigCount via
`_Generic` / function overloading
is a good first step








But we envision a future with MORE



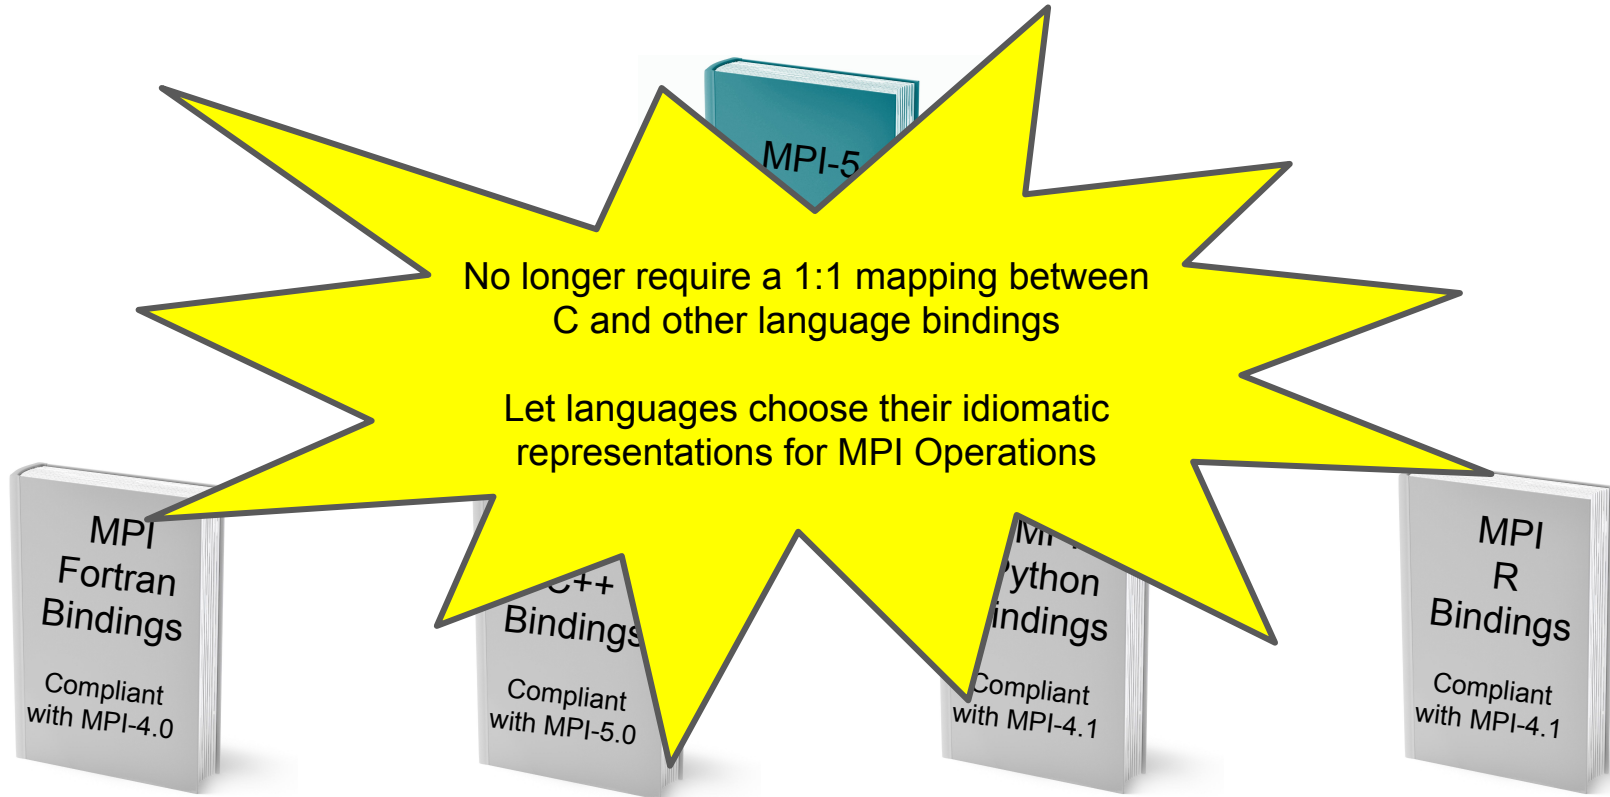
Perhaps something like this:



Perhaps something like this:



Perhaps something like this:



More cool, color-coded books
to publish!



Which option does the Forum want us to pursue?

	C	C++	Fortran
1.		No change / keep all counts as <code>int</code>	
2.		Change all counts to <code>MPI_Count</code>	
3.	“_X” functions	“_X” functions	“_X” functions
4.	Function pointers	Function pointers (?)	Objects
5.	C11 <code>_Generic</code>	Function overloading	Generic functions
6.	Do nothing	Function overloading	Generic functions

Straw poll results

Which way should we go?

- 0: Do nothing / leave int
- 1: Change everything to MPI_Count
- 0: _X functions
- 0: Chattanooga-style function pointers
- 21: (C11 _Generic, C++, Fortran)
- 1: only (C++, Fortran)

Should we add MPI_Count to the “mpi” module or not?

- Yes: 4
- No: 10
- Abstain: 6

Should we make <mpi.hpp> C++ header file?

...we forgot to poll this.

Just in case you wanted to know the differences

	C	C++	Fortran
Mechanism	<code>#define</code> <code>_Generic</code>	overloading, templates	interface
Naming	manual	automagic	manual
Time	compile	compile	compile
Switch	expressions	functions/classes	functions/subroutines
Default	yes	no	no
Glue	manual	automagic	manual
Runtime	no	runtime polymorphism (virtual functions)	no